

Chapter 1: Introduction

Chapter 13: I/O structure

Prof. Li-Pin Chang
CS@NYCU

WHAT IS AN OPERATING SYSTEM?

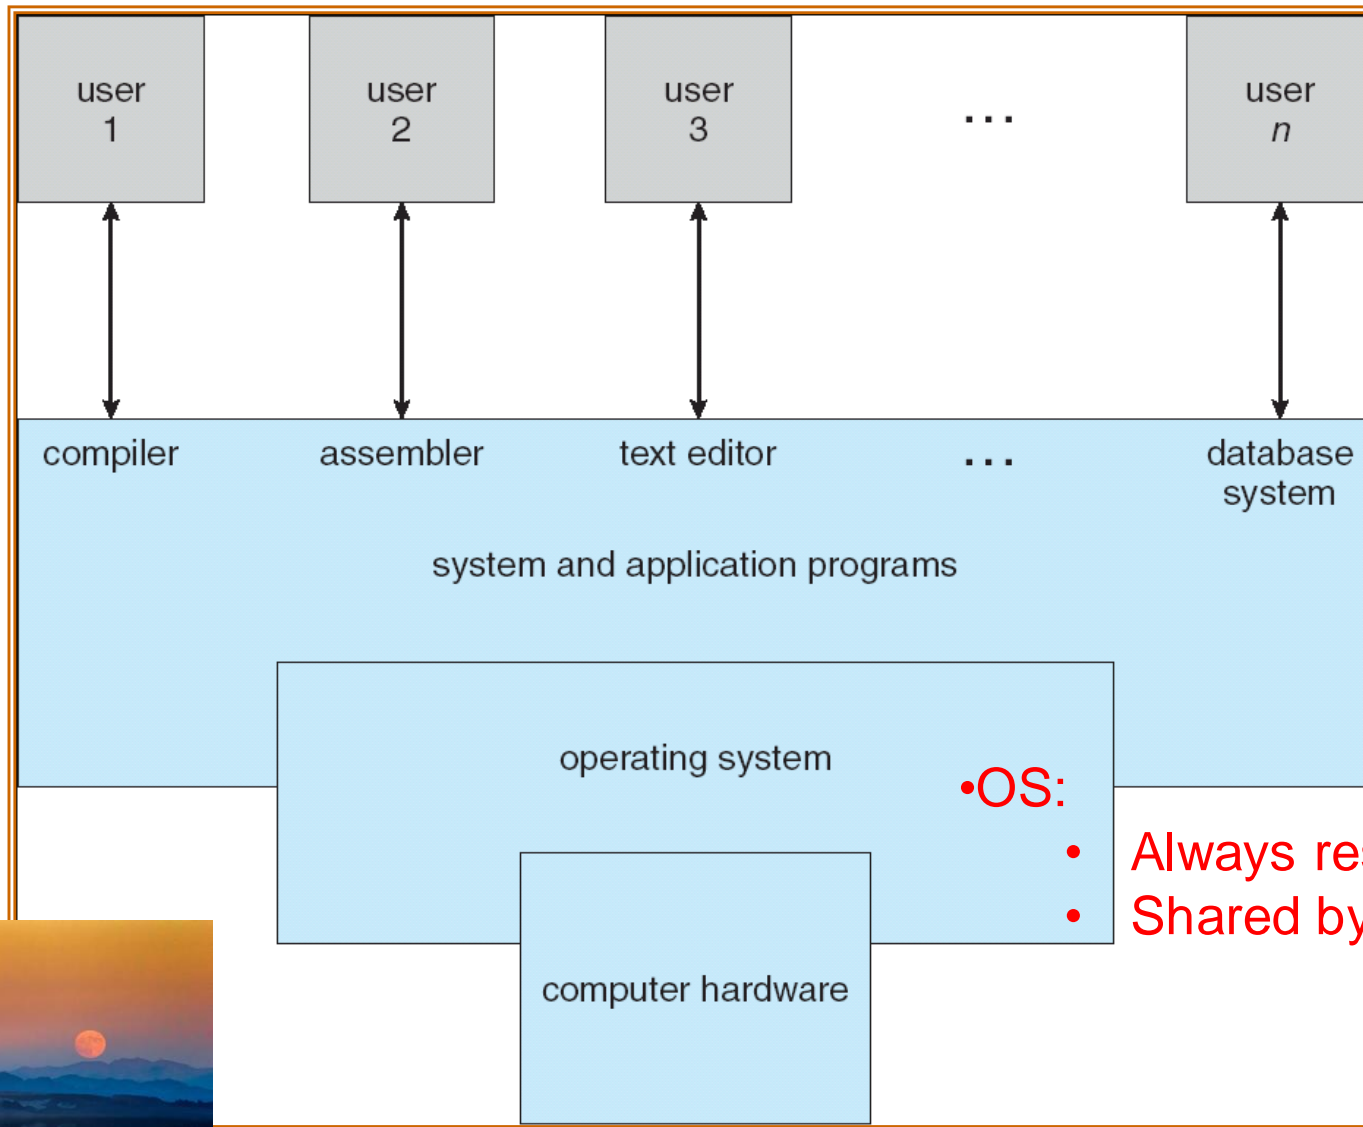
What is an Operating System?

- A program that acts as an intermediary between a user of a computer and the computer hardware
- Operating system goals:
 - **Execute** user programs and make solving user problems easier
 - Make the computer system **convenient** to use
 - Use the computer hardware in an **efficient** manner

Computer System Structure

- Hardware – provides basic computing resources
 - CPU, memory, I/O devices
- Operating system
 - Controls and coordinates use of hardware among various applications and users
- App programs – define the ways in which the system resources are used to solve the computing problems of the users
 - Office suite, compilers, web browsers, database systems, games
- Users
 - Human beings; some systems involve no human (M2M)

Four Components of a Computer System



•OS:

- Always resident
- Shared by users



Operating System Definition

- OS is a resource allocator
 - Manages all resources
 - Decides between conflicting requests for **efficient** and **fair** resource use
- OS is a control program
 - Controls (or monitor) execution of programs to **prevent errors and improper use** of the computer

Operating System Definition (Cont'd)

- No black-and-white definition
 - *“Everything a vendor ships when you order an operating system”* is good approximation
 - Err... how about the browser?
- “The one program running at all times on the computer” is the **kernel**
 - Everything else is either a **system program** or an **application program**
- Example: the UNIX OS
 - POSIX (IEEE 1003.1-2017)
 - A kernel
 - A collection of system programs \sim coreutil + binutil

- Which one(s) of the following are part of operating systems?
 - a) The CPU scheduler
 - b) Device drivers
 - c) Compilers
 - d) Word processors

THE THREE PILLARS

Pillar 1: Process Management

- A process is a program in execution; It is a unit of work within the system; Program is a **passive** entity, process is an **active** entity
- Process needs resources to accomplish its task
 - CPU time for execution
 - I/O for communication
 - Files for data storage
- A system has many processes running concurrently on one or more CPUs

Process Management Activities

- The operating system is responsible to the following activities in connection with process management:
 - Creating and deleting processes
 - Suspending and resuming processes
 - Providing mechanisms for process communication
 - Providing mechanisms for process synchronization

Pillar 2: Memory Management

- All **data** in memory before and after processing
- All **instructions** in memory in order to execute
- Memory management activities
 - **Allocating** and **deallocating** memory space as needed
 - Deciding which processes (or parts thereof) and data to move into and out of memory (**swapping or paging**)
 - Allowing programs allocate much more memory than physical memory (**virtual memory**)
 - Keeping track of which parts of memory are currently being used and by whom (**protection**)

Pillar 3: File/Storage Management

- OS provides uniform, logical view of information storage
 - Abstracts physical properties to logical storage unit - file
- File-System management
 - Files usually organized into directories
 - Access control on most systems to determine who can access what
- File system basic activities
 - **Creating** and **deleting** files and directories
 - **Reading** and **writing** files and directories
 - **Mapping** files from secondary storage to main memory

Storage Management

- Entire speed of computer operation hinges on disk subsystem and its algorithms
- OS activities
 - Storage allocation
 - Free-space management
 - Disk scheduling
- Must deal with medium properties
 - Tape: sequential access
 - Disks: disk head movement
 - SSDs: garbage collection and endurance

I/O STRUCTURE

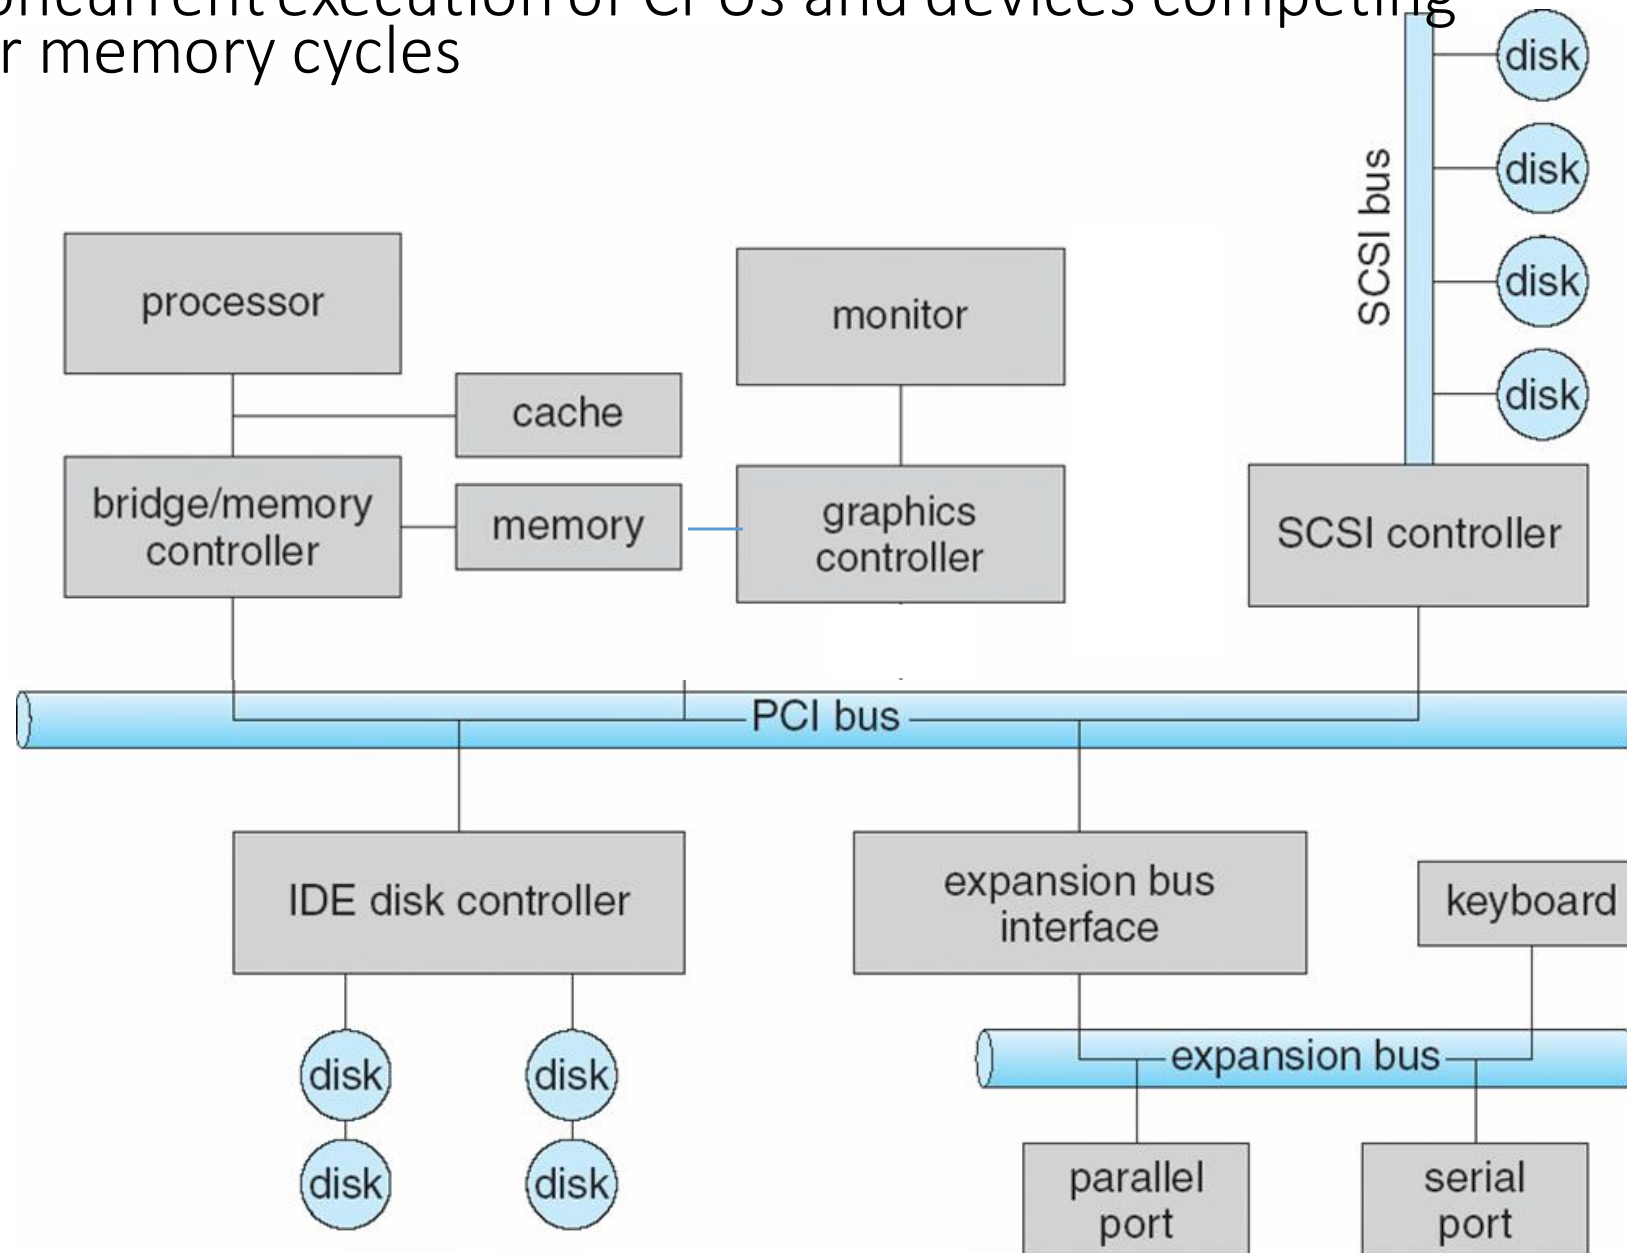
Why Bothering with I/Os?

- Operating system is called synchronously by user programs, but it is also often triggered by asynchronous events, i.e., hardware events from outside of the CPU
- User programs run on the CPU in parallel with operations on I/O devices
- OS strives to optimize **CPU utilization** and **I/O latency** at the same time, and the optimization is closely related to how I/O devices behave

Before OS: Computer Startup


- Bootstrap program is loaded at power-up or reboot
 - Typically stored in ROM or EEPROM, generally known as firmware
 - Initializes all aspects of system
 - Loads operating system kernel and starts execution
- Legacy PC bootloading procedure
 - FFFF:0000
 - BIOS
 - MBR (Master Boot Record)
 - Boot Manager
 - Pre-OS
 - OS

- One or more CPUs, device controllers connect through common bus providing access to shared memory
- Concurrent execution of CPUs and devices competing for memory cycles




Typical PC Organization

- North bridge (memory hub)
 - CPU
 - Memory
 - High-speed PCIe devices, like graphics, NVMe SSDs
- South bridge (IO hub)
 - USB
 - SCSI
 - SATA
 - ...



I/O Hardware

- Incredible variety of I/O devices
- Common concepts
 - Host controller
 - Bus or interface
 - Device controller
- I/O instructions control devices
 - Direct I/O
 - Memory-mapped I/O



Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)

Example: Set the x86 Interval Timer

```
mov al, 0x36
out 0x43, al    ;tell the PIT which channel we're setting

mov ax, 11931
out 0x40, al    ;send low byte
mov al, ah
out 0x40, al    ;send high byte
```


Channel 0	0x40	System Timer
Channel 1	0x41	DRAM refresh (obsolete)
Channel 2	0x42	Buzzer
Command	0x43	Command I/O register

I/O Operation

- I/O devices and CPUs execute concurrently
 - **Polling** I/O: The CPU waits on an I/O device until completion
 - **Interrupt** I/O: The CPU will be notified of I/O completion

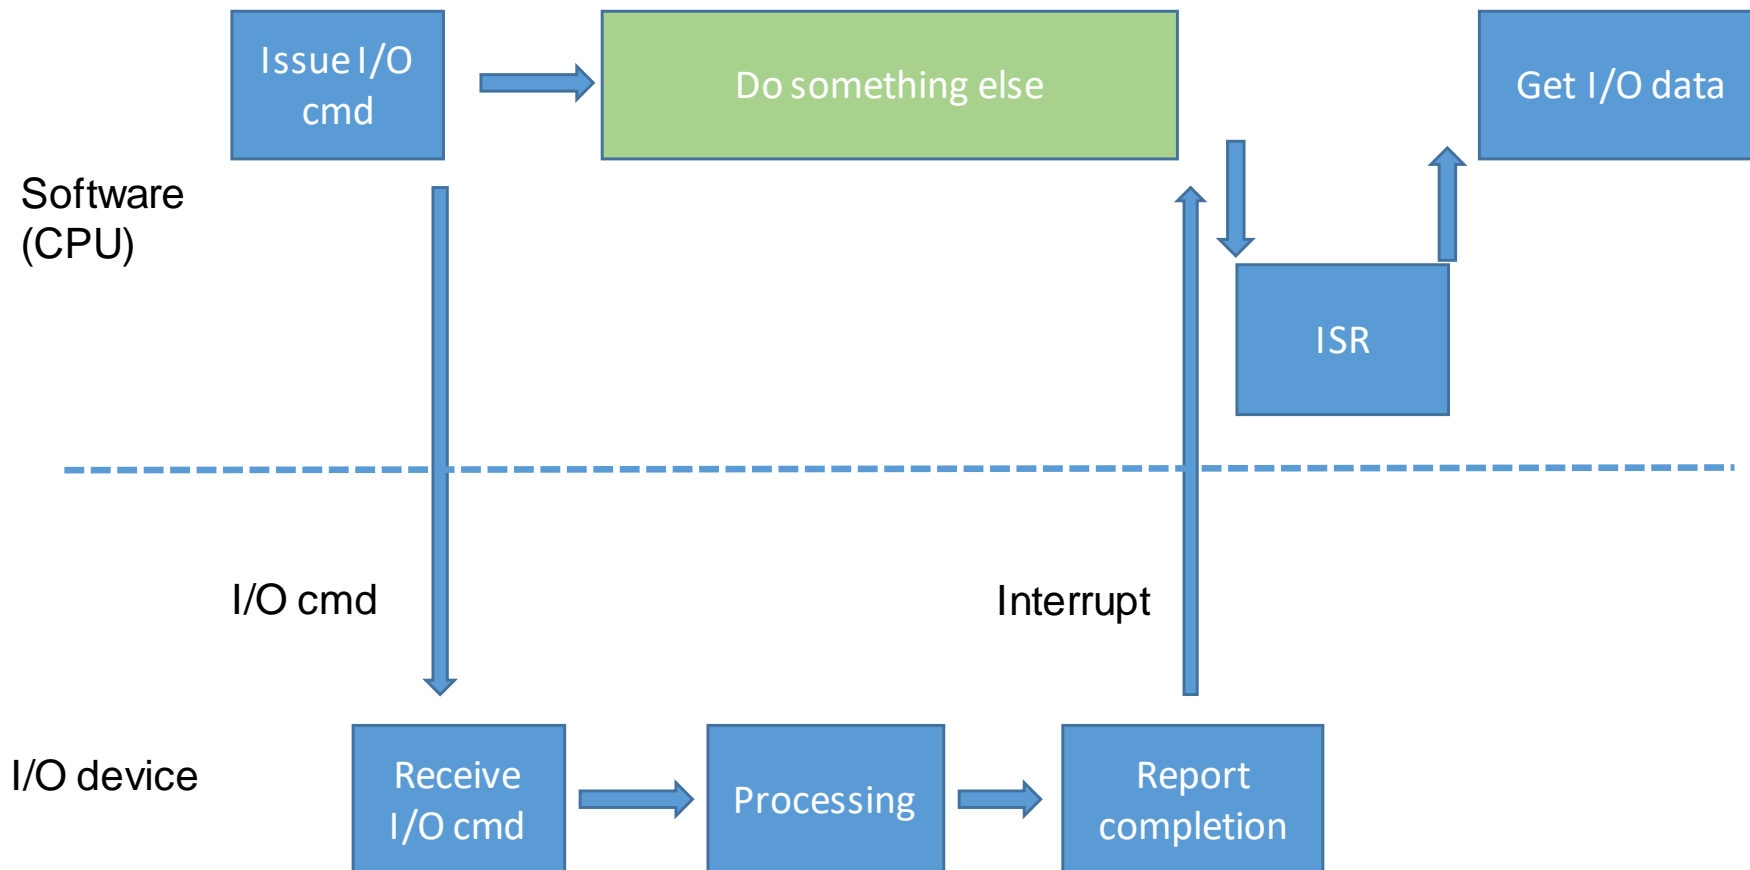
Polling

- CPU determines state of device
 - Busy
 - Data ready
 - Error
- Busy-wait cycle to wait for I/O from device



Interrupts

- CPU Interrupt-request line triggered by I/O device
- Interrupt handler (ISR) receives interrupts
- Interrupt vector to dispatch interrupt to correct handler
- Interrupt mechanism also used for exceptions



Common Functions of Interrupts


- Interrupt transfers control to the interrupt service routine (ISR) generally, through the **interrupt vector**, which contains the addresses of all the service routines
 - A **trap** is a software-generated interrupt caused either by an error or a user request
 - An **IRQ** is generated by hardware
- **An operating system is interrupt driven**
- **Hardware:** interrupt request (IRQ)
 - I/O completion, timer, etc
- **Software:** Trap
 - divide by zero, access violation, system call

Legacy PC interrupts

Interrupt	Function
00h	Divide By Zero
01h	Single Step
02h	Nonmaskable Interrupt (NMI)
03h	Breakpoint Instruction
04h	Overflow Instruction
05h	Print Screen
05h	Bounds Exception (80286, 80386)
06h	Invalid Op Code (80286, 80386)
07h	Math Coprocessor Not Present
08h	Double Exception Error (80286, 80386) (AT Only)
08h	System Timer - IRQ 0
09h	Keyboard - IRQ 1
09h	Math Coprocessor Segment Overrun (80286, 80386) (AT Only)
0Ah	IRQ 2 - Cascade from Second programmable Interrupt Controller
0Ah	Invalid Task Segment State (80286, 80286) (AT Only)
0Ah	IRQ 2 - General Adapter Use (PC Only)
0Bh	IRQ 3 - Serial Communications (COM 2)
0Bh	Segment Not Present (80286, 80386)
0Ch	IRQ 4 - Serial Communications (COM 1)
0Ch	Stack Segment Overflow (80286, 80386)
0Dh	Parallel Printer (LPT 2) (AT Only)
0Dh	IRQ 5 - Fixed Disk (XT Only)
0Dh	General Protection Fault (80286, 80386)
0Eh	IRQ 6- Diskette Drive Controller
0Eh	Page Fault (80386 Only)
0Fh	IRQ 7 - Parallel printer (LPT 1)
⋮	
⋮	
⋮	
⋮	

Interrupt Handling

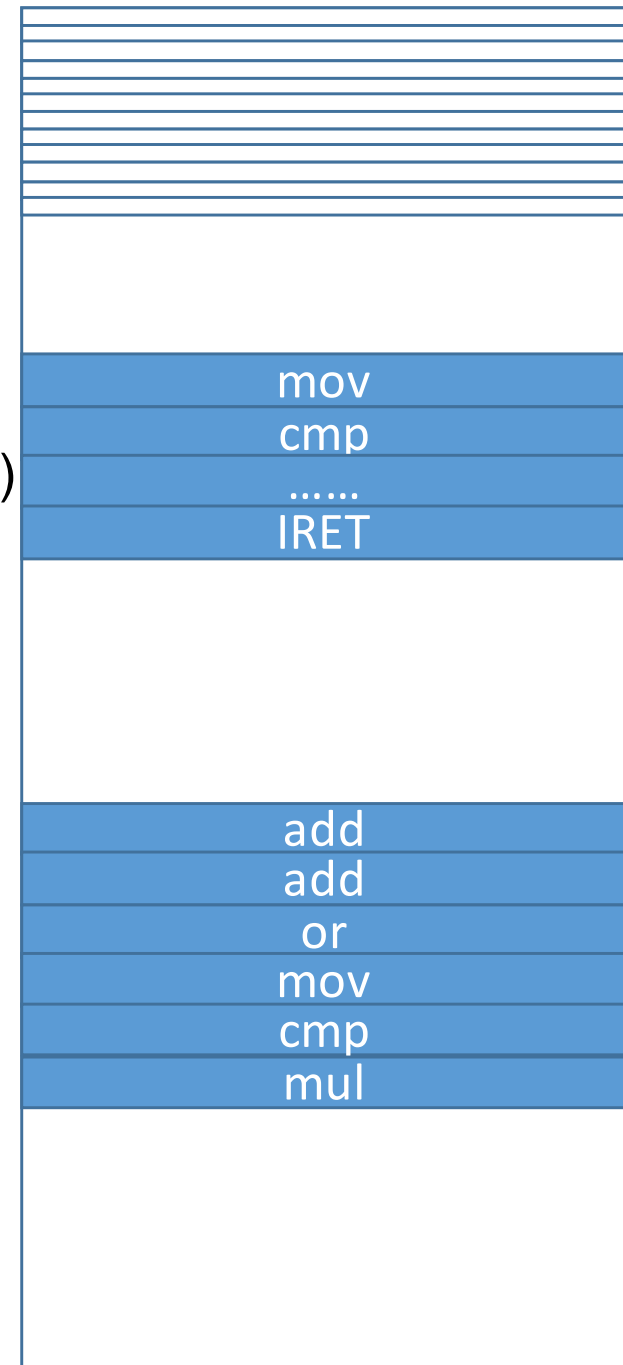
- The operating system preserves the state of the CPU by storing **registers** and the **program counter**
- Determines which type of interrupt has occurred:
 - Reading I/O registers
 - vectored interrupt system
- Separate segments of code determine what action should be taken for each type of interrupt




IVT

ISR
(part of OS)

User
program




Interrupt Timeline

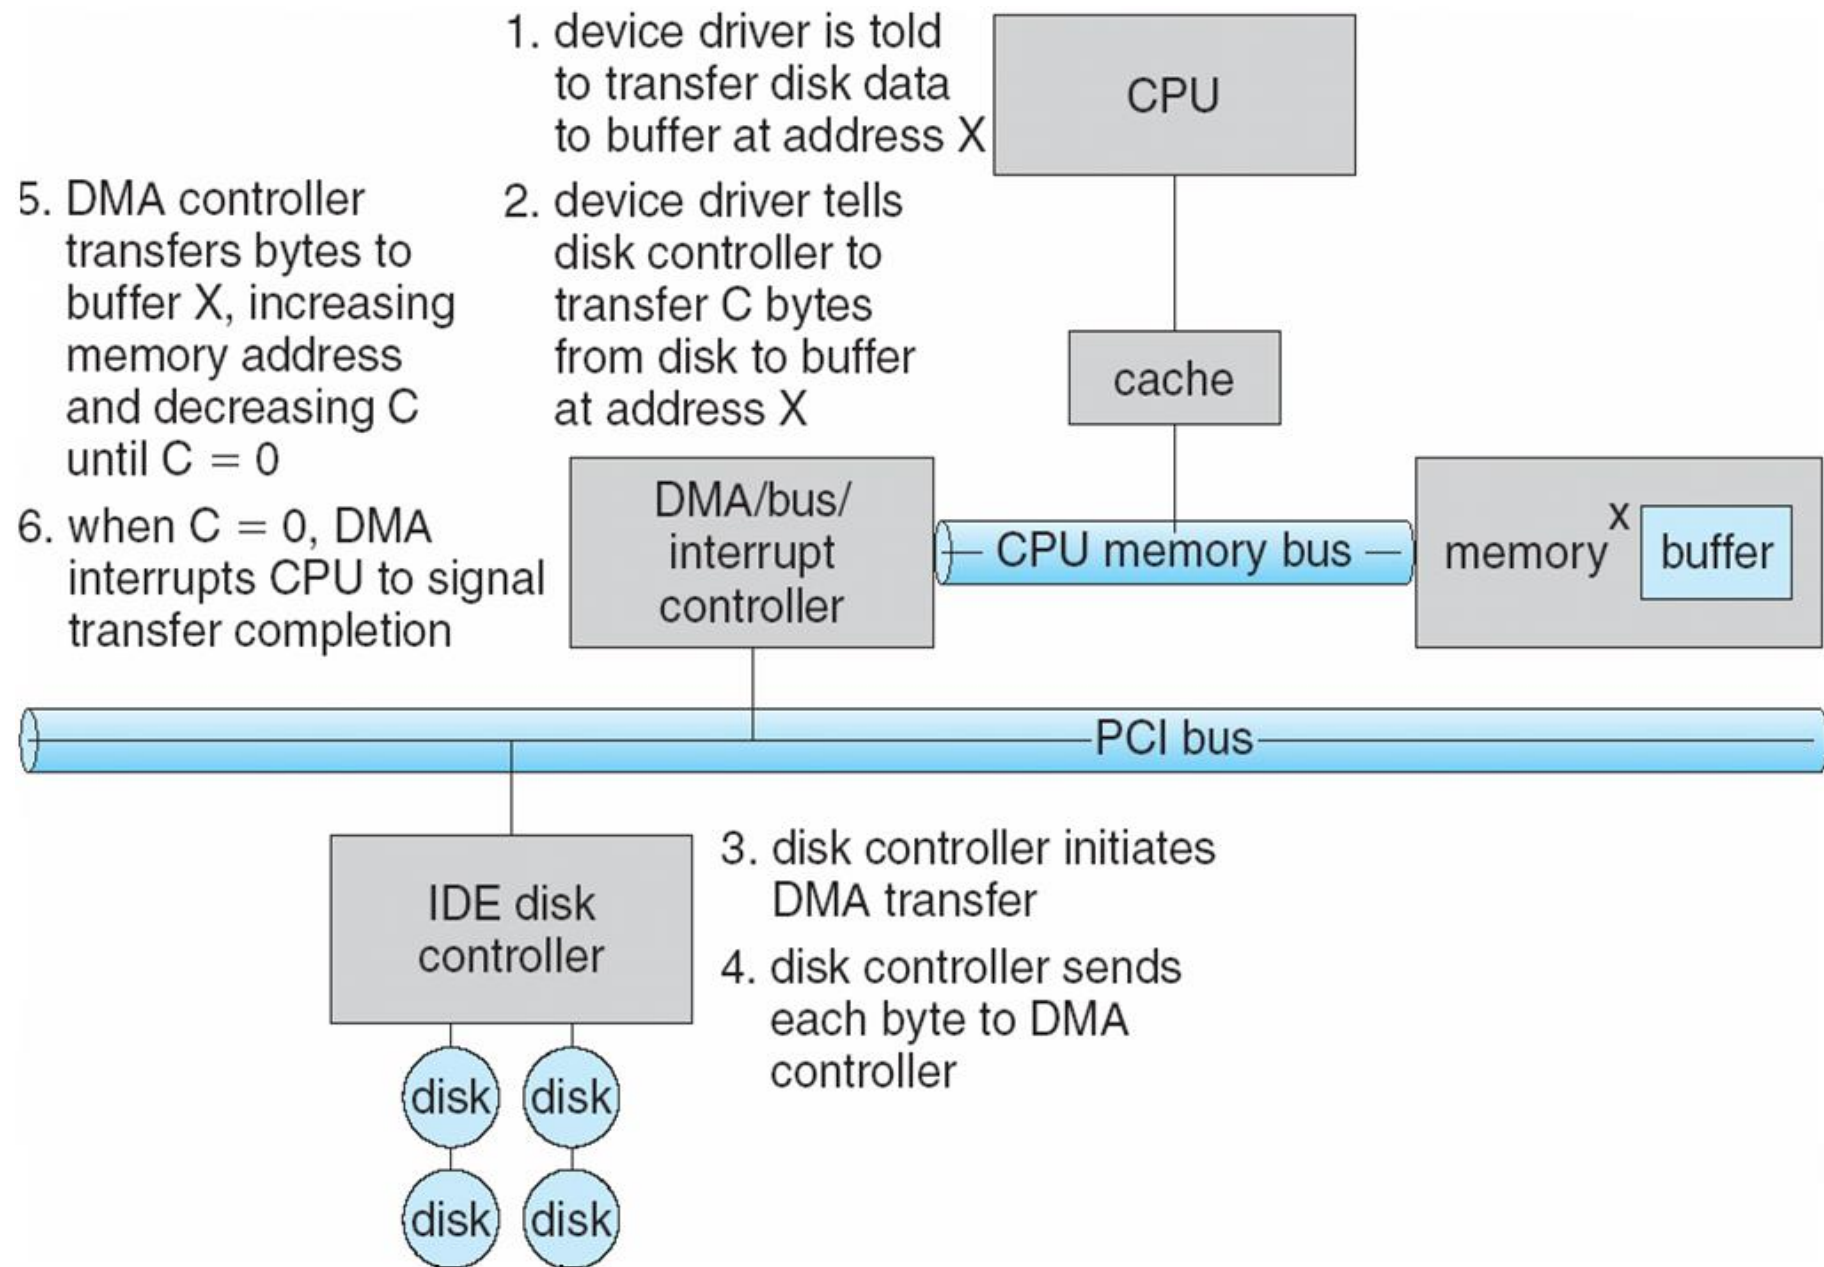


Direct Memory Access

- Data movement of I/O operation
 - Each device controller has a local buffer
 - Option 1: CPU moves data between main memory and I/O local buffers
 - Option 2: CPU offloads data movement to **DMA**
- DMA: Used for high-speed I/O devices able to transmit information at close to memory speeds
- Device controller transfers **blocks of data** from buffer storage directly to main memory without CPU intervention
- Only one interrupt is generated per block, rather than the one interrupt per byte

105





Synchronous I/O vs. Asynchronous I/O

- After I/O starts, control returns to the program only upon I/O completion (**blocking call; sync I/O**)
 - Example: regular I/O reading
`read()` or `fread()`
- After I/O starts, control returns to the program without waiting for I/O completion (**non-blocking call; async I/O**)
 - Example: regular I/O writing
`write()` or `fwrite()`

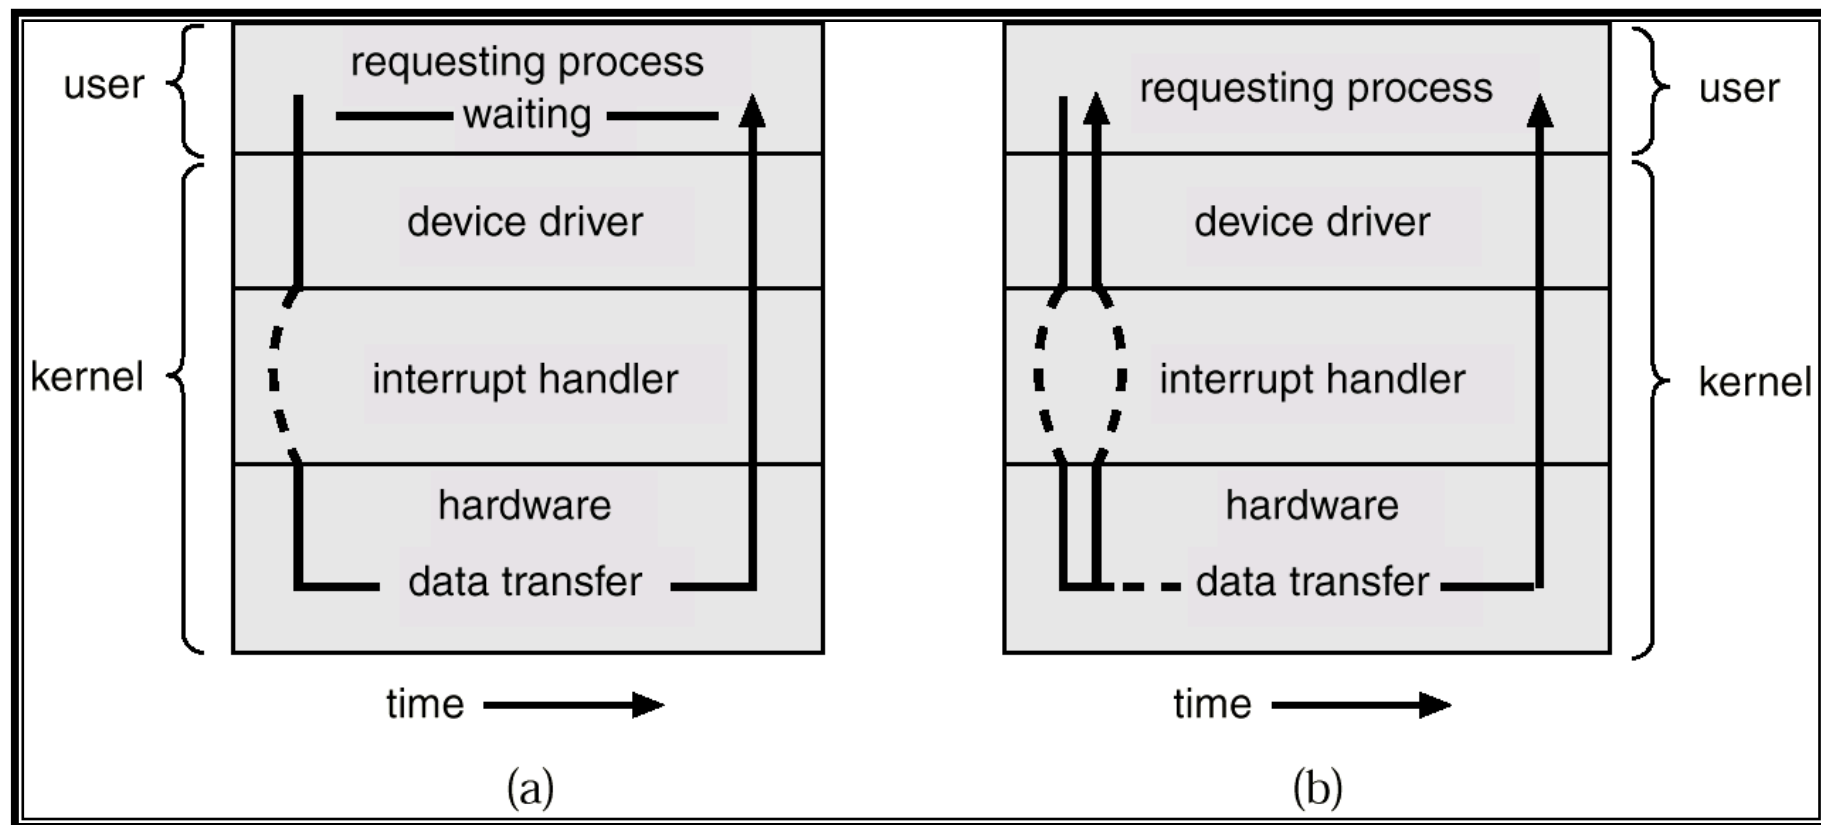
I/O Structure (Async vs Sync)

Synchronous

Blocking I/O

Asynchronous


Non-blocking I/O



- Are the following operations synchronous or asynchronous?
 - `fread()`
 - `fwrite()`
 - `fsync()`
 - `aio_read()`

Multiprogramming

A Possible Memory Layout of Multiprogramming Systems



Multiprogramming

- Single process cannot keep CPU and I/O devices busy at all times
- Multiprogramming organizes processes so CPU always has one to execute
- A subset of all processes in system is kept in memory
- One process selected and run via process scheduling
- When it has to wait (for I/O for example), OS switches to another process

Timesharing

- Timesharing is logical extension in which **CPU switches processes so frequently** that users can interact with each job while it is running, creating **interactive computing**
 - Switching frequency usually ≥ 100 Hz, driven by timer IRQ
 - Each user has at least one program executing in memory
⇒ process (ch3)
 - If several jobs ready to run at the same time ⇒ CPU scheduling (ch5)
 - If processes don't fit in memory, swapping moves them in and out to run (ch5)
 - Virtual memory allows execution of processes not completely in memory (ch8)

Know the Terminology

- Multiprogramming
 - multiple processes are loaded into memory for execution
- Multitasking
 - Multiprogramming with overlapped task execution
- Timesharing
 - multitasking + periodic switch among processes

- A _____ can be used to prevent a user program from never returning control to the operating system

A) portal

B) program counter

C) firewall

D) timer

End of Chapter 1