

Chapter 2: System Structures

Prof. Li-Pin Chang
CS@NYCU

Chapter 2: System Structures

- Operating System Services
 - User Interfaces
 - System Programs
 - **System Calls** and types of System Calls
- Operating System Design and Implementation
 - **Operating System Structure**
 - Virtual Machines

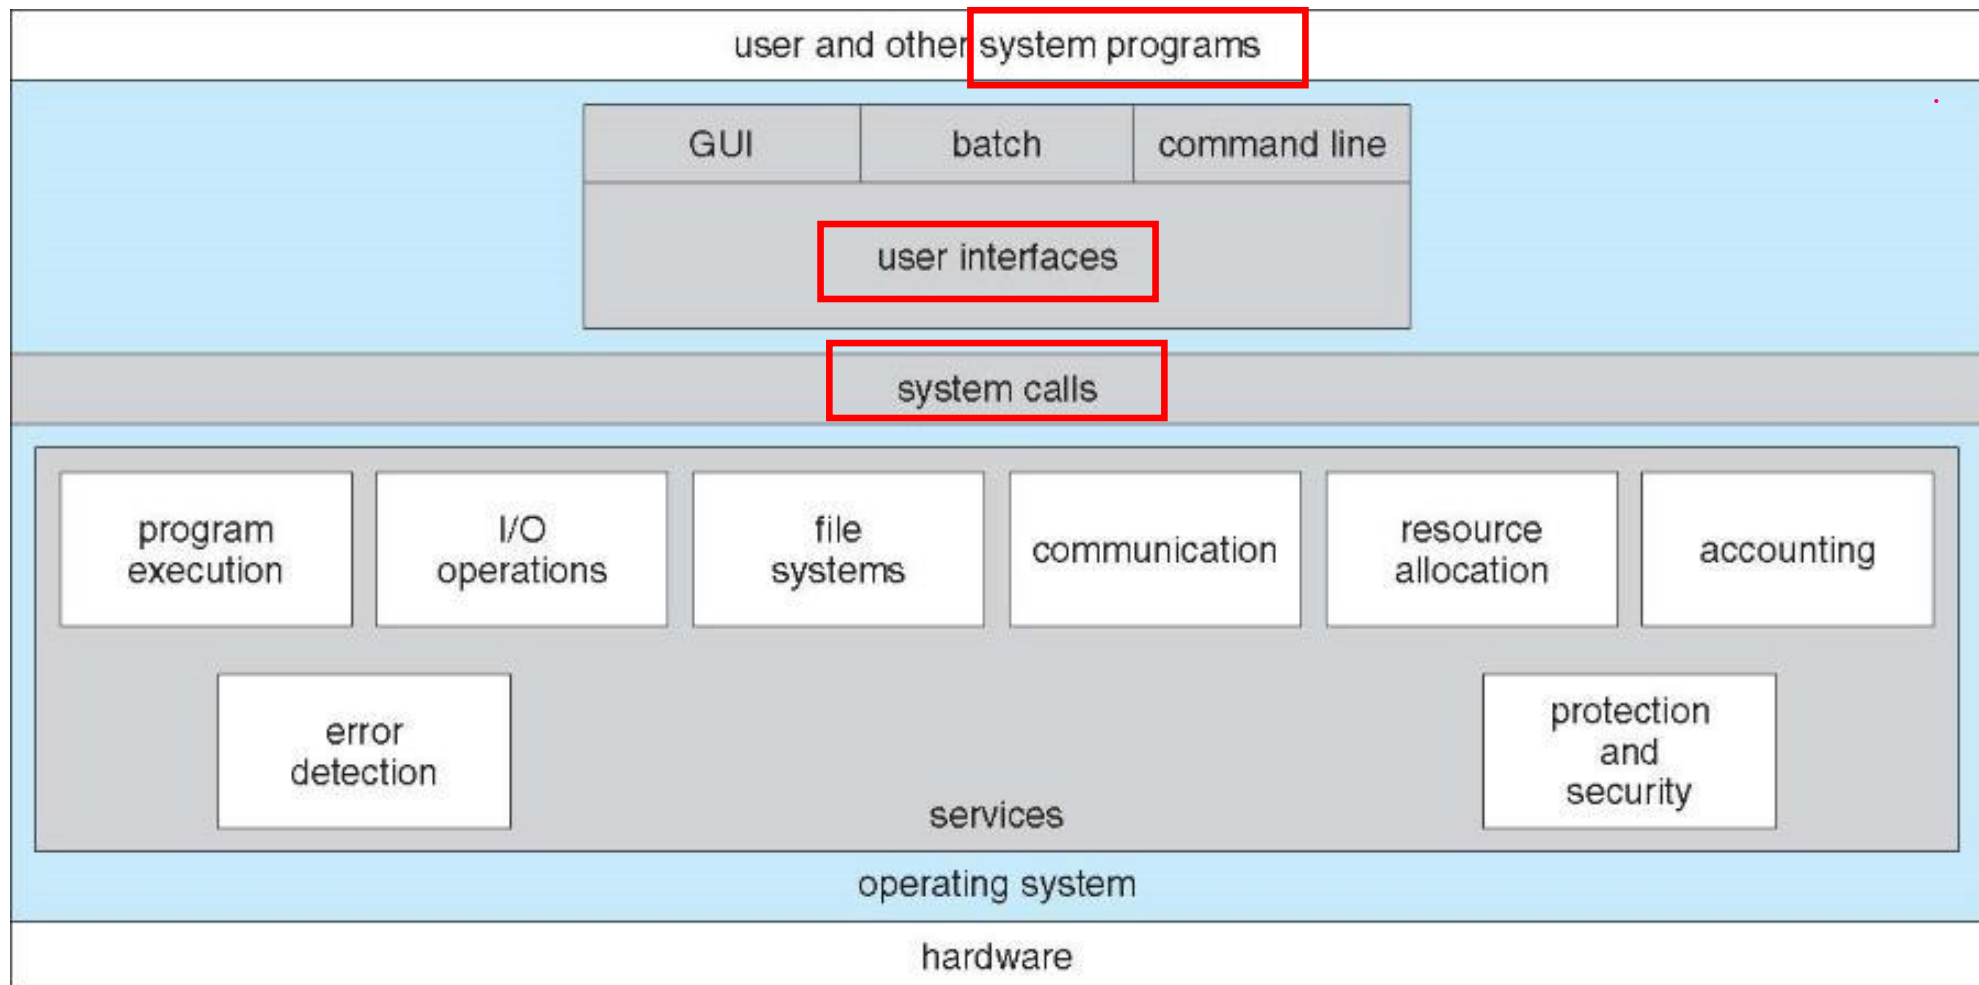
Objectives

- To describe the services an operating system provided to users, processes, and other systems
- To discuss the various ways of structuring an operating system

OPERATING SYSTEM SERVICES

A View of Operating System **Services**

3 levels of services



Operating System Services

- One set of operating-system services provides **functions** that are helpful to the user:
 - **User interface** - Almost all operating systems have a user interface (UI)
 - Varies between Command-Line (CLI), Graphics User Interface (GUI), Batch
 - **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
 - **I/O operations** - A running program may require I/O, which may involve a file or an I/O device.
 - **File-system manipulation** - The file system is of particular interest. Obviously, programs need to read and write files and directories, create and delete them, search them, list file information, permission management.

Operating System Services (Cont.)

- One set of operating-system services provides **functions** that are helpful to the user (Cont):
 - **Communications** – Processes may exchange information, on the same computer or between computers over a network
 - Communications may be via shared memory or through message passing (packets moved by the OS)
 - **Error detection** – OS needs to be constantly aware of possible errors
 - May occur in the CPU and memory hardware, in I/O devices, in user program
 - For each type of error, OS should take the appropriate action to ensure correct and consistent computing
 - Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system

Operating System Services (Cont.)


- Another set of OS functions exists for ensuring the **efficient operation** of the system itself via resource sharing
 - Resource **allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
 - Many types of resources - Some (such as CPU cycles, main memory, and file storage) may have special allocation code, others (such as I/O devices) may have general request and release code.
 - **Accounting** - To keep track of which users use how much and what kinds of computer resources
 - **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
 - Protection involves ensuring that all access to system resources is controlled
 - Security of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts
 - If a system is to be protected and secure, precautions must be instituted throughout it. A chain is only as strong as its weakest link.

SYSTEM PROGRAMS

System Programs

- System programs provide a convenient environment for program development and execution. They can be divided into:
 - File manipulation (cp, mv...)
 - Status information (ls...)
 - File modification (vi...)
 - Programming language support (cc, as, ld, ar...)
 - Program loading and execution
 - Communications (telnet...)
- Most users' view of the operation system is defined by system programs, not the actual system calls

GNU coreutils + binutils



System Programs

- Provide a convenient environment for program development and execution
 - Some of them are simply user interfaces to system calls; others are considerably more complex
- File management - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories
- Programming-language support - **Compilers, assemblers, debuggers** and interpreters sometimes provided
- Program loading and execution- Absolute loaders, relocatable loaders, linkage editors, and overlay-loaders, debugging systems for higher-level and machine language

USER INTERFACE

Operating-System User Interface - CLI

- **Shell** refers to the interface program between users and the kernel
 - Text-driven: Command Line Interface (CLI)
 - Graphics-driven: Graphical User Interface (GUI)
- CLI allows direct command entry
 - Primarily fetches a command from user and executes it
 - Sometimes commands built-in, sometimes just names of programs
 - If the latter, adding new features doesn't require shell modification

CLI in Windows/Linux

```
C:\windows\system32>cd \

C:\>dir/w
磁碟區 C 中的磁碟是 TI30940600B
磁碟區序號: 0A90-10B7

C:\ 的目錄

[BaKoMa TeX]
[Intel]
[PerfLogs]
[Program Files (x86)]
[TOSHIBA]
[UTDService]
[EcpaComponent]
[LJP1100_P1560_P1600_Full_Solution]
[Program Files]
[SSUUpdater.log]
[Users]
[Windows]
1 個檔案 282 位元組
11 個目錄 55,639,674,880 位元組可用

C:\>_
```

```
Loading...

Welcome to JS/Linux (x86)

Use 'vlogin username' to connect to your account.
You can create a new account at https://vfsync.org/signup .
Use 'export_file filename' to export a file to your computer.
Imported files are written to the home directory.

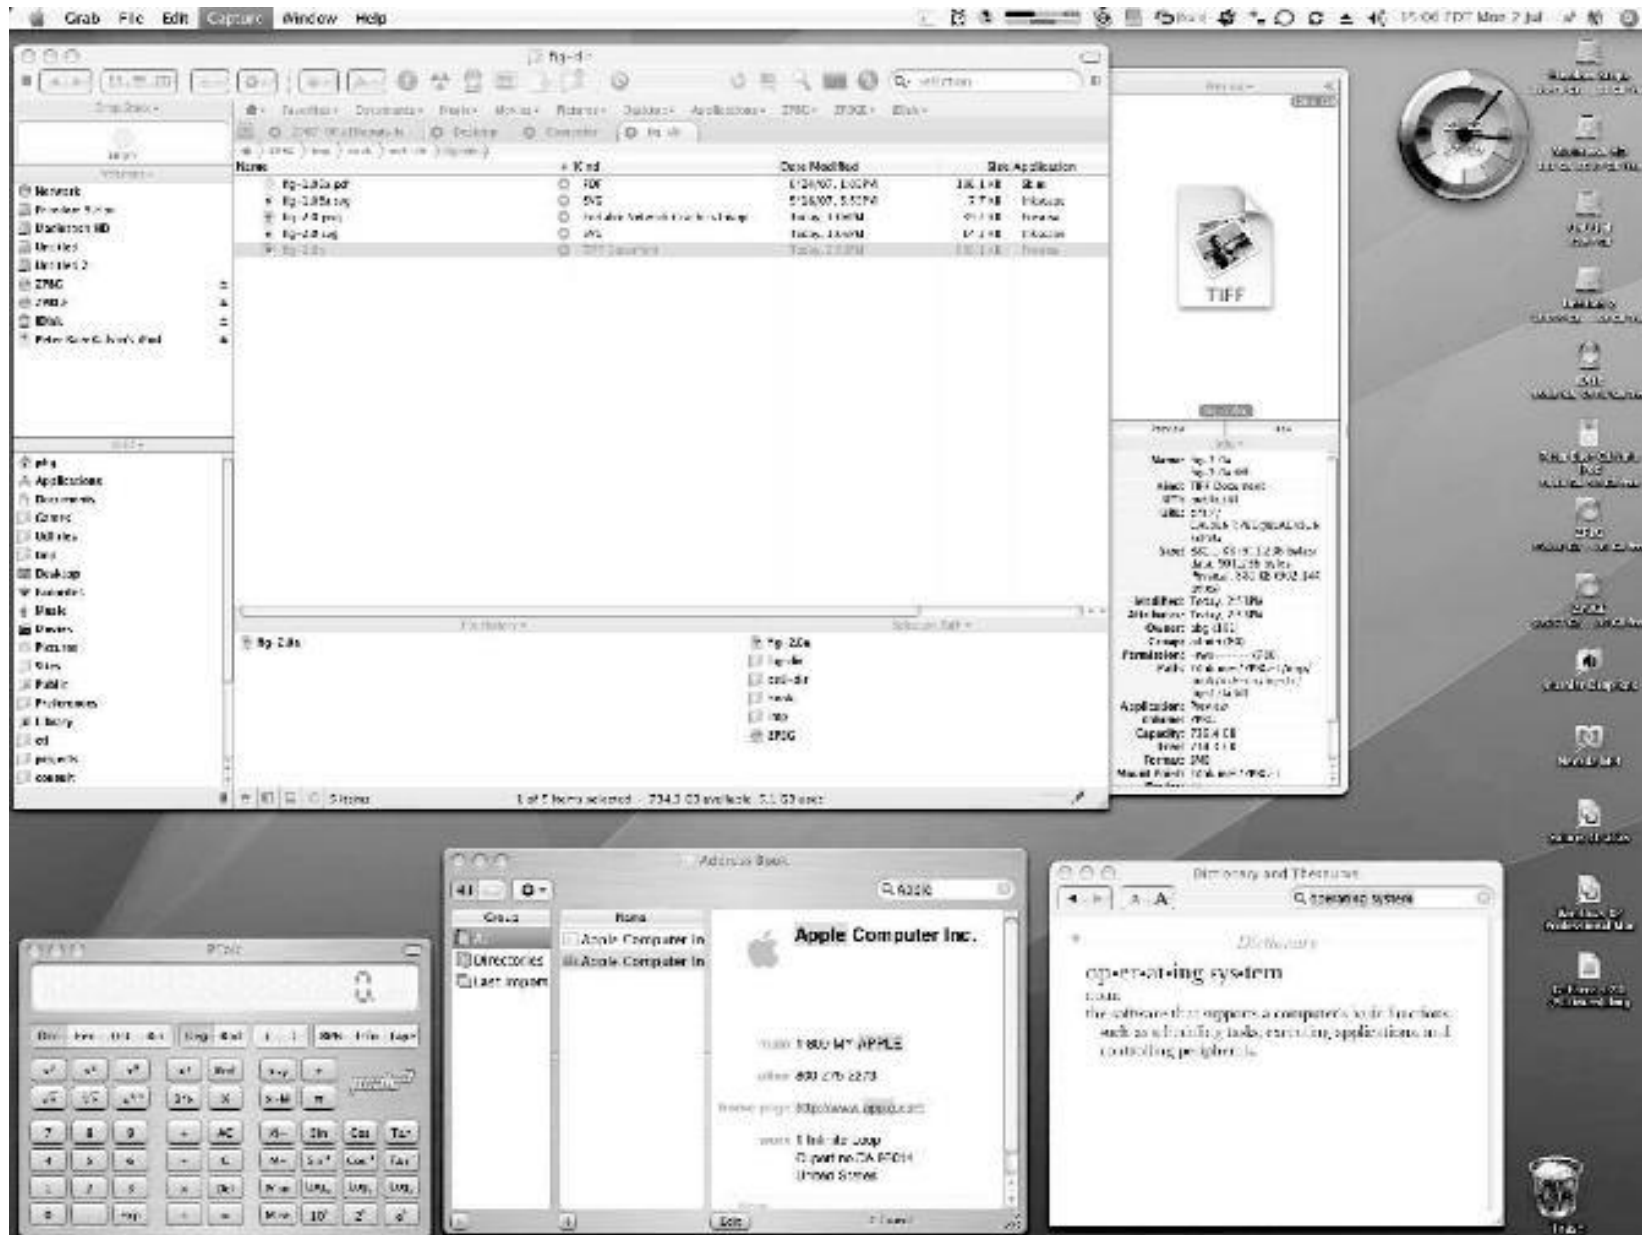
[root@localhost ~]# ls -l
total 8
drwxr-xr-x  3 root  root    163 Aug 21  2011 dos
-rw-r--r--  1 root  root    242 Jul 15  2017 hello.c
[root@localhost ~]# pwd
/root
[root@localhost ~]#
```

User Operating System Interface - GUI

- User-friendly desktop metaphor interface
 - Usually mouse, keyboard, and monitor
 - Icons represent files, programs, actions, etc
 - Invented by Xerox PARC
 - Sarcasm: **W**indow, **I**con, **M**enu, **P**ointer -- WIMP
- Many systems now include both CLI and GUI interfaces
 - Microsoft Windows is GUI with CLI “command” shell
 - Apple Mac OS X as “Aqua” GUI interface with UNIX kernel underneath and shells available
 - Solaris is CLI with optional GUI interfaces (Java Desktop, KDE)

PARC's accomplishments:
mouse, GUI, WYSIWYG editors, postscript, laser printers, Ethernet, small talk

The Mac OS X GUI



Touchscreen Interfaces


- Touchscreen devices require new interfaces
 - Mouse not possible or not desired
 - Actions and selection based on **gestures**
 - Virtual keyboard for text entry



SYSTEM CALLS

System Calls

- **Programming interface** to the services provided by the OS
- Mostly accessed by programs via a high-level Application Program Interface (**API**) rather than direct system call use
 - **Portability and simplicity**
- Three most common APIs are **Win32 API** for Windows, **POSIX API** for UNIX, and **Java API** for the Java VM



Standard C library: fopen("w+"...) C language



WIN32 API: CreateFile() Windows >= win4.0, >=95



Kernel API: NtCreateFile() WinNT, 2k, XP, vista




System Call: int 2e X86 machine instruction

```
int printf ( const char * format, ... );
```



```
BOOL WINAPI WriteFile(  
    _In_      HANDLE hFile,  
    _In_      LPCVOID lpBuffer,  
    _In_      DWORD nNumberOfBytesToWrite,  
    _Out_opt_ LPDWORD lpNumberOfBytesWritten,  
    _Inout_opt_ LPOVERLAPPED lpOverlapped  
);
```


Win32 API



```
NTSTATUS NtWriteFile  
  
(  
    HANDLE          hFile,  
    HANDLE          hEvent,  
    PIO_APC_ROUTINE apc,  
    void*           apc_user,  
    PIO_STATUS_BLOCK io_status,  
    const void*      buffer,  
    ULONG            length,  
    PLARGE_INTEGER   offset,  
    PULONG           key  
)
```

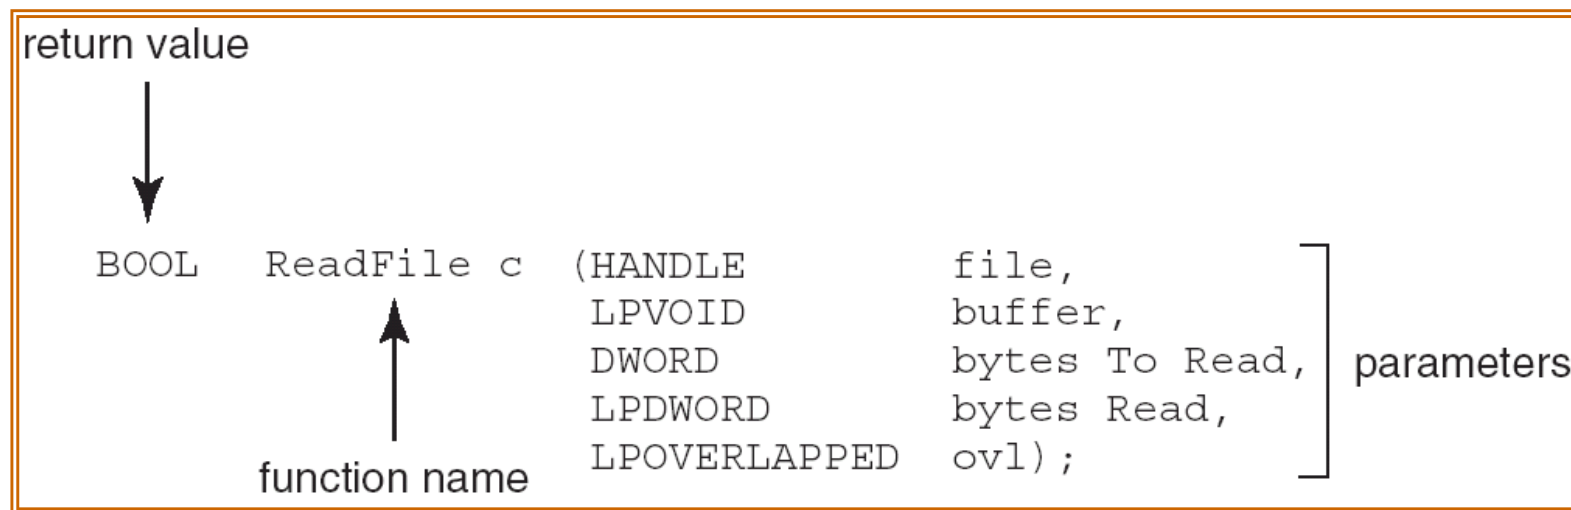
Kernel API

```
mov eax, <service #>  
lea edx, <addr of 1st arg>  
int 2e
```



Example of Standard API (Windows)

- Consider the ReadFile() function in the Win32 API—a function for reading from a file



- A description of the parameters passed to `ReadFile()`
 - `HANDLE file`—the file to be read
 - `LPVOID buffer`—a buffer where the data will be read into and written from
 - `DWORD bytesToRead`—the number of bytes to be read into the buffer
 - `LPDWORD bytesRead`—the number of bytes read during the last read
 - `LPOVERLAPPED ovl`—indicates if overlapped I/O is being used

Example of Standard API (UNIX)

EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the `man` page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

```
#include <unistd.h>

ssize_t  read(int fd, void *buf, size_t count)
```

return value	function name	parameters
-----------------	------------------	------------

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.

System Call Implementation

- Typically, **a number** associated with each system call
 - System-call interface maintains **a table** indexed according to these numbers
- The system call interface invokes intended system call in OS kernel and returns status of the system call and any return values
- The caller need know nothing about how the system call is implemented
 - Just needs to obey API and understand what OS will do as a result call

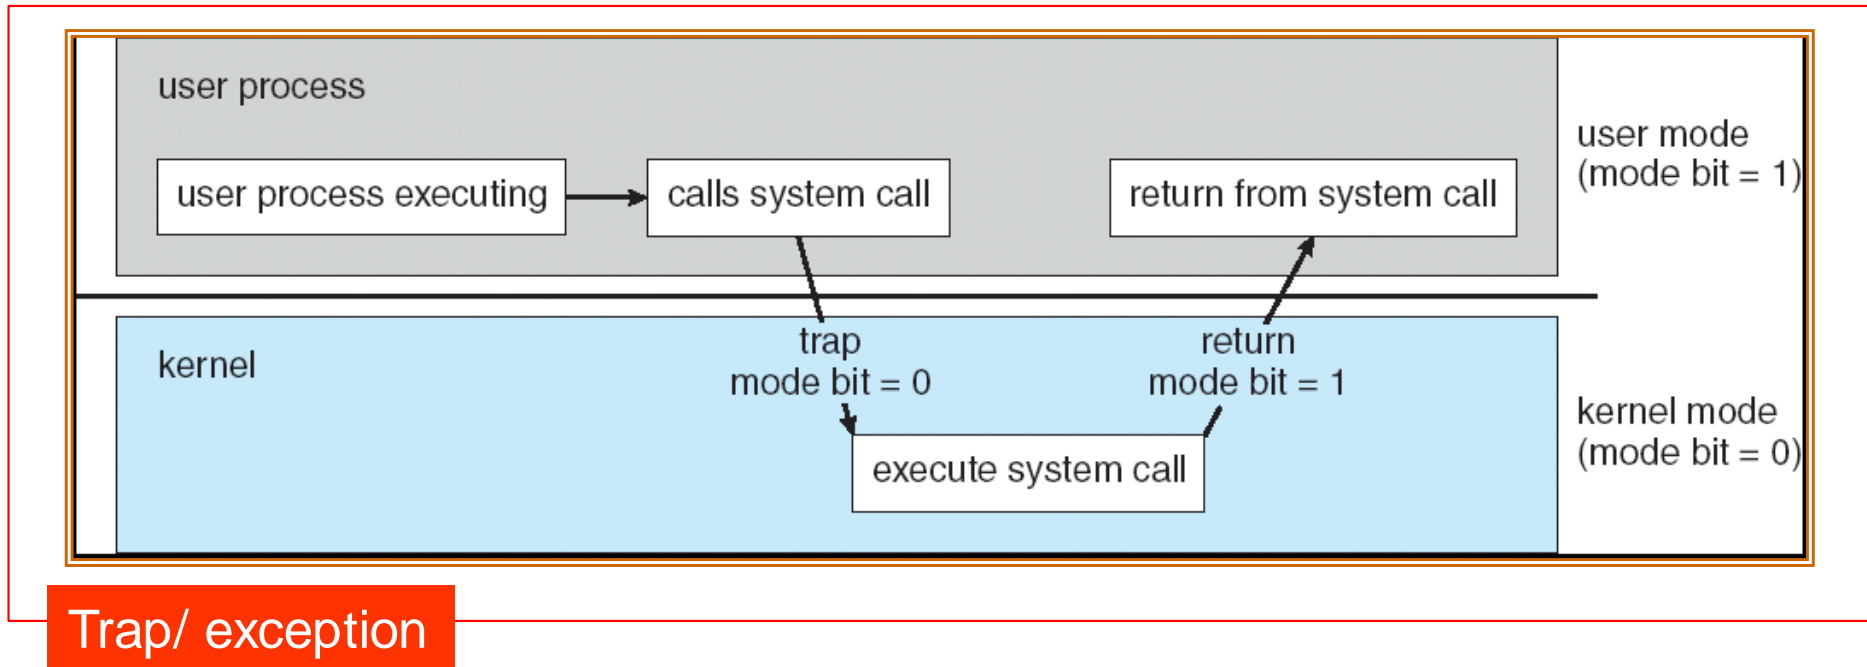
Dual Mode Operations

- Application calls into the kernel through **trap**
- **Interrupt** driven by hardware (IRQ)
- Software error or request creates **trap** or **exception**
 - Division by zero, memory access violation, etc
 - Request for operating system service (system calls)

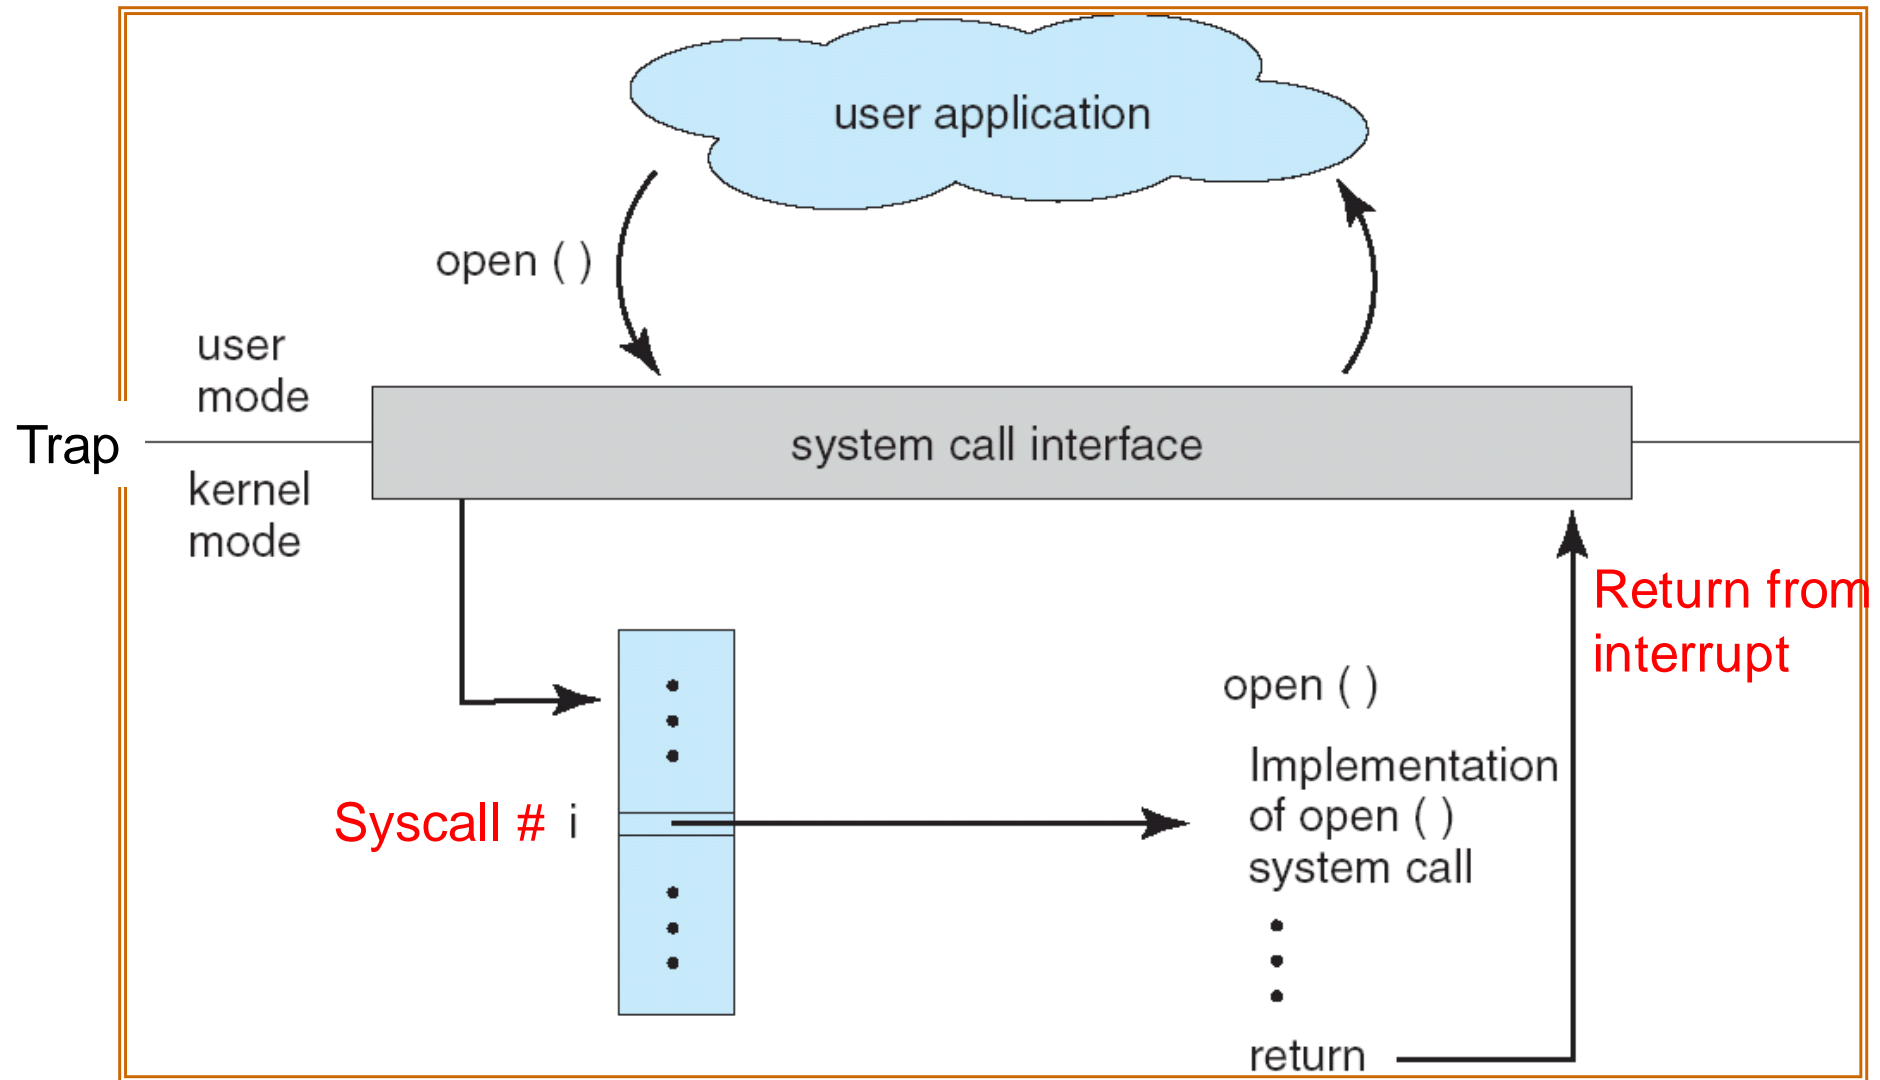
Dual Mode Operations

- **Dual-mode** operation, supported by CPU *hardware*, allows OS to protect itself against un-trusted code
 - **User mode (untrusted)** and **kernel mode (trusted)**
- The purpose of dual-mode design
 - Provides ability to distinguish when system is running user code or kernel code
 - Some instructions designated as **privileged**, only executable in kernel mode
 - System call changes mode to kernel, return from call resets it to user (mode bit supported by the CPU)

Transition from User to Kernel Mode

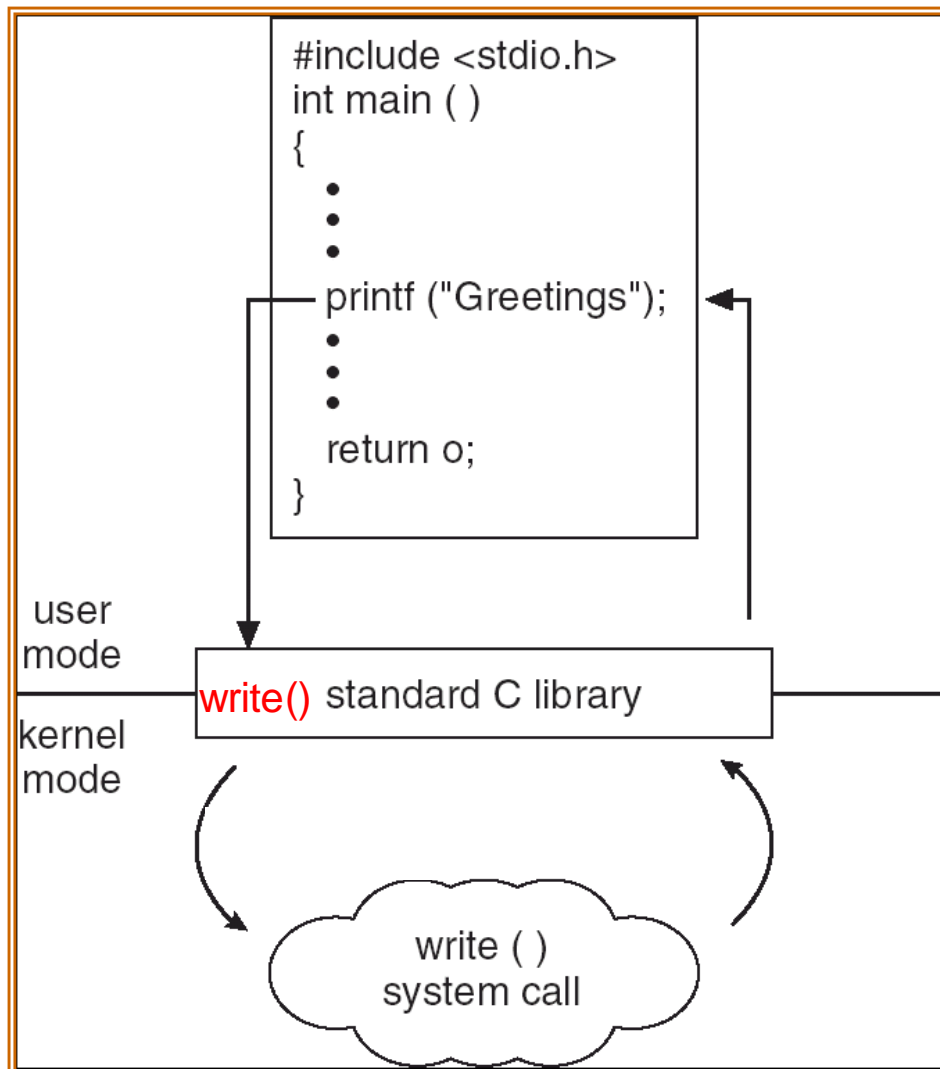


API – System Call – OS Relationship



Standard C Library Example

- C program invoking printf() library call, which calls write() system call



System Call Parameter Passing

- A system call accepts an operation code and a set of parameters
- Three general methods to pass parameters to the OS
 - Pass the parameters in **registers**
 - Parameters stored in a block, or **table**, in memory, and address of block passed as a parameter in a register
 - Parameters placed, or pushed, into the **stack** by the program and popped off the stack by the operating system

Direct System Call in Linux (NASM Syntax)

```
section .data                                ;declare section
msg db  "Hello World! :)",0xa                ;our dear string
len equ $ - msg                              ;length of our dear string

section .text                                ; section declaration

    global _start                            ; exporting entry point
                                           ; to the ELF linker

_start:
; write Hello World string
    mov edx,len ;third arg: message length
    mov ecx,msg ;second arg: pointer to message to write
    mov ebx,1   ;first arg: file handle (stdout)
    mov eax,4   ;system call nr. (sys_write)
    int 0x80    ;call kernel (trigger a trap)
; and exit
    mov ebx,0   ;first syscall args: exit code
    mov eax,1   ;system call no. (sys_exit)
    int 0x80    ;call kernel
```


More on System Calls

- Direct System Call in Windows NT
 - Use the following fragment of assembly code to call the kernel

```
MOV EAX, <service #>  
LEA EDX, <addr of 1st arg>  
INT 2E
```

- Return value is in EAX (if any)
- For modern Intel / AMD CPUs, SYSENTER / SYSCALL are suggested for making system calls, respectively
 - System call entry address is stored in a control register

TYPES OF SYSTEM CALLS

Types of System Calls

- Process control
- File management
- Device management
- Information maintenance
- Communications

Examples of Windows and Unix System Calls

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

Linux System Calls

List by system call number

[00](#) sys_setup [sys_ni_syscall]
[01](#) sys_exit
[02](#) sys_fork
[03](#) sys_read
[04](#) sys_write
[05](#) sys_open
[06](#) sys_close
[07](#) sys_waitpid
[08](#) sys_creat
[09](#) sys_link
[10](#) sys_unlink
[11](#) sys_execve
[12](#) sys_chdir
[13](#) sys_time
[14](#) sys_mknod
[15](#) sys_chmod
[16](#) sys_lchown
[17](#) sys_break [sys_ni_syscall]
[18](#) sys_oldstat [sys_stat]
[19](#) sys_lseek
[20](#) sys_getpid
[21](#) sys_mount

...

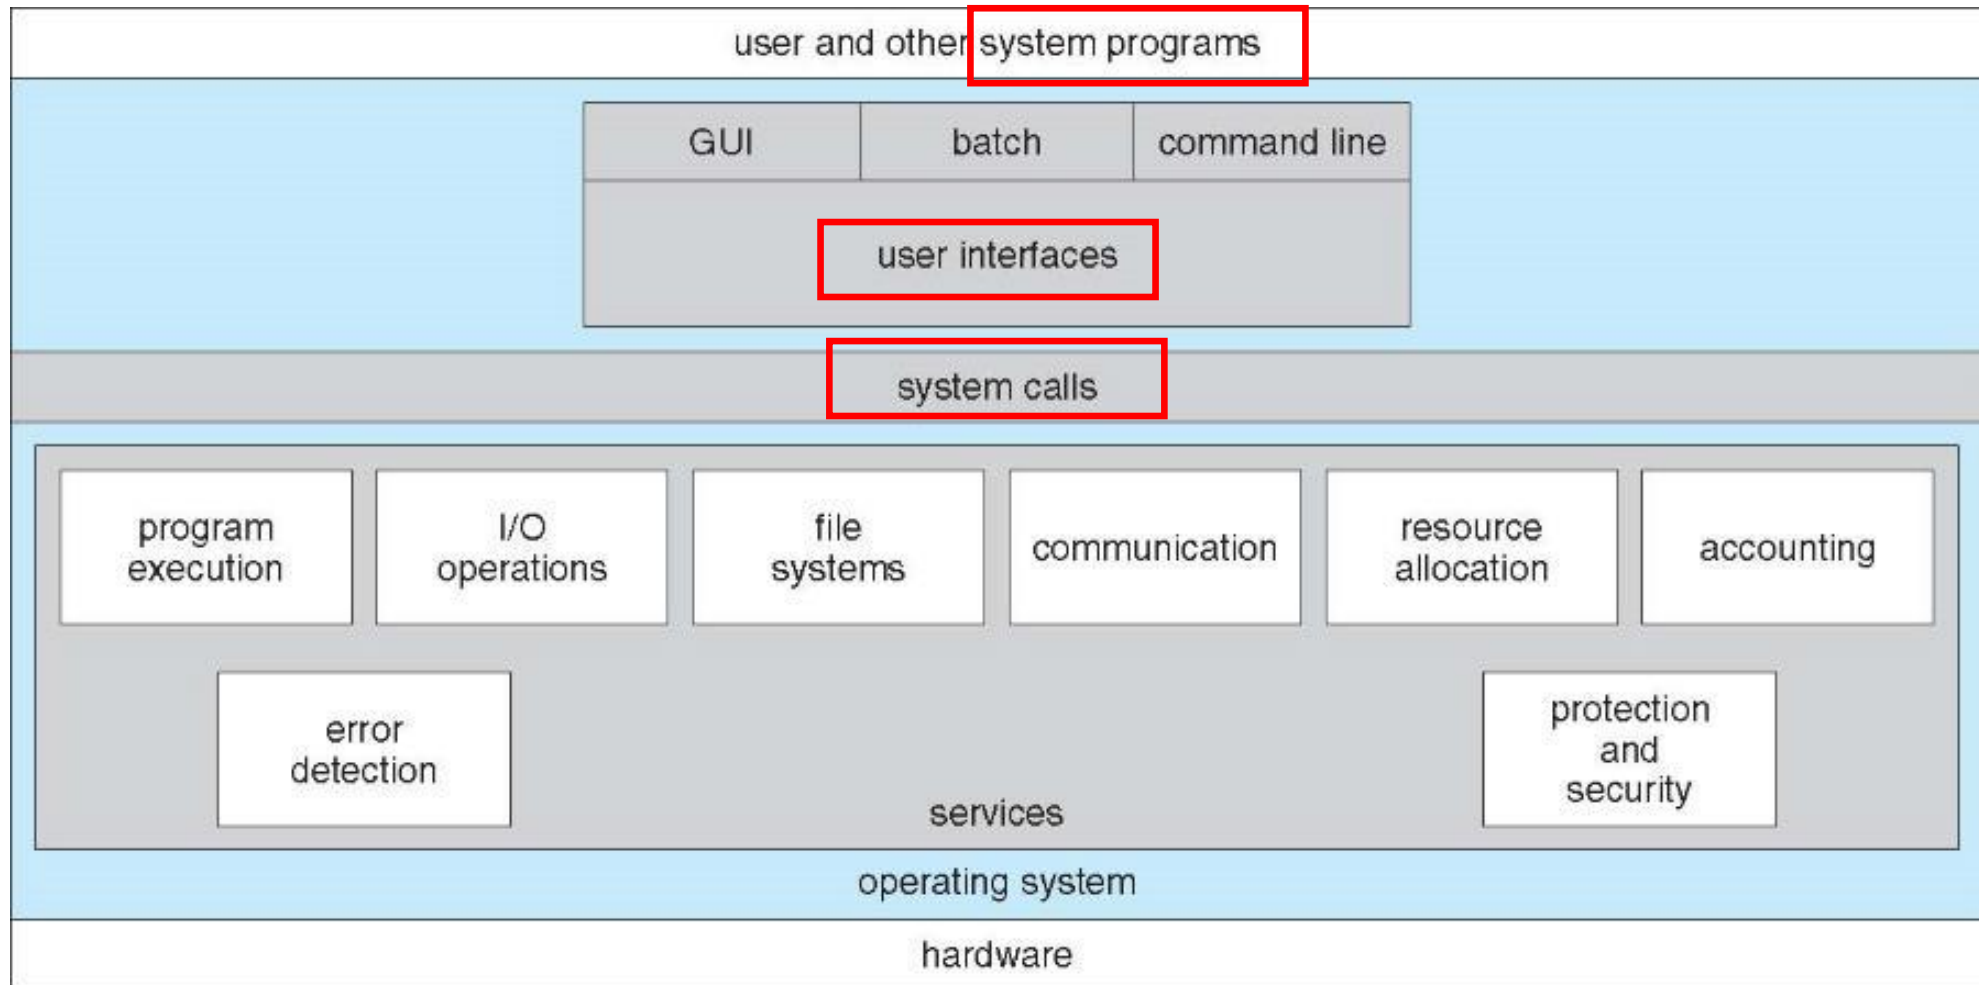
[70](#) sys_setreuid
[71](#) sys_setregid
[72](#) sys_sigsuspend
[73](#) sys_sigpending
[74](#) sys_sethostname
[75](#) sys_setrlimit
[76](#) sys_getrlimit
[77](#) sys_getrusage
[78](#) sys_gettimeofday
[79](#) sys_settimeofday
[80](#) sys_getgroups
[81](#) sys_setgroups
[82](#) sys_select [old_select]
[83](#) sys_symlink
[84](#) sys_oldlstat [sys_lstat]
[85](#) sys_readlink
[86](#) sys_uselib
[87](#) sys_swapon
[88](#) sys_reboot
[89](#) sys_readdir [old_readdir]
[90](#) sys_mmap [old_mmap]
[91](#) sys_munmap

...

[140](#) sys__llseek [sys_lseek]
[141](#) sys_getdents
[142](#) sys__newselect [sys_select]
[143](#) sys_flock
[144](#) sys_msync
[145](#) sys_readv
[146](#) sys_writev
[147](#) sys_getsid
[148](#) sys_fdatasync
[149](#) sys__sysctl [sys_sysctl]
[150](#) sys_mlock
[151](#) sys_munlock
[152](#) sys_mlockall
[153](#) sys_munlockall
[154](#) sys_sched_setparam
[155](#) sys_sched_getparam
[156](#) sys_sched_setscheduler
[157](#) sys_sched_getscheduler
[158](#) sys_sched_yield
[159](#) sys_sched_get_priority_max
[160](#) sys_sched_get_priority_min
[161](#) sys_sched_rr_get_interval

...

Recap: Operating System Services



OPERATING SYSTEM DESIGN AND IMPLEMENTATION

Operating System Design and Implementation

- Design and Implementation of OS not “solvable”, but some approaches have proven successful
- Internal structure of different Operating Systems can vary widely
- Start by defining **goals** and **specifications**
- Affected by choice of hardware, type of system
- User goals and System goals
 - User goals –convenient to use, easy to learn, reliable, safe, and fast
 - System goals –easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient
- Design issues for different types of systems
 - Real-time OS: time predictability, low latency, reliability
 - Mainframe OS: throughput, scalability
 - Desktop OS: responsiveness, user friendly

Operating System Design and Implementation (Cont.)

- Important principle to separate
 - **Mechanism**: How to do it?
 - **Policy**: What will be done?
- Mechanisms determine how to do something, policies decide what will be done next
- The separation of policy from mechanism is a very important principle, it allows maximum flexibility if policy decisions are to be changed later
- Use disk I/O as an example:
 - Mechanism: How to read and write from disk?
 - Policy: Which disk I/O operation should be performed first?

Which one(s) of the following are policies; which are mechanisms?

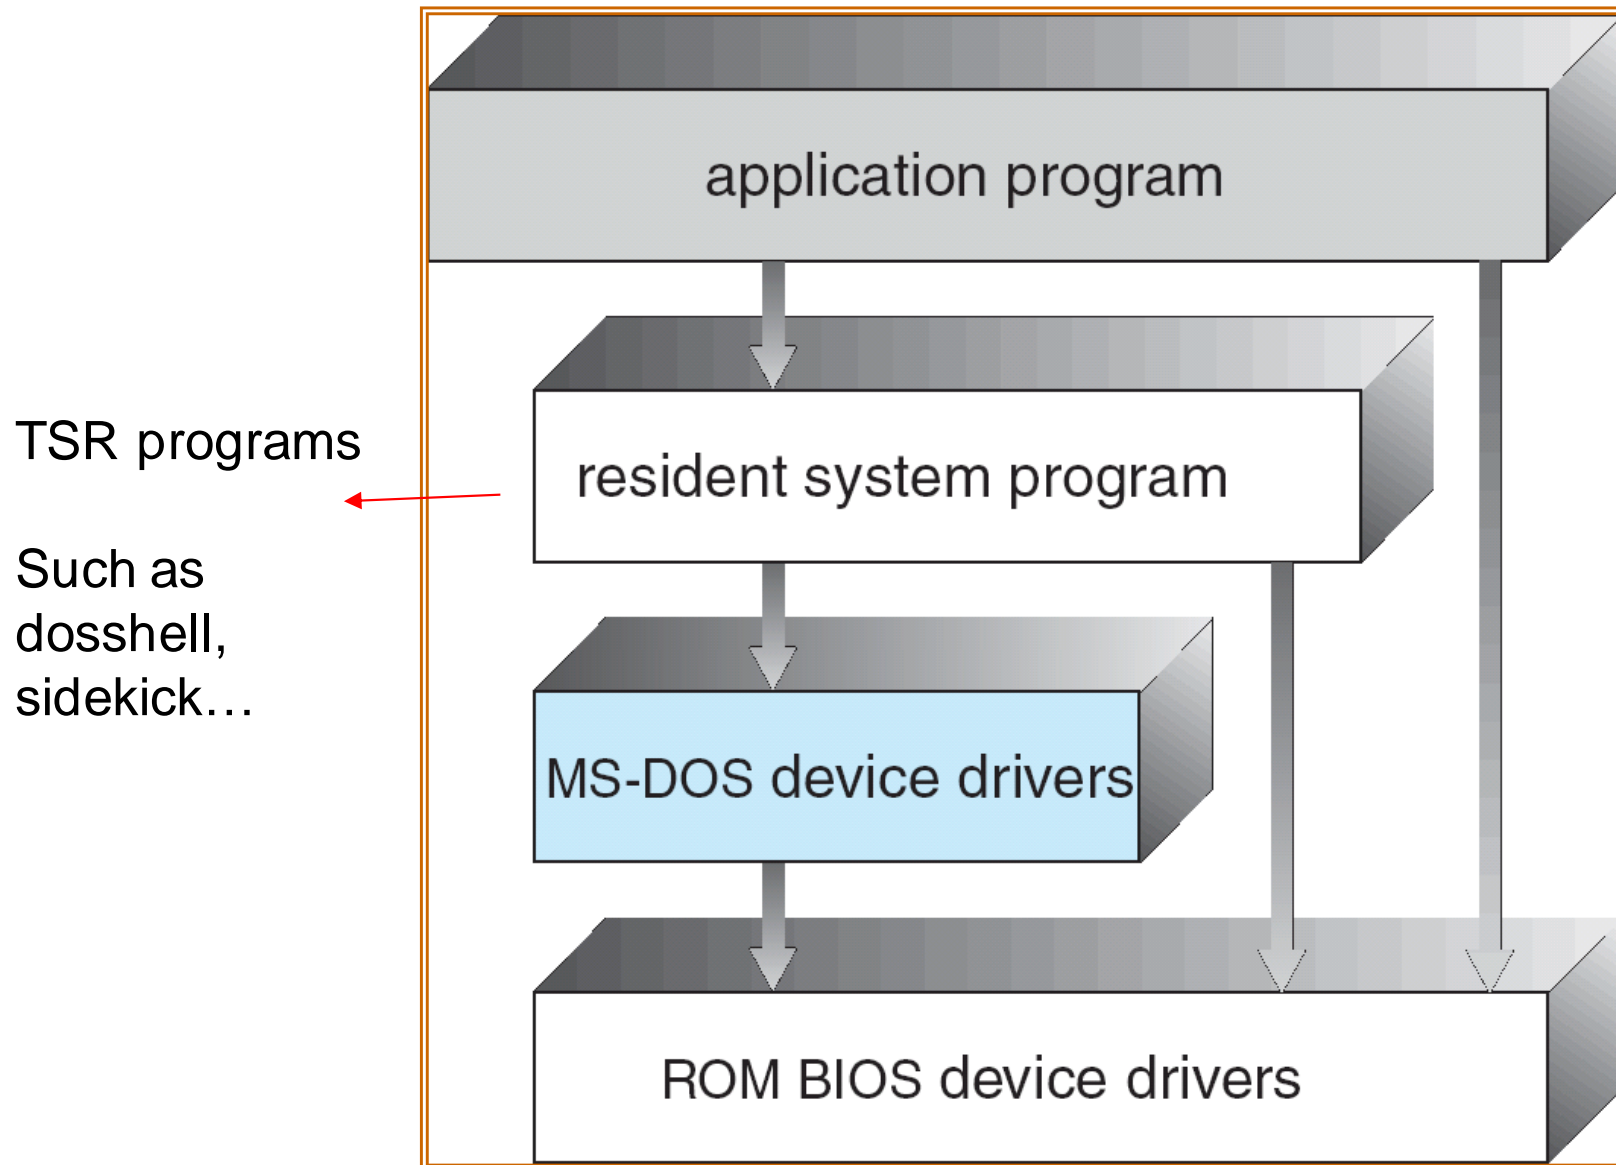
- a) process suspend/resume
- b) allocating the smallest among the memory blocks which are larger than the requested size
- c) marking a disk block as allocated
- d) servicing the disk I/O request which is closest to the disk head

OPERATING-SYSTEM STRUCTURE

Simple Structure

- MS-DOS – written to provide the most functionality in the least space
 - Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated
- No protection. Applications can directly access any memory addresses and control hardware
- MS-DOS is left no choice because the CPU 8086 offers no hardware protection for user-kernel separation

MS-DOS Layer Structure



Booting MS-DOS



UNIX--monolithic

- UNIX – limited by hardware functionality, the original UNIX operating system had “limited” structuring.

The UNIX OS consists of two separable parts:

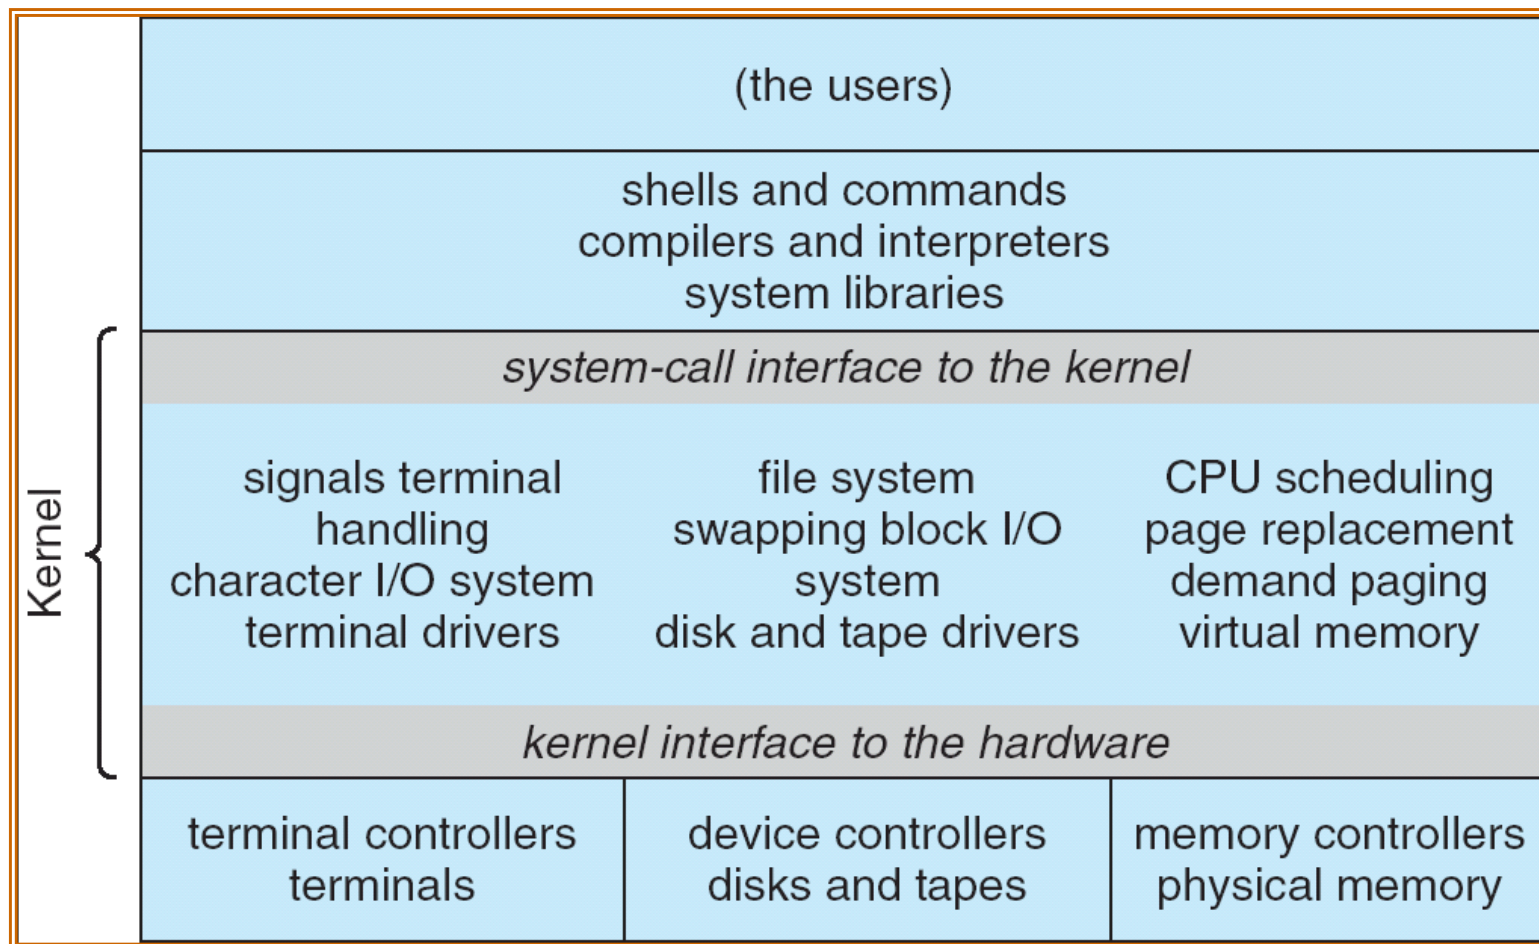
1. Systems programs

- binutils + coreutils: ls, mv, cp, etc

2. The kernel

- Consists of everything below the system-call interface and above the physical hardware
- Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level

UNIX System Structure



UNIX

UN*X is, of course, a huge monolith operating system
Separation of user space from kernel space

Modules

- Most modern operating systems implement kernel modules
 - Uses object-oriented approach
 - Each core component is separate
 - Each talks to the others over known interfaces
 - Each is loadable as needed within the kernel

Common interfaces of a Linux kernel module

Initialize

Clean up

Read (char)

Write (char)

Read (block)

Write (block)

Case Study: Linux Module

```
#include <linux/kernel.h> /* header file for structure pr_info */
#include <linux/init.h>
#include <linux/module.h> /* header file for all modules */
#include <linux/version.h>

MODULE_DESCRIPTION("Hello world !!");
MODULE_AUTHOR("John Doe");
MODULE_LICENSE("GPL");

static int __init hello_init(void)
{
    pr_info("Hello, world\n");
    pr_info("The process is \"%s\" (pid %i)\n", current->comm, current->pid);
    return 0;
}

static void __exit hello_exit(void)
{
    printk(KERN_INFO "Goodbye\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

Source:

<http://blog.wu-boy.com/2010/06/linux-kernel-driver-%E6%92%B0%E5%AF%AB%E7%B0%A1%E5%96%AE-hello-world-module-part-1/>

Case Study: Linux Module

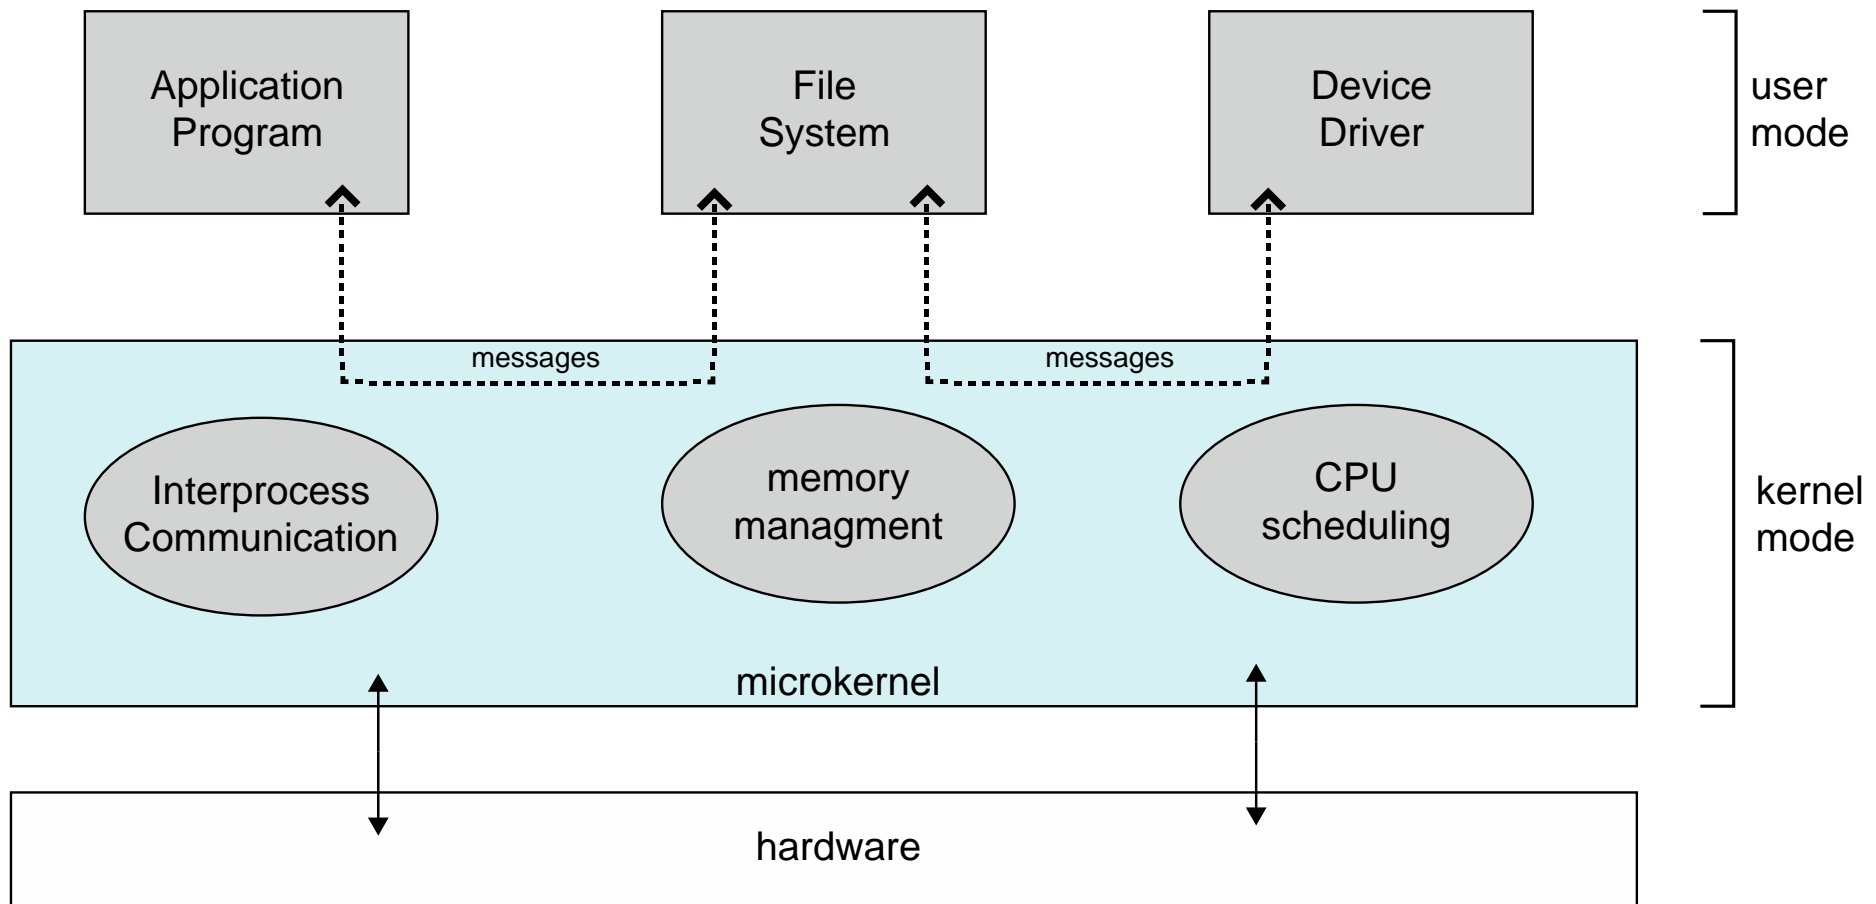
- Insert & init the module
 - insmod ./hello.ko
- Remove & clean up the module
 - rmmod ./hello.ko

```
Jun 21 11:50:00 cvs5 kernel: [8381603.818051] The process is "insmod" (pid 30560)
Jun 21 11:50:08 cvs5 kernel: [8381612.335386] Goodbye
Jun 21 12:07:13 cvs5 rsyslogd: [origin software="rsyslogd" swVersion="4.2.0" x-pid='
.com"] rsyslogd was HUPed, type 'lightweight'.
Jun 21 12:07:13 cvs5 rsyslogd: [origin software="rsyslogd" swVersion="4.2.0" x-pid='
.com"] rsyslogd was HUPed, type 'lightweight'.
Jun 21 14:55:23 cvs5 kernel: [8392723.612597] Hello, world
Jun 21 14:55:23 cvs5 kernel: [8392723.612601] The process is "insmod" (pid 10072)
Jun 21 14:55:37 cvs5 kernel: [8392737.604360] Goodbye
Jun 21 15:05:18 cvs5 kernel: [8393318.795982] Hello, world
Jun 21 15:05:18 cvs5 kernel: [8393318.795985] The process is "insmod" (pid 13127)
Jun 21 15:05:25 cvs5 kernel: [8393325.537903] Goodbye
```

Microkernel System Structure

- Moves as much from the kernel into “user” space
- Communication takes place between user modules (processes, specifically) using **message passing**
- Benefits:
 - Easier to **extend** a microkernel (by adding user-mode modules)
 - Easier to **port** the operating system to new architectures
 - More **reliable and secure** (less code is running in kernel mode)
- Detriments:
 - Performance overhead of user space to kernel space communication and user-kernel mode switches
 - What happened to Windows NT 3.5?

Microkernel System Structure



The Famous Tanenbaum–Torvalds Debate

- *"I'm doing a (free) operating system (just a hobby, won't be big and professional like gnu) for 386(486) AT clones."*

-- Linus Torvalds

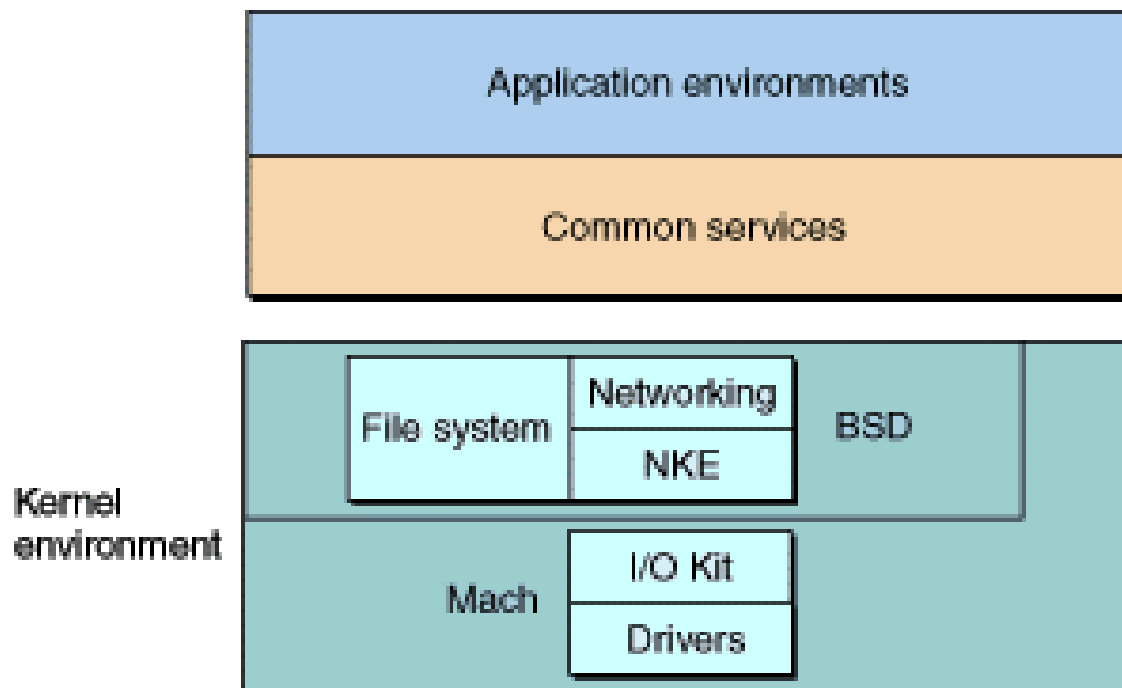
- *"Linux is a monolithic style system. This is a giant step back into the 1970s"*

-- Andrew Tanenbaum

- [Wiki entry](#)

Mac OS X Structure

- Mach (μ -kernel): memory management, RPC, IPC, message passing, thread scheduling
- BSD: networking, file systems, POSIX APIs



Google Fuchsia

- A new operating system developed by Google, based on the Zircon microkernel
- Reportedly designed for IoT devices



Flutter (UI framework) + Dart (language) → UI
Fuchsia → system services and IPC
Zircon → microkernel

Quiz

What are advantages of the micro-kernel approach?

1. Extensibility
2. Robustness
3. Efficiency
4. Security

Summary: Microkernel

- Provide “a minimal set of kernel primitives”
- Pros
 - robust, extensible, secure, portable
- Cons
 - Frequent mode switches
 - High message-passing overhead

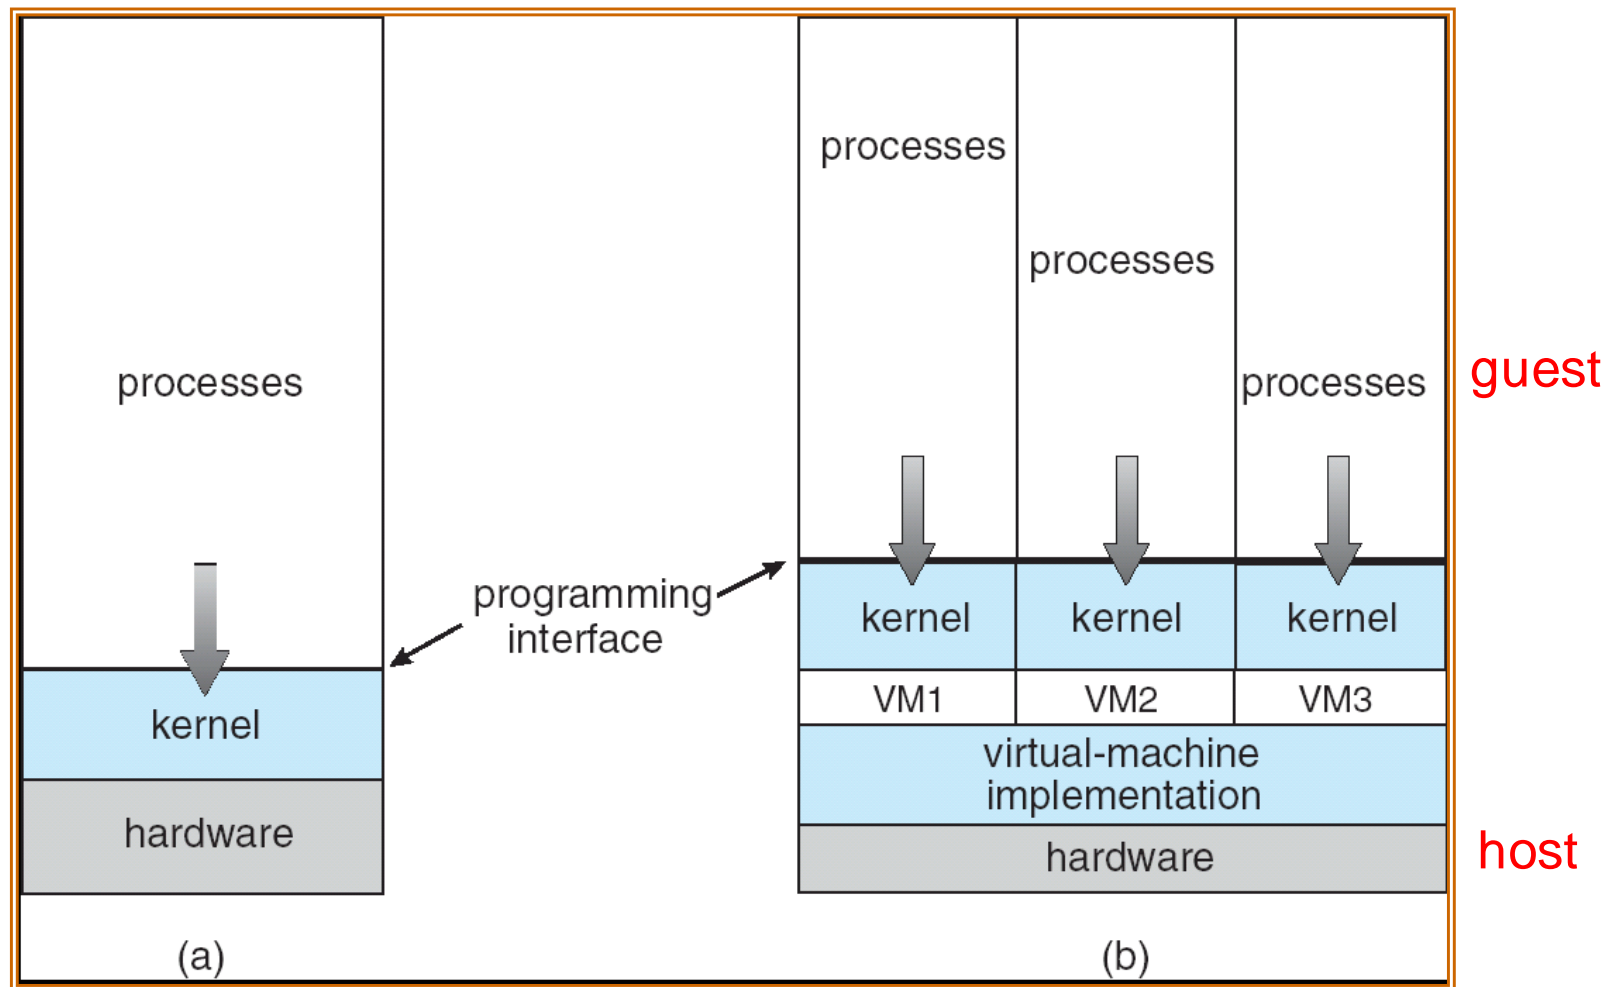
Virtual Machines

- A virtual machine takes the layered approach to its logical conclusion. It treats hardware and the operating system kernel as though they were all hardware
- A virtual machine provides an interface identical to the underlying bare hardware
- The operating system creates the illusion of multiple processes, each executing on its own processor with its own (virtual) memory
- Virtual machine is not a new concept. It has been developed in 197x. VM again become popular because
 - It becomes hard to maintain outdated servers
 - Cloud computing (service virtualization)

Virtual Machines (Cont.)

- The resources of the physical computer are shared to create the virtual machines
 - CPU scheduling can create the appearance that users have their own processor
 - Spooling and a file system can provide virtual card readers and virtual line printers
 - A normal user time-sharing terminal serves as the virtual machine operator's console

Virtual Machines (Cont.)

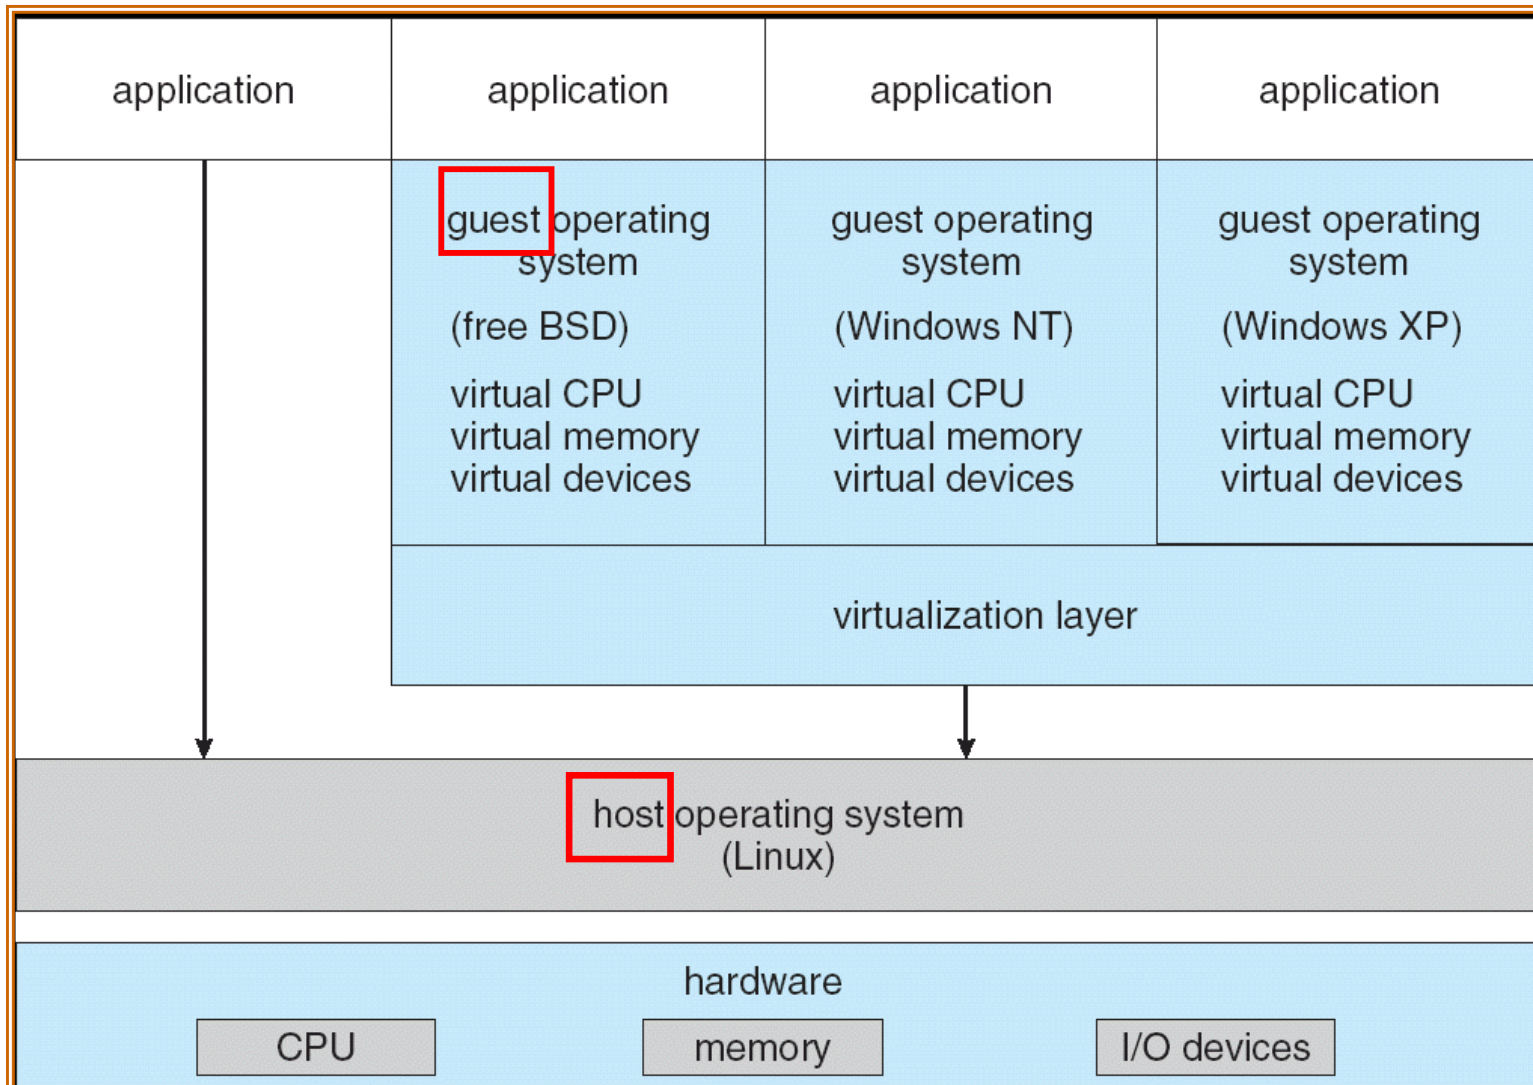


(a) bare-metal (b) virtual machines

Virtual Machines (Cont.)


- The virtual-machine concept provides **complete protection of system resources since each virtual machine is isolated from all other virtual machines**. This isolation, however, permits no direct sharing of resources.
- A virtual-machine system is a **perfect vehicle for operating-systems research and development**. System development is done on the virtual machine, instead of on a physical machine and so does not disrupt normal system operation.
- The virtual machine concept is **difficult** to implement due to the effort required to provide an exact duplicate to the underlying machine

VMware Architecture



Native execution

The Java Virtual Machine



Java programs: the source code

Byte code: the compiled binary for JVM

JVM or java runtime: a hardware-independent virtual machine

*Non-native execution


Java Bytecode

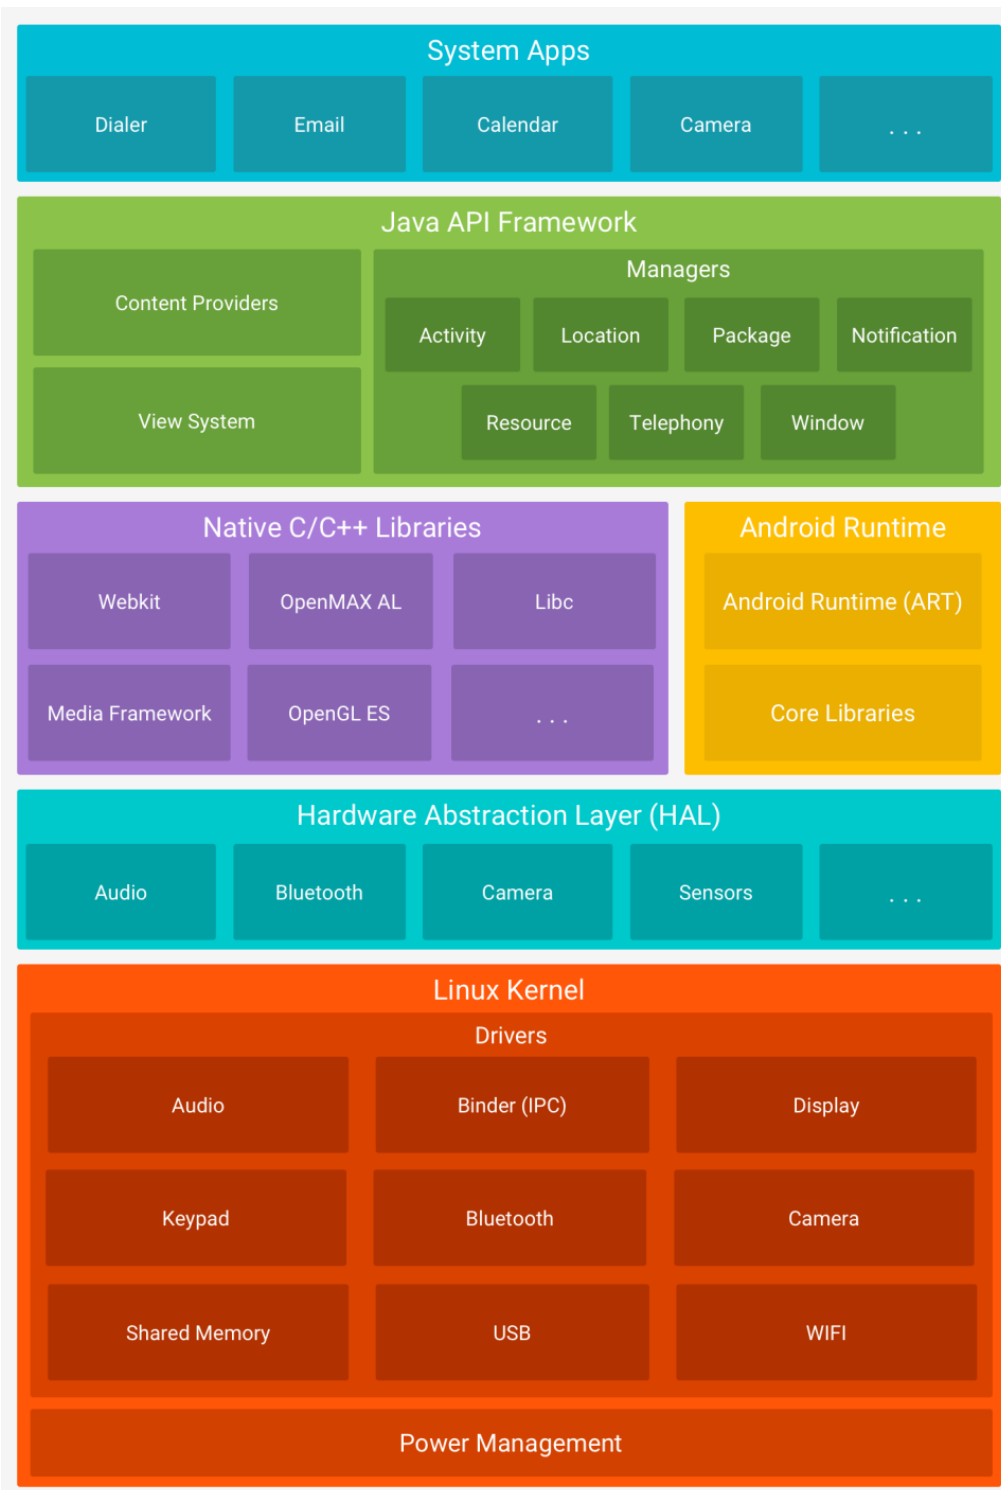
iload_1

iload_2

iadd

istore_3





Android Structure

- **Java API Framework:** The entire feature-set of the Android OS is available to you through APIs written in the Java language.
- **ART** is written to run multiple virtual machines on low-memory devices by executing DEX files
- The [hardware abstraction layer \(HAL\)](#) provides standard interfaces that expose device hardware capabilities to the higher-level [Java API framework](#).

Quiz

The virtual machine approach is suitable to which one(s) of the following scenarios?

1. OS development
2. Cloud computing
3. Performance-critical gaming
4. Writing an application for heterogeneous hardware platforms

Summary: Virtual Machines

- Virtualizes hardware
- Pros
 - Guest operating systems run without modifications
 - Perfect resource partition and fault isolation
 - Resource sharing among VMs
- Cons
 - Inefficient mapping between emulated hardware and the underlying hardware
 - Hard to implement hardware virtualization

Review

- Simple structure
 - Pros: simple, cons: poorly structured
- Monolithic
 - Pros: efficient, cons: hard to scale
- Microkernel
 - Pros: robust and scalable, cons: inefficient
- Virtual machine
 - Pros: perfect resource isolation, cons: inefficient hardware emulation

End of Chapter 2