

# Chapter 5: CPU Scheduling

Prof. Li-Pin Chang  
CS@NYCU

# Chapter 5: Process Scheduling

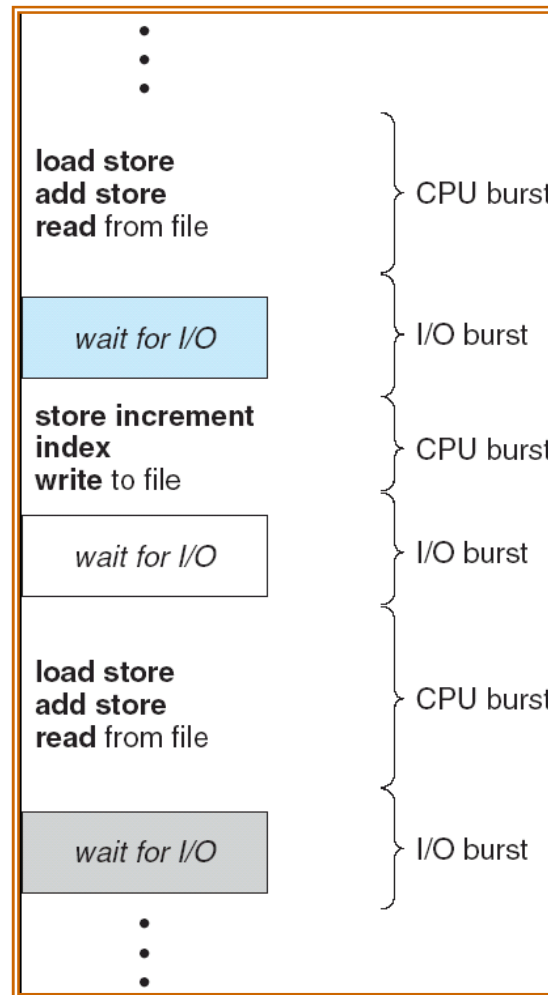
- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Multiple-Processor Scheduling
- Thread Scheduling
- Operating Systems Examples
- Algorithm Evaluation

# BASIC CONCEPTS


# Basic Concepts

- Maximum CPU utilization obtained with multiprogramming
- CPU–I/O Burst Cycle – Process execution consists of a cycle of CPU execution and I/O wait
- CPU burst distribution

# Alternating Sequence of CPU And I/O Bursts



## Histogram of CPU-burst Times



In timesharing systems, the majority of the CPU bursts are short

# CPU Scheduler

- Selects from among the processes in memory that are **ready** to execute, and allocates the CPU to one of them
- CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state
  2. Switches from running to ready state (←?)
  3. Switches from waiting to ready (←?)
  4. Terminates
- Scheduling under **only** 1 and 4 is non-preemptive or cooperative
- All other scheduling is **preemptive**

# Cooperative scheduling

- Easy to implement and requires no extra hardware (e.g., the timer)
- A process may voluntarily gives up the CPU
  - Call Sleep(0) or yield()
  - A blocking call also causes a context switch
- An ill-behaved process can take over the entire system
- Example
  - Windows 3.1
  - Old versions of Mac OS
  - Sensor-node OS (e.g., TinyOS)



# Scheduler Preemptivity

- Preemptive scheduling
  - Higher responsiveness
  - Higher cxtsw overheads
  - Must deal with race conditions
- Cooperative scheduling
  - Easy to implement
  - Poor responsiveness
  - Ill-behaved processes will take over the system

# Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - switching context
  - switching to user mode
  - jumping to the proper location in the user program to restart that program
- Dispatch latency – time it takes for the dispatcher to stop one process and start another running

the context-switch overhead

# SCHEDULING CRITERIA

# Scheduling Criteria

- (+)CPU utilization – keep the CPU as busy as possible
- (+) Throughput – # of processes that complete their execution per time unit
- (-)Turnaround time – amount of time to execute a particular process
- (-) Waiting time – amount of time a process has been waiting in the ready queue
- (-) Response time – amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)

Good average performance vs. Predictable worst-case performance

# Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time
  
- There are conflicts among the objectives
  - E.g., throughput vs. waiting time

# SCHEDULING ALGORITHMS

# First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:




- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$

## FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than the previous case



- *Convoy effect* : short process behind long process
  - FCFS is non-preemptible -> poor average waiting time
  - Harmful to I/O-bound processes and result in poor I/O utilization!! (why?)
- Question: Which one(s) of the following could be the performance issue of FCFS?
  - Lengthy response
  - Poor I/O utilization
  - Low CPU utilization


# Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
  - Non-preemptive – once CPU given to the process it cannot be preempted until completes its CPU burst
  - Preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- **SJF is optimal** – gives **minimum average waiting time** for a given set of processes (must all be ready at time 0)
  - How to prove it?

SJF: All ready at time 0

Process	Arrival Time	Burst Time
P1	0.0	6
P2	0.0	8
P3	0.0	7
P4	0.0	3

SJF (non-preemptive)




- Average waiting time =  $(0 + 3 + 9 + 16)/4 = 7$
- Compared to FCFS?  $(0+6+14+21)/4=10.25$

## Non-Preemptive SJF with Arbitrary Arrival Times

Process	Arrival Time	Burst Time
P1	0.0	7
P2	2.0	4
P3	4.0	1
P4	5.0	4

- SJF (non-preemptive)




- Average waiting time =  $(0 + 6 + 3 + 7)/4 = 4$

## Preemptive SJF with Arbitrary Arrival Times

Process	Arrival Time	Burst Time
P1	0.0	7
P2	2.0	4
P3	4.0	1
P4	5.0	4

- SJF (preemptive)



- Average waiting time =  $(9 + 1 + 0 + 2)/4 = 3$

- Question: Which one(s) is the characteristic of SJF?
  - Short average waiting time
  - Small waiting time variation

# Fairness vs. Efficiency

- With SJF, long jobs may be indefinitely delayed if short jobs keep arriving
  - Starvation
- SJF is thus not a fair scheduling algorithm
  - Processes have long CPU bursts may starve
  - By contrast, with FCFS, a process's waiting time is always bounded
- The long-standing dilemma: efficiency vs. fairness

# Determining Length of Next CPU Burst


- It is nearly impossible to know the actual job execution time in advance
- Using the exponential moving average method to predict the length of the next CPU burst based on the lengths of previous CPU bursts

1.  $t_n$  = actual length of  $n^{th}$  CPU burst
2.  $\tau_{n+1}$  = predicted value for the next CPU burst
3.  $\alpha, 0 \leq \alpha \leq 1$
4. Define :

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$



# Prediction of the Length of the Next CPU Burst



Alpha=0.5

# Examples of Exponential Moving Average

- $\alpha = 0$ 
  - $\tau_{n+1} = \tau_n$
  - Recent history does not count
- $\alpha = 1$ 
  - $\tau_{n+1} = \alpha t_n$
  - Only the actual last CPU burst counts
- If we expand the formula, we get:

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)^1 \alpha t_{n-1} + \dots + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^{n+1} \tau_0$$


- Since both  $\alpha$  and  $(1 - \alpha)$  are less than or equal to 1, **each successive term has less weight than its predecessor**

# Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - Preemptive or non-preemptive
- Problem: **Starvation** – low priority processes may never execute
- Solution: **Aging** – as time progresses increase the priority of the process

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

Using priority scheduling, we would schedule these processes according to the following Gantt chart:




## Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum**), usually 10-100 milliseconds. After this time has elapsed, the process is **preempted** and **added to the end of the ready queue**.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. **No process waits more than  $(n-1)q$  time units.**
- Performance
  - $q$  large  $\Rightarrow$  FIFO
  - $q$  small  $\Rightarrow$   $q$  must be large with respect to context switch, otherwise overhead is too high

## Example of RR with Time Quantum = 20


<u>Process</u>	<u>Burst Time</u>
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24

- The Gantt chart is:




- Processes with RR have longer turnaround time but better response

# Time Quantum and Context Switch Time



# Turnaround Time Varies With The Time Quantum



Large quantum: better turnaround  
Small quantum: better interactivity



## RR vs. FCFS


- Processes have better response (interactivity) with RR
  - Timesharing
- Processes (often) have shorter turnaround times with FCFS
  - Consider the case that the time quantum is extremely small. With RR, two processes of the same burst length will complete at the same time

# Multilevel Queue

- Ready queue is partitioned into separate queues:  
foreground (interactive)  
background (batch)
- Each queue has its own scheduling algorithm
  - foreground – RR
  - background – FCFS
- Scheduling must be done between the queues
  - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
  - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR and 20% to background in FCFS

(Rumor has it that, when they shut down the IBM 7094 at MIT in 1973, they found a low-priority process that had been submitted in 1967 and had not yet been run.)

# Multilevel Queue Scheduling




# Multilevel Feedback Queue

- A process can move between the various queues; **aging** can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - scheduling algorithms for each queue
  - method used to determine when to **upgrade** a process
  - method used to determine when to **demote** a process
  - method used to determine which queue a process will enter when that process needs service

# Example of Multilevel Feedback Queue

- Three queues:
  - Q0 – RR with time quantum 8 milliseconds
  - Q1 – RR time quantum 16 milliseconds
  - Q2 – FCFS
- Scheduling
  - A new job enters queue Q0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not **finish (?)** in 8 milliseconds, job is moved to queue Q1.
  - At Q1 job is again served FCFS and receives 16 additional milliseconds. If it still does not **complete (?)**, it is preempted and moved to queue Q2.

# Multilevel Feedback Queues



We should add a promotion policy

# Multilevel Feedback Queues

- Real designs also employ a promotion policy
  - See the Solaris priority table
- Why do we promote I/O-bound processes?
  - To improve I/O utilization (recall the convoy effect)
  - To favor interactive processes
- Why do we demote CPU-bound processes?
  - Assign a large time quantum
  - To improve throughput and turnaround


# MULTIPLE PROCESSOR SCHEDULING




# Multiple-Processor Scheduling

- Process scheduling
- Symmetric multiprocessing (i.e., **partitioned scheduling**) – each processor is self-scheduling all ~~processes in common ready queue, or each~~ and has its own private queue of ready processes (e.g., Linux)
- Asymmetric multiprocessing (i.e., **global scheduling**) – only one processor accesses the system data structures, ~~alleviating the need for data sharing~~ and acts as the dispatcher

Partitioned Scheduling



Global Scheduling




# Multiple-Processor Scheduling

- Architecture
- Homogeneous architecture -- all CPUs share global memory (**SMP**)
- Heterogeneous architecture -- CPUs have local memory and memory reference is either local or remote (**NUMA**)


# Symmetric multiprocessing (SMP)

- Multiple CPUs share the same main memory and contend for memory cycles via the common bus
- Good performance with a few CPUs; high costs of bus congestion and cache coherence with many CPUs



# Non-Uniform Memory Access (NUMA)

- Memory access time depends on the location of memory
- More scalable than SMP
- Common topology: ring, mesh, hypercube\*, ad-hoc



\* <https://lwn.net/Articles/254445/>

# Multiple-Processor Scheduling

- **Processor affinity**

- The cost of process migration: cache re-population, pipeline re-start, and housekeeping data transfer
- Sometimes it's better to stick processes with processors
- Soft affinity, hard affinity


- **Load balancing**

- Push migration – push a process away from a heavily loaded processor to an idle processor
  - Pull migration – pull a waiting process from a heavily loaded processor into an idle processor
  - Linux – runs push migration every 200ms and runs pull migration whenever a processor queue is empty
- The two goals **conflict** with each other!

# Multicore and Multithreading Processors

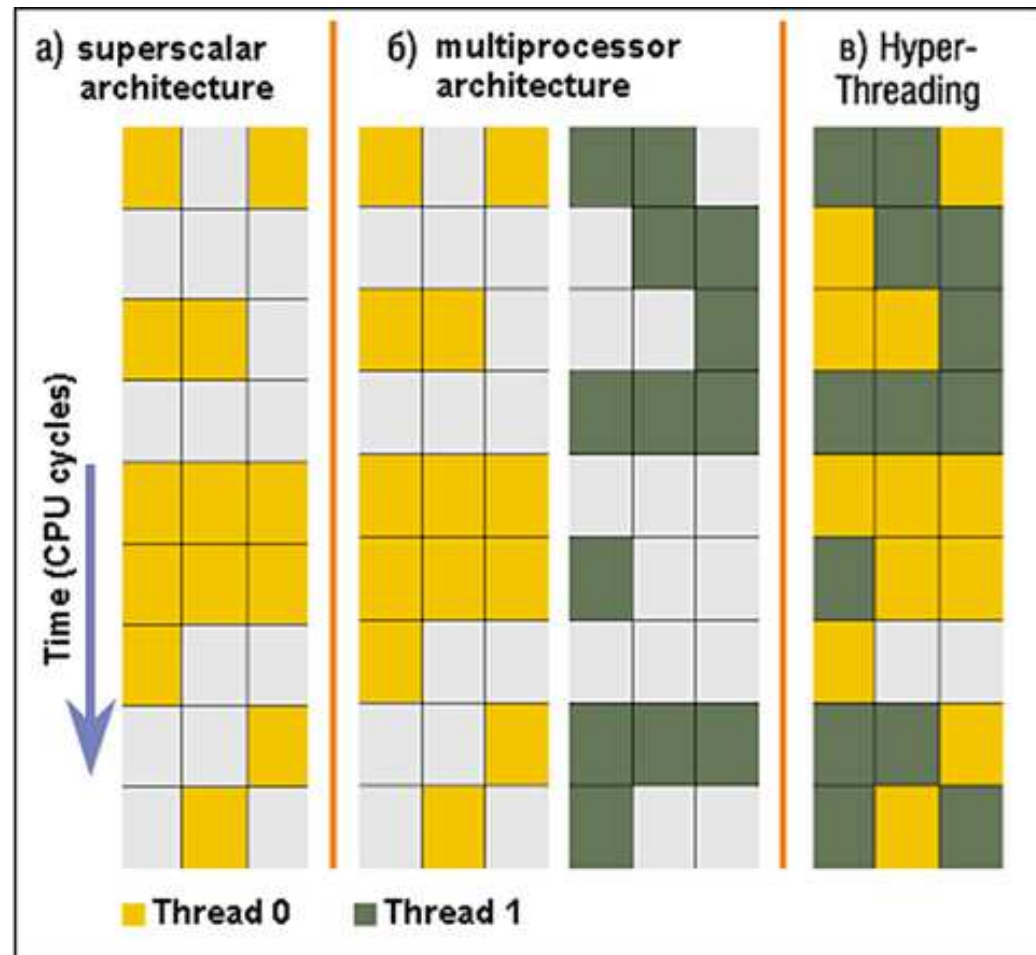
- Multiple **physical** processor cores on the same physical chip
  - Faster communication
  - More power efficient
  - Better cache sharing
- Multiple (hardware) **threads** per core also growing
  - Takes advantage of memory stall to make progress on another thread while memory retrieve happens
  - A thread here is hardware-oriented, **different** from “threads” in Chapter 4
  - A thread corresponds to a **logical processor**, emulated by an **independent set of registers**

# Improving CPU and Memory Cycle Utilizations




Higher utilizations of both CPU and memory

# Better Utilization of Processing Units in Superscalar (Multi-Issue) Processors





# Hierarchical Scheduling Domains



Assume: 2 PP, each has 2 LP.

$[1\ 1][1\ 1]$  is better than  $[2\ 0][2\ 0]$

$[1\ 0][1\ 0]$  is better than  $[1\ 1][0\ 0]$

Linux: **hierarchical** scheduling domains for load balancing

Firstly, evenly distribute process among physical CPUs

In each physical CPU, evenly distribute processes among logical CPUs


# Real-Time Scheduling

# Real-Time CPU Scheduling

- IEEE definition of real time systems:
  - “A real-time system is a system whose correctness includes its response time as well as its functional correctness.”
- **Soft real-time systems** – no guarantee as to when critical real-time process will be scheduled
- **Hard real-time systems** – task must be serviced by its deadline


# Priority-based Scheduling

- For real-time scheduling, scheduler must support preemptive, priority-based scheduling
  - But only guarantees soft real-time
- For hard real-time must also provide ability to meet deadlines
- Processes have new characteristics: **periodic** ones require CPU at constant intervals
  - Has processing time  $t$ , deadline  $d$ , period  $p$
  - $0 \leq t \leq d \leq p$
  - **Rate**



# Rate Monotonic Scheduling


- A priority is assigned based on the inverse of its period
- Shorter periods = higher priority;
- Longer periods = lower priority
- P1 is assigned a higher priority than P2.



Computation:  $c_1=20$ ,  $c_2=35$

Periods:  $P_1=50$ ,  $P_2=100$

# Missed Deadlines with Rate Monotonic Scheduling




Computation:  $c_1=25$ ,  $c_2=35$

Periods:  $P_1=50$ ,  $P_2=80$

# Earliest Deadline First Scheduling (EDF)

- Priorities are assigned according to deadlines:

the earlier the deadline, the higher the priority; the later the deadline, the lower the priority



Computation:  $c_1=25$ ,  $c_2=35$


Periods:  $P_1=50$ ,  $P_2=80$

# THREAD SCHEDULING



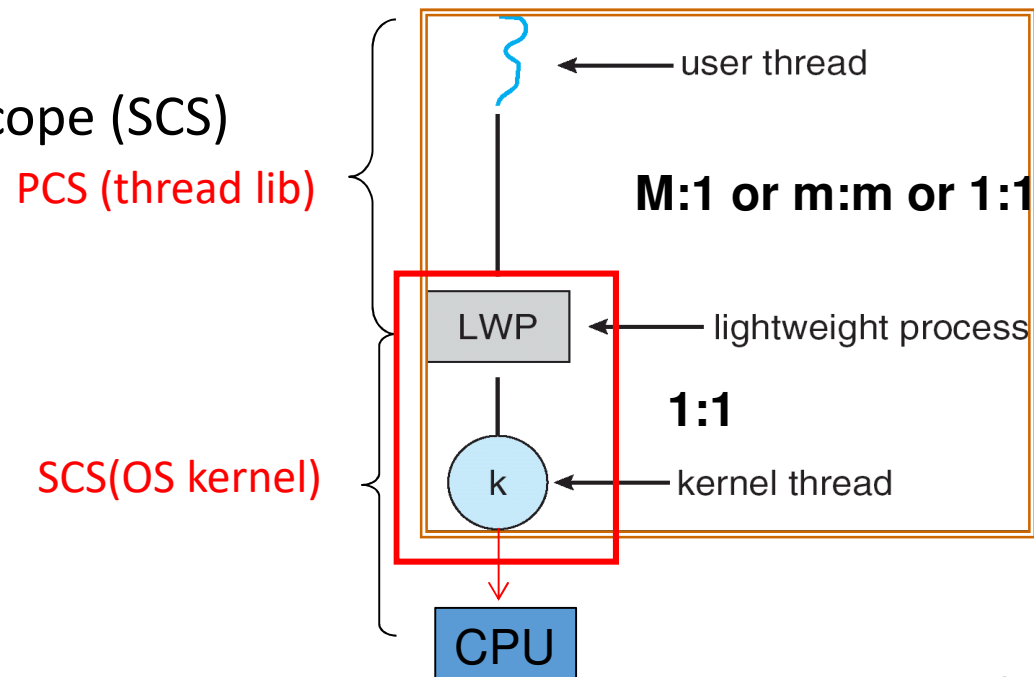
# Light-Weight Process

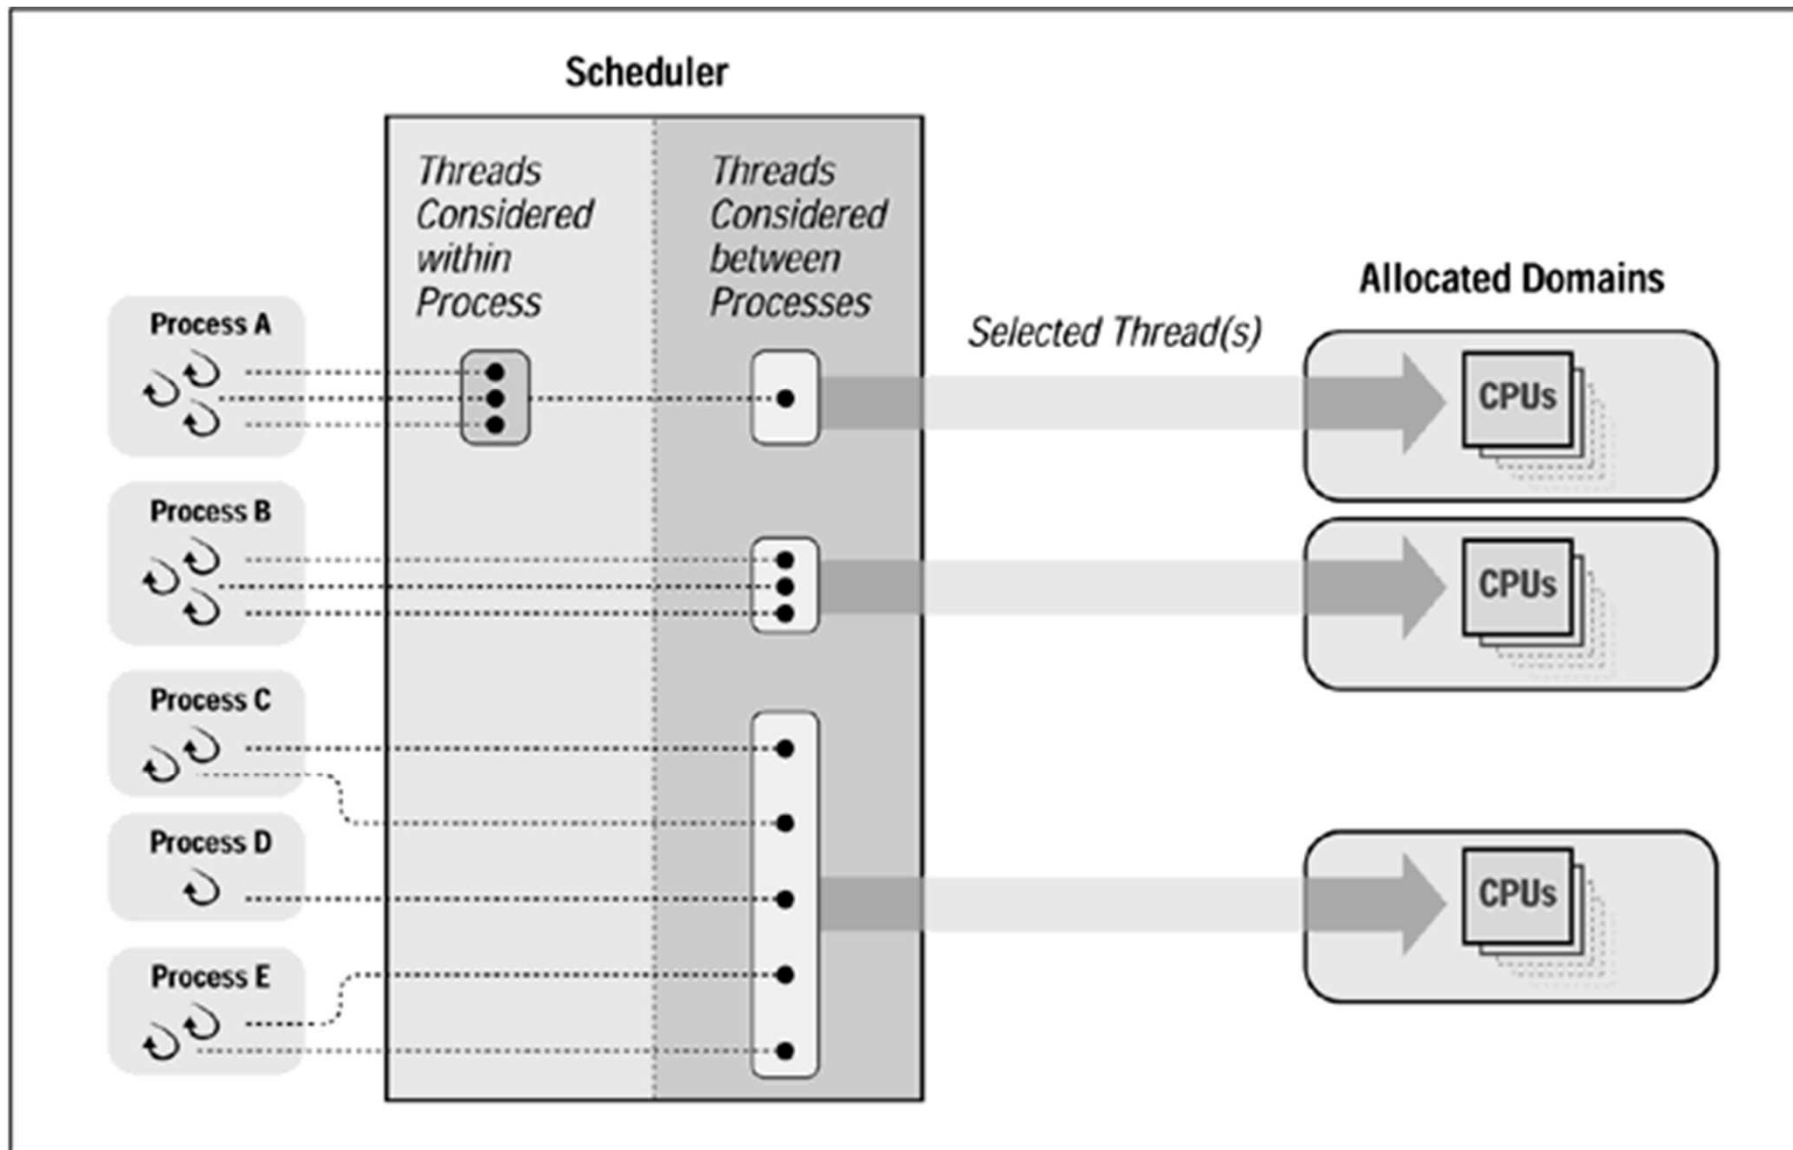
- LWP is an **optional** abstraction of kernel scheduling units
- An LWP is like a virtual processor on which user threads are scheduled
- Basically the mapping of LWPs to kernel threads is 1-1
- The mapping of user threads to LWP is 1-1, M-1, or M-M



# Thread Scheduling

- Local Scheduling – How the threads library decides which thread to put onto an available LWP
  - Process contention scope (PCS)
- Global Scheduling – How the kernel decides which kernel thread to run next
  - System contention scope (SCS)





```

#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
int main(int argc, char *argv[])
{
    int i;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_t init_attr;
    /* set the scheduling algorithm to PROCESS or SYSTEM */
    pthread_attr_t attr;
    /* set the scheduling policy - FIFO, RR, or OTHER */
    pthread_attr_t attr;
    /* create the threads */
    for (i = 0; i < NUM_THREADS; i++)
        pthread_create(&tid[i], &attr, runner, NULL);
    /* now join on each thread */
    for (i = 0; i < NUM_THREADS; i++)
        pthread_join(tid[i], NULL);
}
/* Each thread will begin control in this function */
void *runner(void *param)
{
    printf("I am a thread\n");
    pthread_exit(0);
}

```

# Pthread Scheduling API

Linux Pthread is 1-1, and Linux supports only PTHREAD\_SCOPE\_SYSTEM. Changing the scheduling policy for system contention scope effective changes the thread scheduling policy

PTHREAD\_SCOPE\_SYSTEM  
PTHREAD\_SCOPE\_PROCESS

SCHED\_FIFO  
SCHED\_RR  
SCHED\_DEADLINE  
SCHED\_SPORADIC  
SCHED\_OTHER

# OPERATING-SYSTEM EXAMPLES


# Operating System Examples

- Solaris scheduling
- Windows XP scheduling
- Linux scheduling

# Solaris

- Priority-based scheduling
  - RR on threads of the same priority
- Six classes available
  - Time sharing (default) (TS)
  - Interactive (IA)
  - Real time (RT)
  - System (SYS)
  - Fair Share (FSS)
  - Fixed priority (FP)
- Given thread can be in one class at a time
- Each class has its own scheduling algorithm
- Time sharing is multi-level feedback queue
  - Loadable table configurable by sysadmin

# Solaris Scheduling





# Solaris Dispatch Table

priority	time quantum	time quantum expired	return from sleep
0	200	0	50
5	200	0	50
10	160	0	51
15	160	5	51
20	120	10	52
25	120	15	52
30	80	20	53
35	80	25	54
40	40	30	55
45	40	35	56
50	40	40	58
55	40	45	58
59	20	49	59

- Priority + → time quantum –
- Processes run out their time quantum are demoted
- Processes return from I/O operations are promoted

High

# Windows Scheduling

- Windows uses priority-based preemptive scheduling
- Highest-priority thread runs next
- **Dispatcher** is scheduler
- Thread runs until (1) blocks, (2) uses time slice, (3) preempted by higher-priority thread
- Real-time threads can preempt non-real-time
- 32-level priority scheme
- **Variable class** is 1-15, **real-time class** is 16-31
- Priority 0 is memory-management thread
- Queue for each priority
- If no run-able thread, runs **idle thread**

# Windows Priorities

## Six priority classes

Fixed  
priority

Variable priority

	real-time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

Relative priority with respect to a priority class


Priorities in each individual variable-priority classes are adjusted as they are in the feedback scheduling

# Linux Scheduling in Version 2.6.23 +

- Completely Fair Scheduler (CFS)
- Virtual runtime (vruntime)
  - Increases as a process executes on the CPU
  - The process of the smallest vruntime is selected for running
  - CFS favor IO-bound processes over CPU-bound processes
    - IO-bound processes increase their vruntime slower than CPU-bound processes
- Nice value
  - -20~+19 (high priority ~ low priority)
  - The increasing rate of vruntime
    - E.g., with -20, after a process runs 200ms, its vruntime increases less than 200ms
    - with +19, after a process runs 200ms, its vruntime increases larger than 200ms
  - Processes with small nice values increase their vruntime slower, and thus, receive larger portions of CPU time
- CFS determines the increasing rate of vruntime for a process based on its **nice value** and its **recent execution history** (IO- or CPU-bound)

# CFS Performance


The Linux CFS scheduler provides an efficient algorithm for selecting which task to run next. Each runnable task is placed in a red-black tree—a balanced binary search tree whose key is based on the value of `vruntime`. This tree is shown below:



When a task becomes runnable, it is added to the tree. If a task on the tree is not runnable (for example, if it is blocked while waiting for I/O), it is removed. Generally speaking, tasks that have been given less processing time (smaller values of `vruntime`) are toward the left side of the tree, and tasks that have been given more processing time are on the right side. According to the properties of a binary search tree, the leftmost node has the smallest key value, which for the sake of the CFS scheduler means that it is the task with the highest priority. Because the red-black tree is balanced, navigating it to discover the leftmost node will require  $O(\lg N)$  operations (where  $N$  is the number of nodes in the tree). However, for efficiency reasons, the Linux scheduler caches this value in the variable `rb_leftmost`, and thus determining which task to run next requires only retrieving the cached value.

# Linux Scheduling (Cont.)

- Real-time processes in POSIX.1b
  - Real-time tasks have static priorities
  - Real-time plus normal map into global priority scheme
- Normal processes
  - Nice value of -20 maps to global priority 100
  - Nice value of +19 maps to priority 139



End of Chapter 5