

Chapter 10: File-System Interface

Prof. Li-Pin Chang
CS@NYCU

Chapter 10: File-System Interface

- File Concept
- Directory Structure
- File and directory operations
- File aliasing
- File-System Mounting
- File permission and protection

Objectives

- To explain the function of file systems
- To describe the interfaces to file systems
- To discuss file-system design tradeoffs, including access methods, file sharing, file locking, and directory structures
- To explore file-system protection

File Concept

- A computer resource to write data to and read data from storage device
- A contiguous logical address space

UNIX file types:

- Regular files
- Device files (device node)
- Directory files
- Links

File Attributes

- **Name** – only information kept in human-readable form
- **Identifier** – unique tag (number) identifies file within file system
- **Type** – needed for systems that support different types
 - Regular, directory, device, link (system functionality)
 - .c , .exe, .bat (user purpose)
- **Location** – pointer to file location on device
- **Size** – current file size
- **Protection** – controls who can do reading, writing, executing
- **Time, date, and user identification** – data for protection, security, and usage monitoring

File Types – Name, Extension

file type	usual extension	function
executable	exe, com, bin or none	ready-to-run machine- language program
object	obj, o	compiled, machine language, not linked
source code	c, cc, java, pas, asm, a	source code in various languages
batch	bat, sh	commands to the command interpreter
text	txt, doc	textual data, documents
word processor	wp, tex, rtf, doc	various word-processor formats
library	lib, a, so, dll	libraries of routines for programmers
print or view	ps, pdf, jpg	ASCII or binary file in a format for printing or viewing
archive	arc, zip, tar	related files grouped into one file, sometimes com- pressed, for archiving or storage
multimedia	mpeg, mov, rm, mp3, avi	binary file containing audio or A/V information

File Operations

Function	description
fopen()	create a new file or open a existing file
fclose()	closes a file
getc()	reads a character from a file
putc()	writes a character to a file
fscanf()	reads a set of data from a file
fprintf()	writes a set of data to a file
fread()	reads a number of bytes from a file
fwrite()	writes a number of bytes to a file
fseek()	set the position to desire point
link()	make a new name for a file
unlink()	decrement the reference count of a file (delete on ref=0)

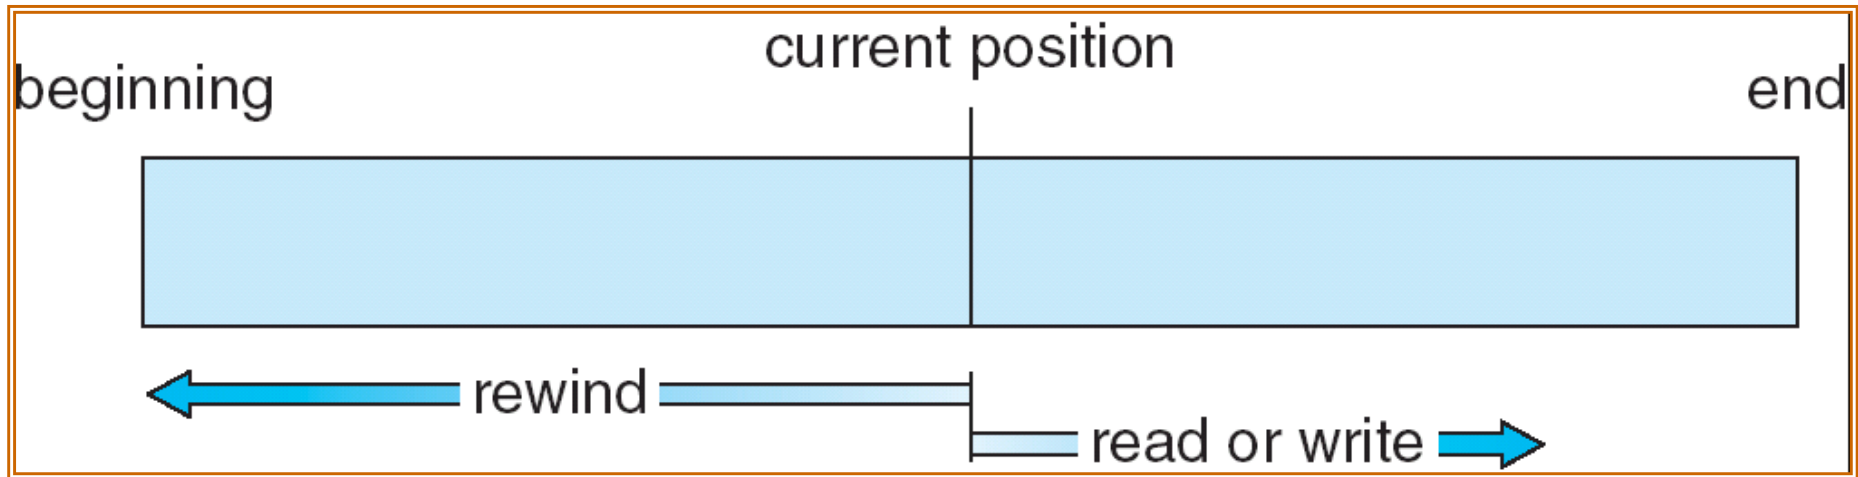
Why Opening/Closing Files

- Information required to manage opened files:
 - File pointer: pointer to last read/write location, per process that has the file open
 - File-open count: counter of number of times a file is open – to allow removal of data from open-file table when last processes closes it (e.g., removal of USB drives)
 - Disk location of the file
 - Access rights: per-process access mode information
- These are called “**metadata**”, i.e., data of data
- The file system caches metadata when opening files for efficient operations; it also flushes modified metadata to disk when closing files

fopen(): Binary or Text?

- `fopen("abc.txt","r+t");`
- `fopen("xyz.mp3","rb");`
- Text mode
 - Translate Ctrl-Z (1A) into EOF
 - Translation between `\r\n` and `\n` for different OSes
 - UNIX: `\n` (0A) Windows `\r\n` (0D 0A)
 - Possibility of filtering out the MSB (only 7 LSBs used)
- Binary mode
 - Raw input

File Accessing Model




Device Node

- Commonly appear under the `/dev` directory
- Can be manually created using `mknode` command, with proper device major-minor #'s assigned
- Device drivers register themselves to the kernel using the device major-minor #'s
- `open()`, `close()`, `read()`, `write()` a device node will communicate with the device driver registered with the same major-minor #'s as that of the device node
- Example: `open ()` on `/dev/sda` → #M8m0

Directory

- A collection of nodes containing information about all files



Directory itself is a file, too

Directory Operations

- Search for a file
- Create a file
- Delete a file
 - If the deleted file is a directory?
 - Recursively delete all its files and sub-directories?
 - If the deleted file is a regular file?
- Directory enumeration (listing)
- Rename a file

Open and read a directory

```
DIR *opendir(const char *name);
struct dirent *readdir(DIR *dirp);

struct dirent {
    ino_t      d_ino;      /* Inode number */
    off_t      d_off;      /* Not an offset; see below */
    unsigned short d_reclen; /* Length of this record */
    unsigned char d_type;   /* Type of file; not supported
                             by all filesystem types */
    char       d_name[256]; /* Null-terminated filename */
};

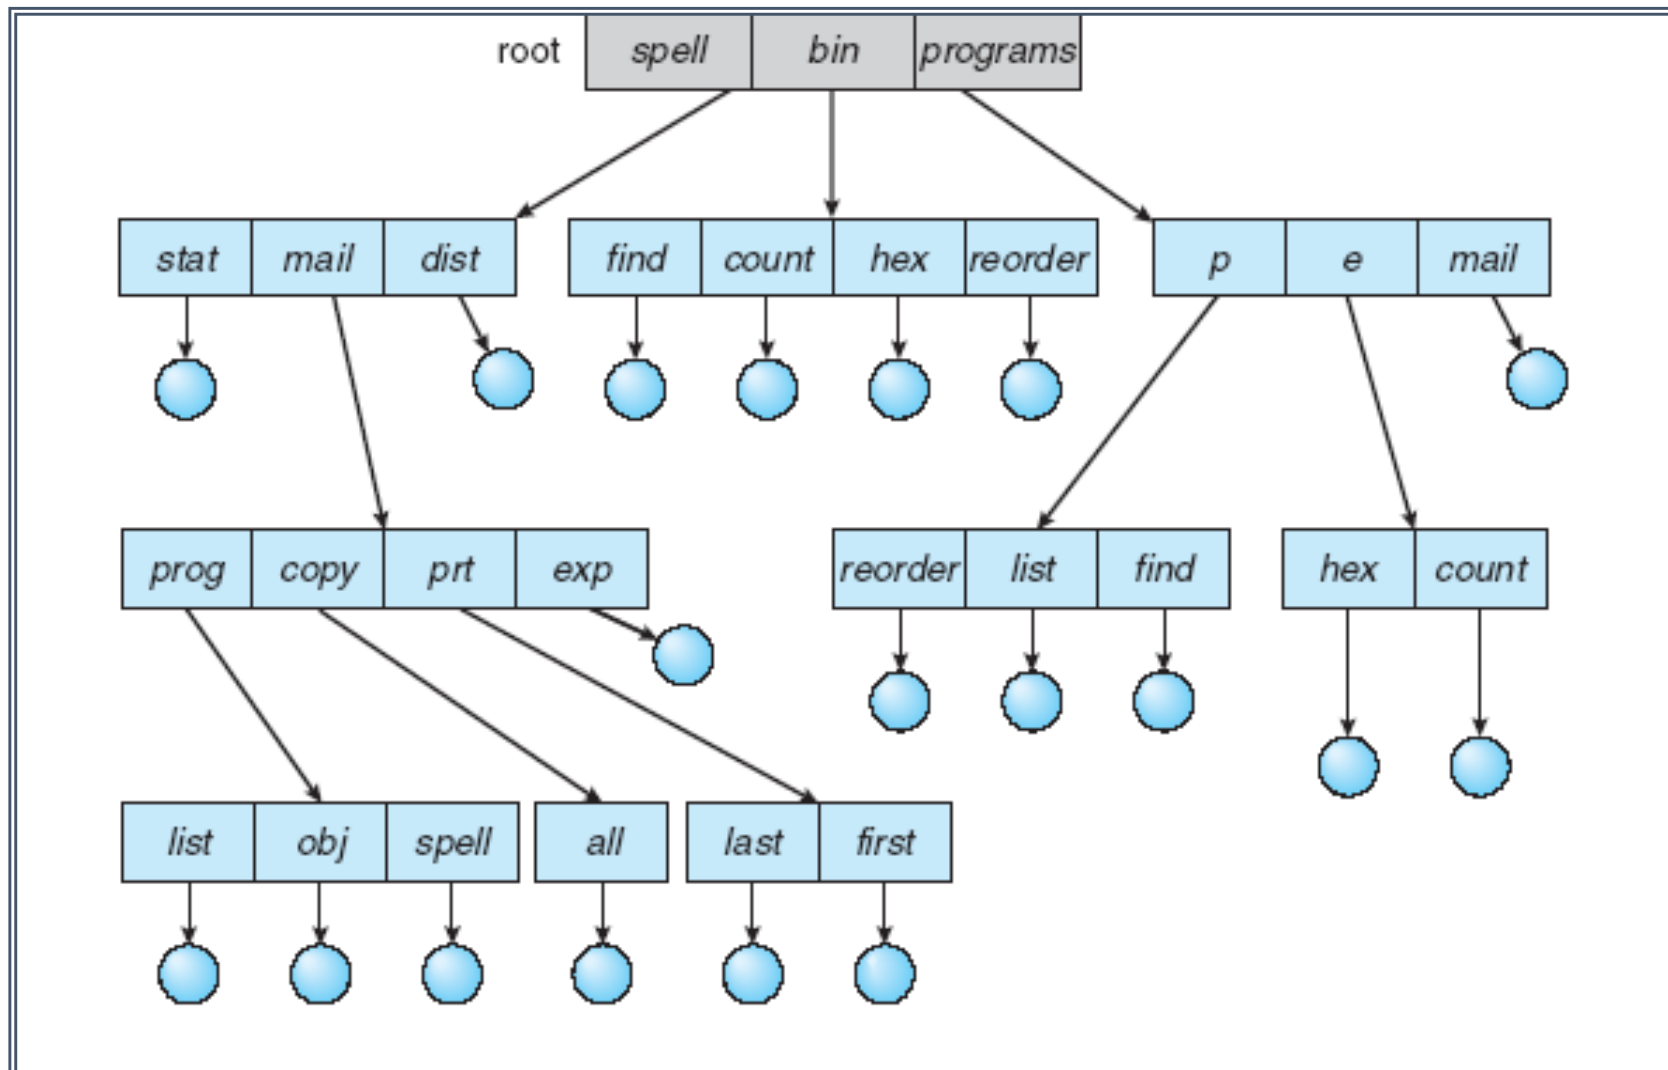
// -----

#include <sys/types.h>
#include <dirent.h>

DIR *dir;
struct dirent *dirp;

dir = opendir("foo");
dirp = readdir(dir);
dirp = readdir(dir);
dirp = readdir(dir);
dirp = readdir(dir);
```

Tree-Structured Directories



Tree-Structured Directories (Cont)

- The *current working directory (CWD)* environment variable (per process)
 - “.” and “..”
- Absolute or relative path name
- Traverse the file system

```
char *getcwd(char *buf, size_t size);  
int chdir(const char *path);
```


Tree-Structured Directories (Cont)

- Creating a new file is done in current directory
- Delete a file


`rm <file-name>`

- Creating a new subdirectory is done in current directory

`mkdir <dir-name>`

Example: if in current directory `/mail`

`mkdir count`

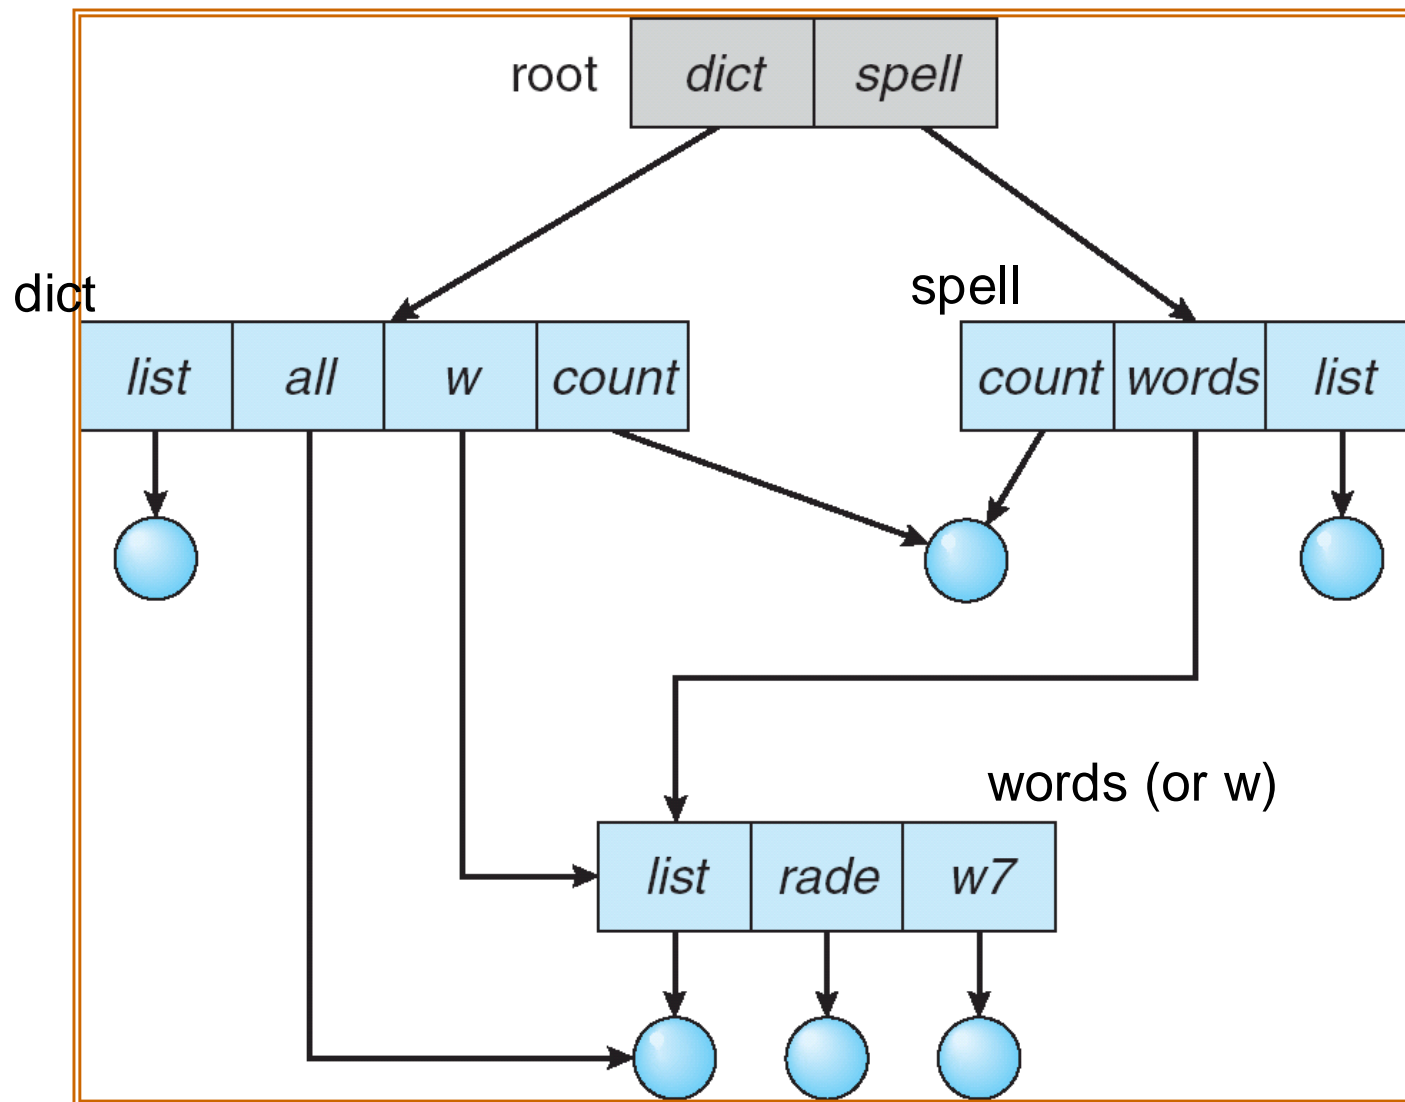


Deleting “mail” \Rightarrow deleting the entire subtree rooted by “mail”
`rm -r` or `del /s`

File Aliasing (Link)

- A file may have two different names (alias)
- A file link
 - Another name of (pointer to) an existing file
 - Resolve the link – follow pointer to locate the file

Acyclic-Graph Directories



Softlinks

- Softlinks (symbolic link)
 - String substitution
 - Independent of file system
 - Appearing as a link file
- Usage
 - UNIX: `ln -s [target] [link]`
 - Windows (NTFS): `junction.exe [link] [target]`

<https://tw.arip-photo.org/736330-how-to-list-symbolic-link-WLWBSA>

```
[root@localhost ~]# ln -s ./test/simpleText.txt ./simpleText
[root@localhost ~]# ls -l
total 16
drwxr-xr-x  3 root    root      163 Aug 21  2011 dos
-rw-r--r--  1 root    root      242 Jul 15  2017 hello.c
lrwxrwxrwx  1 root    root       21 Feb 21 22:22 simpleText -> ./test/si
pleText.txt
drwxr-xr-x  2 root    root       68 Feb 21 22:13 test
[root@localhost ~]#
```

Hardlinks

- Hardlinks
 - A link file that refers to the target file using file system internal location information
 - File-system-dependent
 - Nothing different from a regular file
 - The target file has a link count > 1 ; use `unlink()` to delete files
- Usage
 - UNIX: `ln [target] [link]`
 - Windows (NTFS): `fsutil hardlink create [link] [target]`

Problems with Aliasing

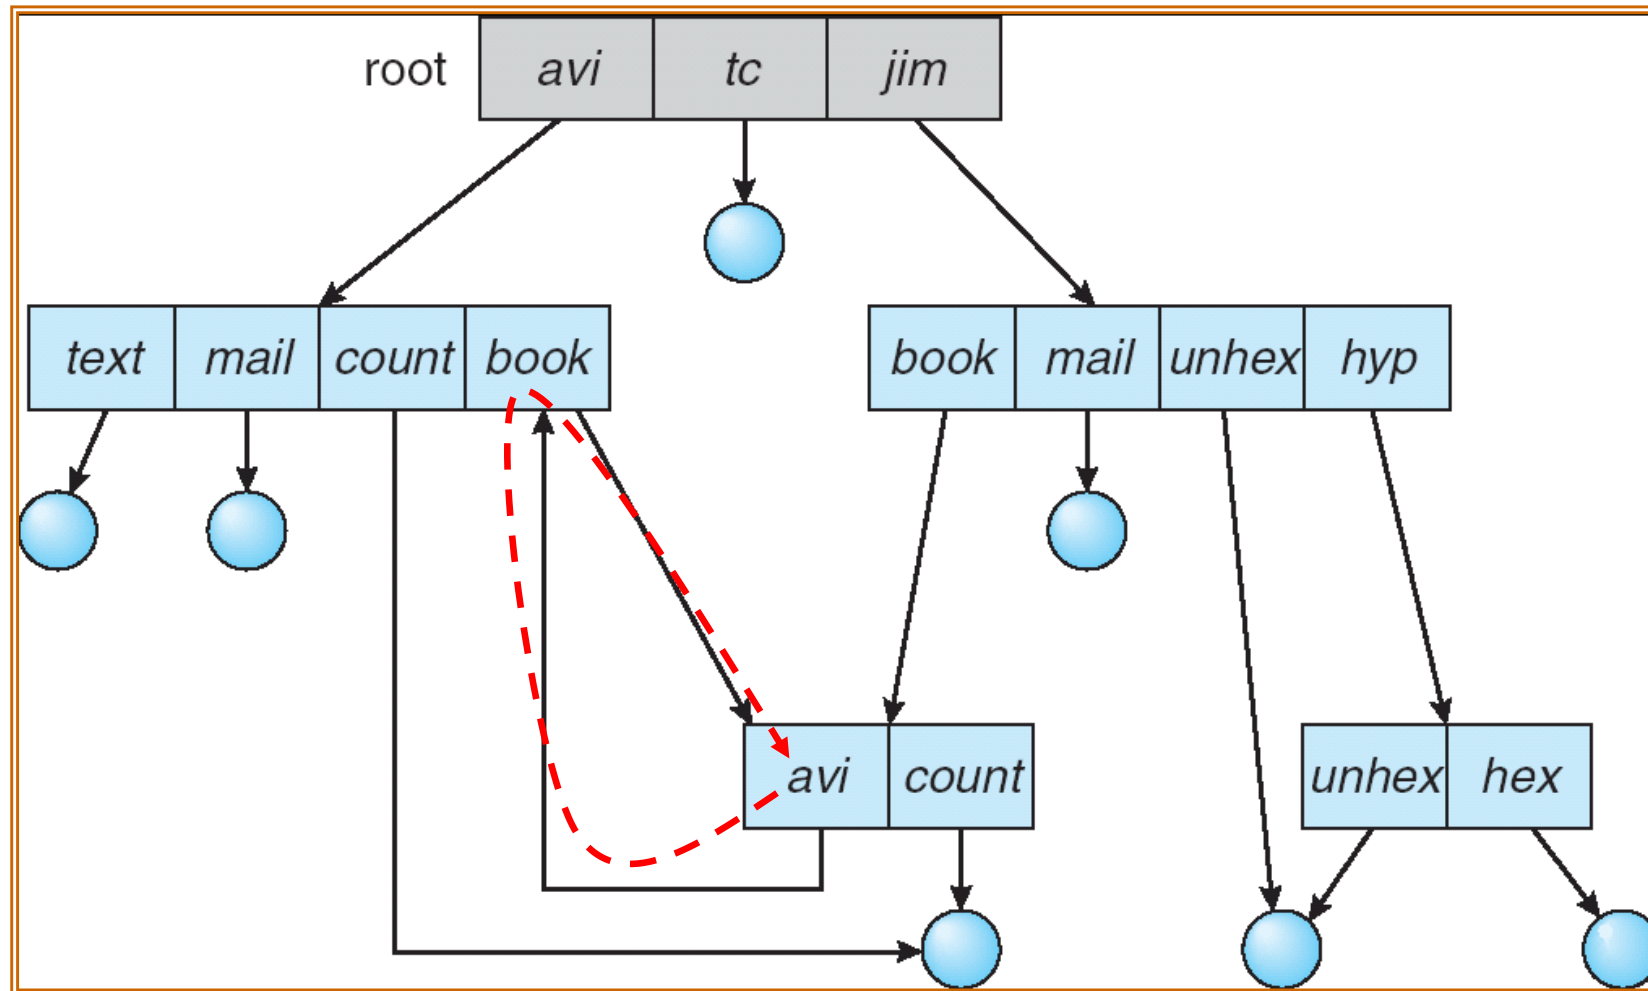
- **Backup**— Duplication problem
 - May duplicate files during backup
 - “cp -a” or “rsync” to preserve hard links as many as possible

Problems with Aliasing

- **Loop** – Endless file path
- Loops caused by hard links
 - Hard links to directories are forbidden in recent UNIX implementations
 - Every time a new link is added use a cycle detection algorithm to determine whether it is OK (less practical)
- Loops caused by soft links
 - Soft links to directories are still possible
 - Linux: Keep a time-to-live counter (e.g., 40)
 - Windows: Limiting the pathname length (~ 260 chars)

Modern UNIX implementations do not allow *hard links* to directories! (soft links are still possible)

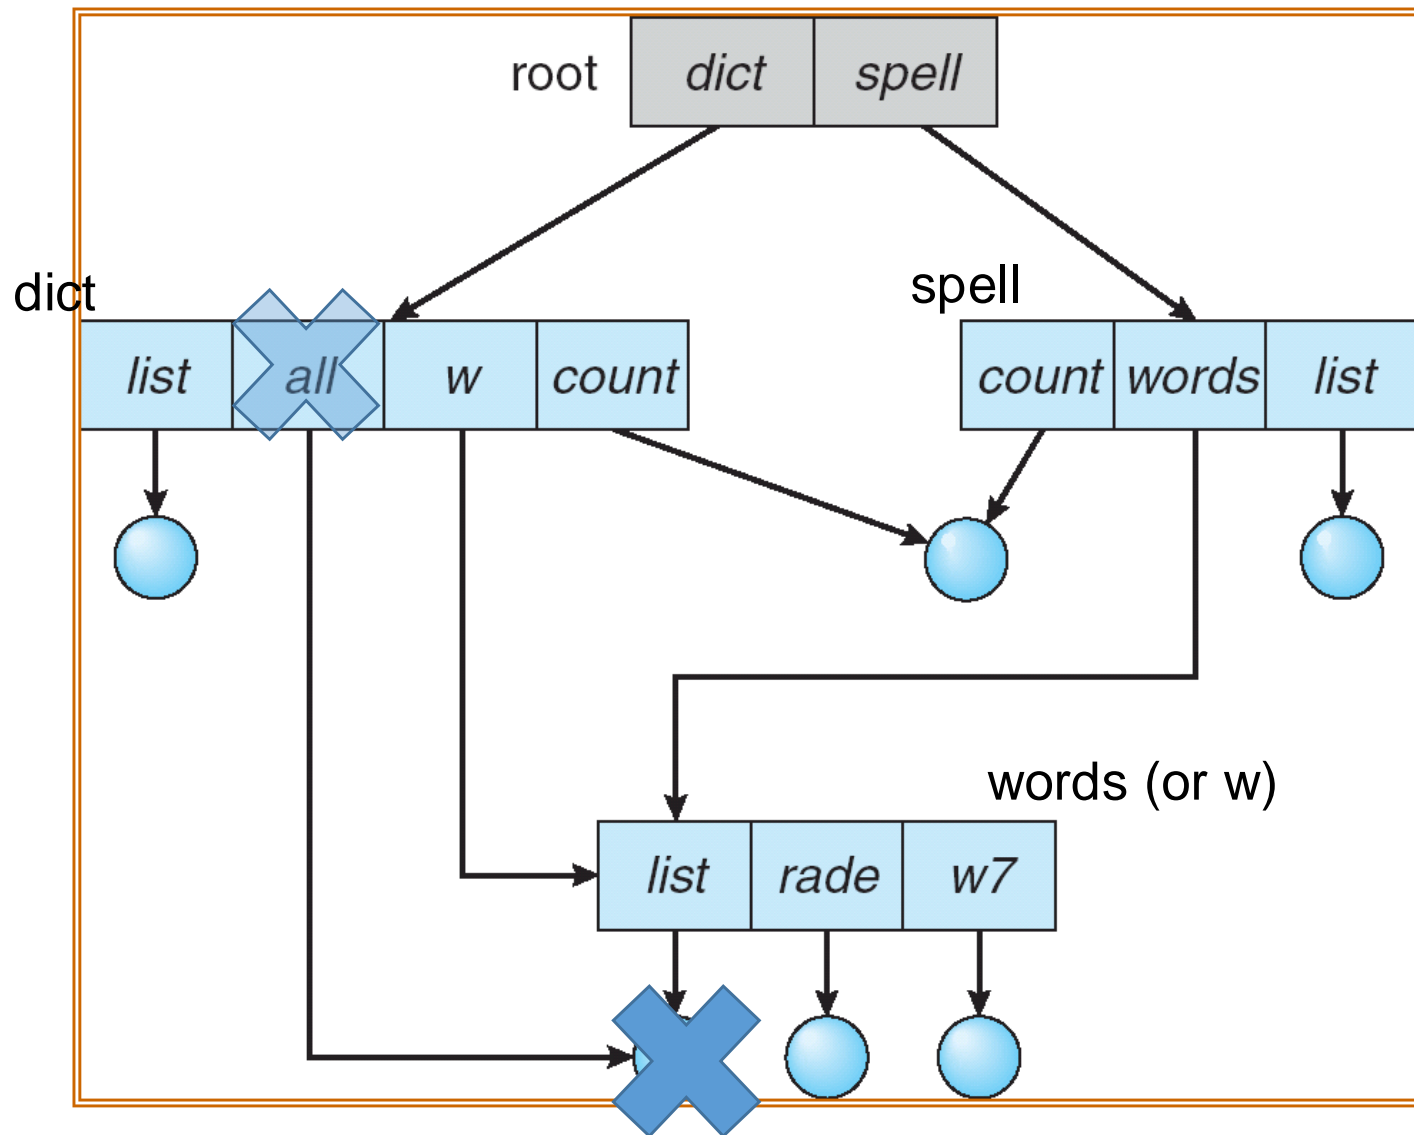
Loop in directories



Problems with Aliasing

- **Deletion**— Dangling pointer problem
 - Deleting “all” in dict makes the symbolic link “list” dangling
- Solutions:
 - Hard links require proper management of dangling pointers as referring to undefined storage address may expose security issues
 - ✓ (hard link) Backpointers, so we can delete all pointers
 - ✓ (hard link) Entry-hold-count solution (unlink() in UNIX)
 - Soft links are less problematic
 - ✓ (sym link) Leave a symbolic link dangling

Acyclic-Graph Directories



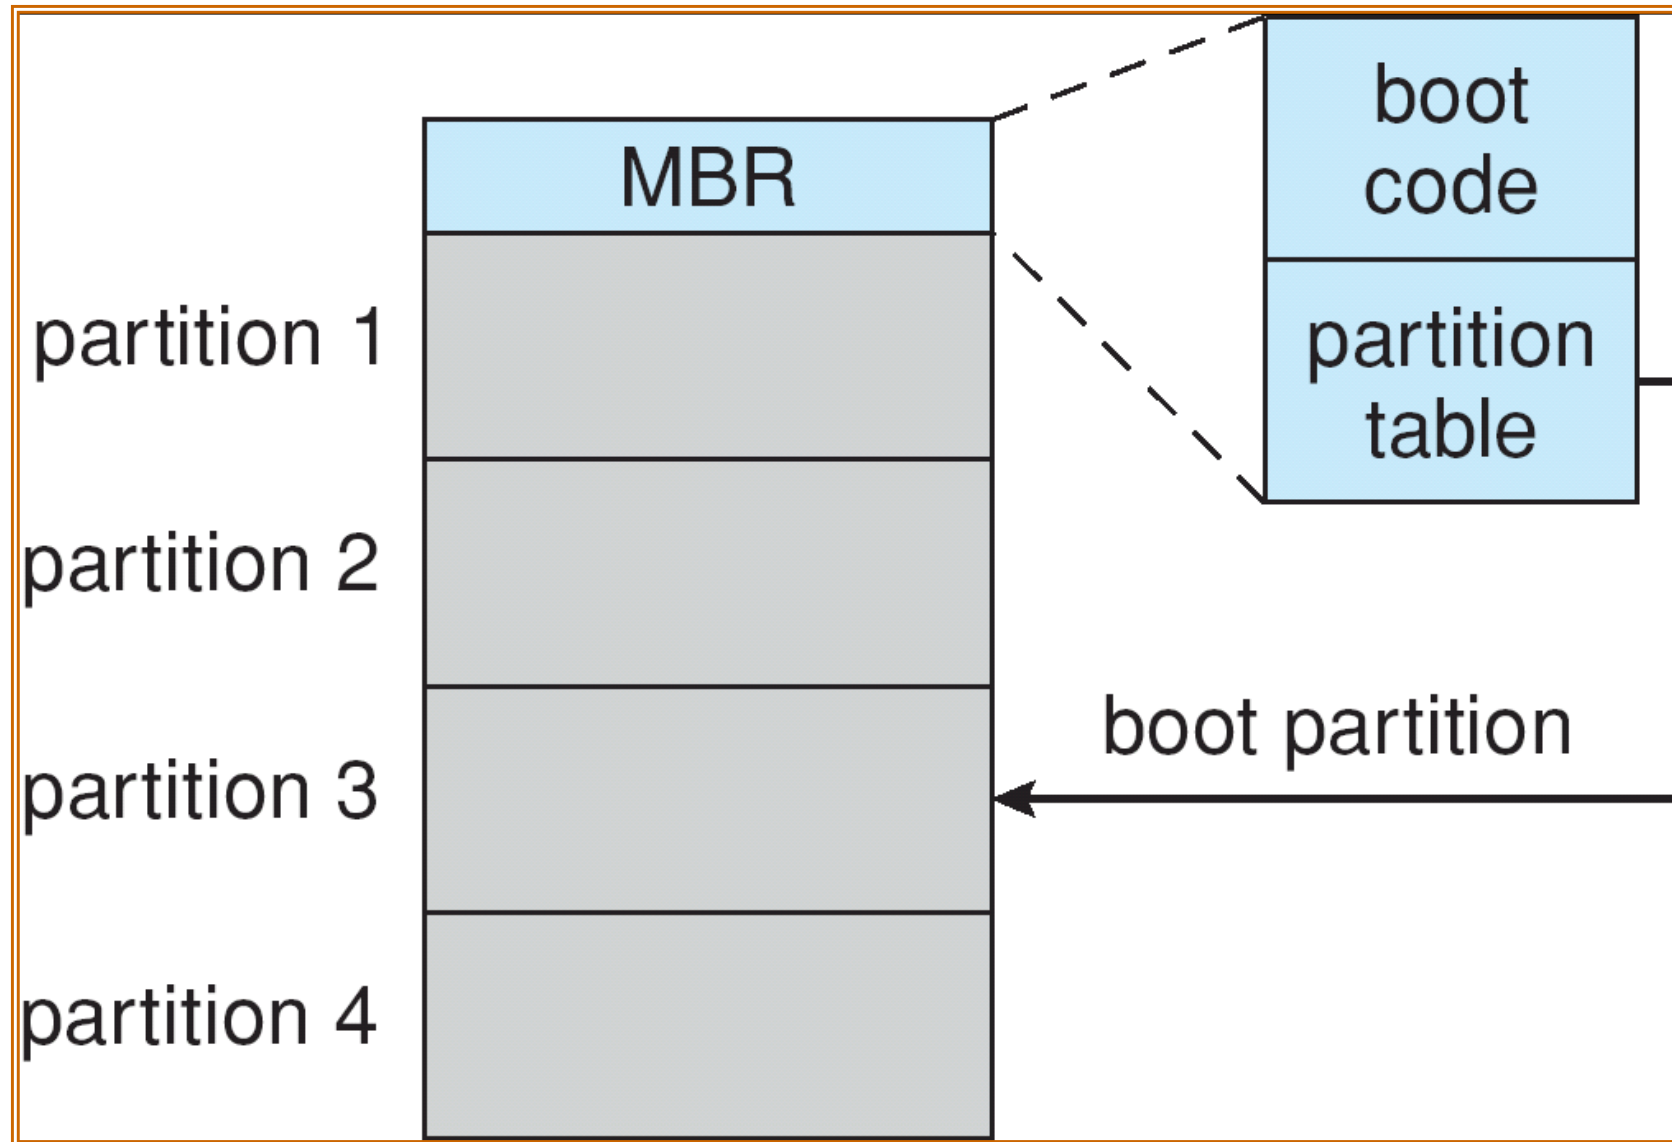
Problems with Aliasing

- Dangling pointers
- Softlink (symbolic link)
 - Simply leave the symbolic link dangling
 - `/bin/l`s \rightarrow `/sbin/l`s
- Hardlink
 - link is established inside the file system
 - Keep a reference count
 - Creating hardlink to the file: `+count`
 - Removing a hardlink to the file: `-count`
 - When `count==0`: remove the file

Soft link vs. Hard link: Revisit

- Softlink
 - Can span over different file systems
 - Dangling pointer problem
- Hardlink
 - More efficient than soft links
 - Can not span over different file systems
 - Often confusing, cannot tell which file is the “original one”

Legacy MBR Partitions



Partitioning a Disk

- The very first step of preparing a hard drive
- Use fdisk or other GUI utilities
- Partitions can be formatted into different file systems or used as a swap device
- The bios loads the MBR, which in turn loads the next loader in the boot partition
 - An OS or a boot manager
- MBR partition tables are being replaced by GPT, which allows larger partition sizes and unlimited partitions

Formatting a Partition

- To use a disk to hold files, the operating system still needs to record its own data structures on the disk.
 - Logical formatting, high-level formatting or “making a file system”.
 - Writing file system metadata
- Low-level formatting, or physical formatting — Dividing a disk into sectors that the disk controller can read and write.
 - Remapping bad tracks to spare tracks
 - Zoned-bit encoding

Mounting a File System

- A file system must be **mounted** before it can be accessed
- A unmounted file system is mounted at a **mount point**
- Mounting a file system
 - `mount -t ext4 /users /dev/hda1`
 - Specify the **file system type**
 - Find the file-system superblock in the **partition device node**
 - Specify the **mounting point** of the file-system naming space

Protection

- File owner/creator should be able to control:
 - what can be done
 - by whom
- Types of access
 - Read
 - Write
 - Execute
 - Append (regards to disk space)
 - Delete
 - List

FTP

File Sharing – Multiple Users


- **User IDs** identify users, allowing permissions and protections to be per-user
- **Group IDs** allow users to be in groups, permitting group access rights

Access Lists and Groups

- Mode of access: read, write, execute
- Three classes of users

			RWX
a) owner access	7	⇒	1 1 1 RWX
b) group access	6	⇒	1 1 0 RWX
c) public access	1	⇒	0 0 1

- Ask manager to create a group (unique name), say G, and add some users to the group.
- For a particular file (say *game*) or subdirectory, define an appropriate access.
- Attach a group to a file: `chgrp G game`



UNIX File Permission Management Utilities

- adduser: create a user
 - mkgrp: create a group
 - addgrp: add a user to a group
 - chown: change the owner of a file
 - chgrp: change the group of a file
 - chmod: change file permissions
-
- Users are managed by /etc/password
 - Groups are managed by /etc/group

A Sample UNIX Directory Listing

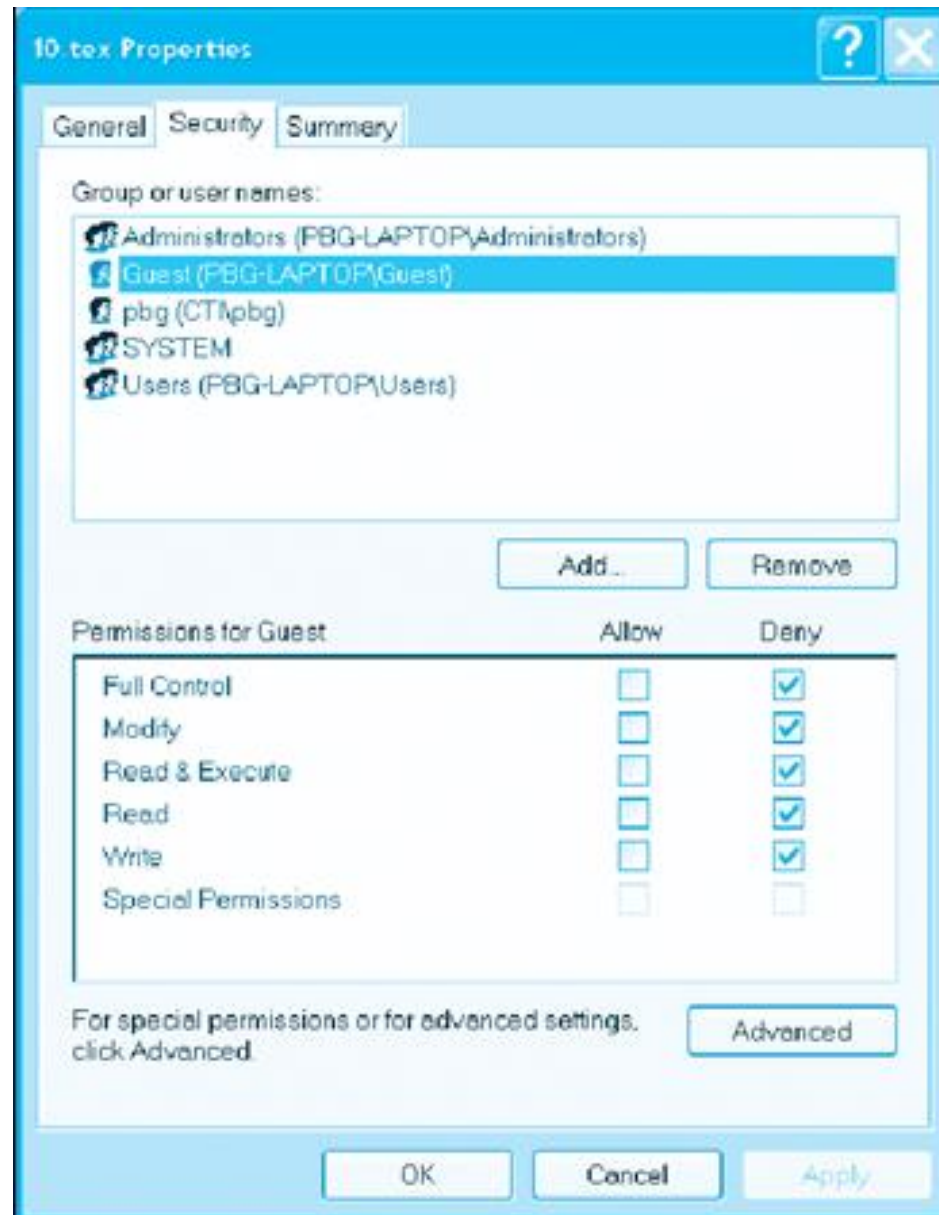
-rw-rw-r--	1	pbg	staff	31200	Sep 3 08:30	intro.ps
drwx-----	5	pbg	staff	512	Jul 8 09:33	private/
drwxrwxr-x	2	pbg	staff	512	Jul 8 09:35	doc/
drwxrwx---	2	pbg	student	512	Aug 3 14:13	student-proj/
-rw-r--r--	1	pbg	staff	9423	Feb 24 2003	program.c
-rwxr-xr-x	1	pbg	staff	20471	Feb 24 2003	program
drwx--x--x	4	pbg	faculty	512	Jul 31 10:31	lib/
drwx-----	3	pbg	staff	1024	Aug 29 06:52	mail/
drwxrwxrwx	3	pbg	staff	512	Jul 8 09:35	test/

[Permission] [hard link count][Owner] [group] [filesize] [date] [filename]

- Regular file: link count ≥ 1 , file is deleted when link count = 0
- A directory: link count is $2+n$
 - 1 from its own directory entry + 1 from “.” of itself
 - n from “..” of all its sub-directories

A directory with the permission “x” = the directory can be searched/entered

Windows XP Access-control List Management



End of Chapter 10