

(1)

- Preprocess the `x_test` data by divided by 255 to normalize it.
- I use 4 convolution layers to build the model, which in order is 32, 64, 64, 128.
- The other hyperparameters are the same: `kernel_size=(3, 3)`, `padding=same`, `kernel_regularizer=l2(5e-4)`, `activation=relu` and add `batch_normalization` after each of the layer. I also add `dropout(0.25)` and add `maxpooling(2, 2)` to reduce overfitting and lower the spatial dimension.
- `Batch_size` is altered to 64 with 50 epochs and `learning_rate=1e-3`.
- Instead of using SGD, I use Adam as optimizer with `beta_1`(serve as momentum)=0.9, `beta_2`=0.999, `epsilon=1e-4` and `decay=learning_rate/epochs` (which decreases `learning_rate` gradually).
- `ImageDataGenerator` is added as data augmentation (preprocessing).
- Train data is divided into 9:1 as training set and validation set.

Test Result:

```
✓ 1 秒 [10] y_pred = model.predict(x_test)  
      y_pred = np.argmax(y_pred, axis=1)
```

```
[ ] assert y_pred.shape == (10000,)
```

```
(10000,)
```

```
✓ 0 秒 [11] y_test = np.load("y_test.npy")  
      print("Accuracy of my model on test set: ", accuracy_score(y_test, y_pred))
```

```
Accuracy of my model on test set: 0.8337
```

(2) Used packages are attached in “requirements.txt”

(3) unzip the RAR file, then click on “inference.py”, then you can reproduce the result.