



# Pattern Recognition

## Dimensionality Reduction

林彥宇 教授

Yen-Yu Lin, Professor


國立陽明交通大學 資訊工程學系

Computer Science, National Yang Ming Chiao Tung University

Some slides are modified from S.-J. Wang  
and T.-L. Liu



# Dimensionality reduction

- Many data sets have the property that the data points all lie close to **a manifold of much lower dimensionality** than that of the original data space
- A synthetic data set obtained by taking one of the off-line digit images and creating multiple copies via random displacements and rotations



- Dimension of the data space:  $100 \times 100 = 10,000$
- **Degrees of freedom**: vertical translation, horizontal translation, rotation
- Intrinsic dimension: 3

# Data dimension vs. intrinsic dimension



# Principal component analysis (PCA)

- Principal component analysis (PCA) is one of the most widely used **dimensionality reduction** techniques
- PCA is used for various applications, such as data compression, feature extraction, and data visualization
- PCA linearly projects data from a high-dimensional input space to a low-dimensional feature space
  - **Principal subspace**: the lower dimensional space obtained by PCA
  - **Principal component**: a projection direction/vector found by PCA
  - A projection is composed of one or a few principal components

# Definitions of PCA

- Two commonly used definitions:
  - **Maximum variance of the projected data** (Hotelling 1933): The orthogonal projection of the data onto a lower dimensional linear space such that the variance of the projected data is maximized
  - **Minimum projection error** (Pearson 1901): The linear projection that minimizes the average projection cost, defined as the mean squared distance between the data points and their projections
- The two different definitions lead to **the same algorithm, PCA**

# Definitions of PCA

- Maximum variance vs. minimum projection error


**Points:** data

**Points:** projected data

**Line:** 1D principal space

Definition 1: Maximize the variance of **green points**

Definition 2: Minimize the sum of squared projection errors, indicated by **blue lines**



# Maximum variance formulation

- Consider a data set of observations  $\{\mathbf{x}_n\}$  where  $n = 1, 2, \dots, N$  and  $\mathbf{x}_n$  in an Euclidean space of dimensionality  $D$
- The goal of PCA is to project the data onto a space having dimensionality  $M < D$  while maximizing the variance of the projected data
- Consider the projection onto a one-dimensional space ( $M = 1$ )
- We can define the direction of this space using a  $D$ -dimensional unit vector  $\mathbf{u}_1$  and  $\mathbf{u}_1^T \mathbf{u}_1 = 1$ 
  - Data point:  $\mathbf{x}_n$
  - Projected data point:  $\mathbf{u}_1^T \mathbf{x}_n$

# Maximum variance formulation

- Data mean and projected data mean

$$\bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n \quad \text{and} \quad \mathbf{u}_1^T \bar{\mathbf{x}}$$

- The variance of the projected data is

$$\frac{1}{N} \sum_{n=1}^N \{ \mathbf{u}_1^T \mathbf{x}_n - \mathbf{u}_1^T \bar{\mathbf{x}} \}^2 = \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1$$

where  $\mathbf{S}$  is the covariance matrix

$$\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T.$$



# Maximum variance formulation: Optimization

- Constrained optimization problem

$$\begin{aligned} & \text{maximize} \quad \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 \\ & \text{subject to} \quad \mathbf{u}_1^T \mathbf{u}_1 = 1 \end{aligned}$$

- Introduce a Lagrange multiplier to convert the constrained optimization problem to an unconstrained one, where the **Lagrangian function** is

$$\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 + \lambda_1 (1 - \mathbf{u}_1^T \mathbf{u}_1) .$$

- By setting the derivative w.r.t.  $\mathbf{u}_1$  to zero, we have

$$\mathbf{S} \mathbf{u}_1 = \lambda_1 \mathbf{u}_1$$

- $\mathbf{u}_1$  is an eigenvector of  $\mathbf{S}$

# Maximum variance formulation: Optimization

- The objective in the constrained optimization problem of PCA is

$$\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 = \lambda_1$$

- The maximum variance of the projected data is equal to the largest eigenvalue  $\lambda_1$  of the covariance matrix  $\mathbf{S}$
- The projection vector  $\mathbf{u}_1$  is the eigenvector of  $\mathbf{S}$  corresponding to the largest eigenvalue  $\lambda_1$
- The eigenvector  $\mathbf{u}_1$  is known as the first principal component, which is used to project a data point  $\mathbf{x}_n$  via  $\mathbf{u}_1^T \mathbf{x}_n$

# Maximum variance formulation: Multi-dimensional extension

- Consider the general case of an  $M$ -dimensional projection space
- The optimal linear projection is defined by the  $M$  eigenvectors  $\mathbf{u}_1, \mathbf{u}_2, \dots$ , and  $\mathbf{u}_M$  of the data covariance matrix  $S$  corresponding to the  $M$  largest eigenvalues  $\lambda_1, \lambda_2, \dots$ , and  $\lambda_M$ .
- Computational cost of PCA
  - Full eigenvector decomposition:  $O(D^3)$
  - Find the first  $M$  eigenvectors and eigenvalues via the power method:  $O(MD^2)$

# Minimum error formulation

- We introduce a complete orthonormal set of  $D$ -dimensional basis vectors  $\{\mathbf{u}_i\}$  where  $i = 1, 2, \dots, D$
- Each data point can be represented exactly by a linear combination of the basis vectors

$$\mathbf{x}_n = \sum_{i=1}^D \alpha_{ni} \mathbf{u}_i$$

➤ where  $\alpha_{ni} = \mathbf{x}_n^T \mathbf{u}_i$

- We have

$$\mathbf{x}_n = \sum_{i=1}^D (\mathbf{x}_n^T \mathbf{u}_i) \mathbf{u}_i.$$

# Minimum error formulation

- We aim to approximate this data point using a representation involving a restricted number  $M < D$  of variables, which corresponds to a projection onto a lower-dimensional subspace
- Without loss of generality, we approximate each data point  $\mathbf{x}_n$  by

$$\tilde{\mathbf{x}}_n = \sum_{i=1}^M z_{ni} \mathbf{u}_i + \sum_{i=M+1}^D b_i \mathbf{u}_i$$

- Each point  $\mathbf{x}_n$  has its own coefficients  $\{z_{ni}\}$  for the first  $M$  basis vectors
- $\{b_i\}$  are constants that are the same for all data points
- In minimum error formulation, we optimize  $\{\mathbf{u}_i\}$ ,  $\{z_{ni}\}$ , and  $\{b_i\}$

# Minimum error formulation: Optimization

- In this formulation, we minimize the squared projection error, the squared distance, i.e.,

$$J = \frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \tilde{\mathbf{x}}_n\|^2.$$

- Setting the derivative of  $J$  with respect to  $\{z_{nj}\}$  to zero, we get

$$z_{nj} = \mathbf{x}_n^T \mathbf{u}_j \quad \text{where } j = 1, 2, \dots, M$$

- Setting the derivative of  $J$  with respect to  $\{b_j\}$  to zero, we get

$$b_j = \bar{\mathbf{x}}^T \mathbf{u}_j \quad \text{where } j = M + 1, M + 2, \dots, D$$



# Minimum error formulation: Optimization

- If we substitute for  $\{z_{ni}\}$  and  $\{b_i\}$ , the displacement vector between  $\mathbf{x}_n$  and its approximated point  $\tilde{\mathbf{x}}_n$

$$\mathbf{x}_n - \tilde{\mathbf{x}}_n = \sum_{i=M+1}^D \left\{ (\mathbf{x}_n - \bar{\mathbf{x}})^T \mathbf{u}_i \right\} \mathbf{u}_i$$

- The displacement vector lies in the space orthogonal to the principal space
- The squared projection error

$$J = \frac{1}{N} \sum_{n=1}^N \sum_{i=M+1}^D \left( \mathbf{x}_n^T \mathbf{u}_i - \bar{\mathbf{x}}^T \mathbf{u}_i \right)^2 = \sum_{i=M+1}^D \mathbf{u}_i^T \mathbf{S} \mathbf{u}_i.$$

# Minimum error formulation: Optimization

$$J = \sum_{i=M+1}^D \mathbf{u}_i^T \mathbf{S} \mathbf{u}_i$$

- Consider the example of a two-dimensional data space with one-dimensional projection space, i.e.,  $D = 2$  and  $M = 1$
- We want to optimize  $\mathbf{u}_2$  so as to minimize  $J = \mathbf{u}_2^T \mathbf{S} \mathbf{u}_2$ , subject to  $\mathbf{u}_2^T \mathbf{u}_2 = 1$ , i.e.,

$$\begin{aligned} & \text{minimize} \quad \mathbf{u}_2^T \mathbf{S} \mathbf{u}_2 \\ & \text{subject to} \quad \mathbf{u}_2^T \mathbf{u}_2 = 1 \end{aligned}$$

- The Lagrangian function is

$$\tilde{J} = \mathbf{u}_2^T \mathbf{S} \mathbf{u}_2 + \lambda_2 (1 - \mathbf{u}_2^T \mathbf{u}_2) .$$

- Setting the derivative with respect to  $\mathbf{u}_2$  to zero leads to

$$\mathbf{S} \mathbf{u}_2 = \lambda_2 \mathbf{u}_2$$

- $\mathbf{u}_2$  is an eigenvector of  $\mathbf{S}$  with eigenvalue  $\lambda_2$





# Minimum error formulation: Optimization

- The objective in this formulation is to minimize

$$\mathbf{u}_2^T \mathbf{S} \mathbf{u}_2 = \mathbf{u}_2^T \lambda_2 \mathbf{u}_2 = \lambda_2$$

- The basis  $\mathbf{u}_2$  is the eigenvector of covariance matrix  $\mathbf{S}$  with the smallest eigenvalue
- Recall that

$$\tilde{\mathbf{x}}_n = \sum_{i=1}^M z_{ni} \mathbf{u}_i + \sum_{i=M+1}^D b_i \mathbf{u}_i$$

- The projection vector  $\mathbf{u}_1$  is the eigenvector of  $\mathbf{S}$  corresponding to the largest eigenvalue  $\lambda_1$



# Minimum error formulation: Multi-dimensional extension

- For arbitrary  $D$  and arbitrary  $M < D$ , the optimal linear projection is defined by the  $M$  eigenvectors  $\mathbf{u}_1, \mathbf{u}_2, \dots$ , and  $\mathbf{u}_M$  of the data covariance matrix  $S$  corresponding to the  $M$  largest eigenvalues  $\lambda_1, \lambda_2, \dots$ , and  $\lambda_M$
- The approximated point

$$\begin{aligned}\tilde{\mathbf{x}}_n &= \sum_{i=1}^M (\mathbf{x}_n^T \mathbf{u}_i) \mathbf{u}_i + \sum_{i=M+1}^D (\bar{\mathbf{x}}^T \mathbf{u}_i) \mathbf{u}_i \\ &= \bar{\mathbf{x}} + \sum_{i=1}^M (\mathbf{x}_n^T \mathbf{u}_i - \bar{\mathbf{x}}^T \mathbf{u}_i) \mathbf{u}_i\end{aligned}$$

# Algorithm summary

- 1. We are given a set of  $N$  unlabeled data  $\{\mathbf{x}_n\}$ . Data lie in  $D$ -dimensional space. The desired dimension of the projection space is  $M$
- 2. Compute the covariance matrix

$$\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T.$$


- 3. Perform eigen-decomposition
  - Get  $M$  eigenvectors  $\mathbf{u}_1, \mathbf{u}_2, \dots$ , and  $\mathbf{u}_M$  of  $\mathbf{S}$  corresponding the largest  $M$  eigenvalues
- 4. Approximate the data  $\{\mathbf{x}_n\}$  via

$$\tilde{\mathbf{x}}_n = \bar{\mathbf{x}} + \sum_{i=1}^M (\mathbf{x}_n^T \mathbf{u}_i - \bar{\mathbf{x}}^T \mathbf{u}_i) \mathbf{u}_i$$




# Applications of PCA: Data compression

- Off-line digit data set: handwritten digit images, each of which is of resolution  $28 \times 28$  (784)
- Because each eigenvector of the covariance matrix is a vector in the original  $D$ -dimensional space, it can be displayed as an image
- The mean vector (image) and the first four PCA eigenvectors  $\mathbf{u}_1$ ,  $\mathbf{u}_2$ ,  $\mathbf{u}_3$ , and  $\mathbf{u}_4$ , together with the corresponding eigenvalues




# Applications of PCA: Data compression

- The spectrum of eigenvalues, sorted into decreasing order



- The distortion measure (squared projection error)  $J$  associated with a particular value of  $M$

$$J = \sum_{i=M+1}^D \lambda_i$$



# Applications of PCA: Data compression


- PCA is used to approximate each input data point  $\mathbf{x}_n$  via

$$\tilde{\mathbf{x}}_n = \bar{\mathbf{x}} + \sum_{i=1}^M (\mathbf{x}_n^T \mathbf{u}_i - \bar{\mathbf{x}}^T \mathbf{u}_i) \mathbf{u}_i$$

- The coordinate of  $\mathbf{x}_n$  in the low-dimensional space is  $[a_i] = (\mathbf{x}_n^T \mathbf{u}_i - \bar{\mathbf{x}}^T \mathbf{u}_i)$ , for  $1 \leq i \leq M$
- The reconstructed data point in the original space is  $\tilde{\mathbf{x}}_n$
- Storage space before PCA:  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$
- Storage space using PCA:  $\bar{\mathbf{x}}, \mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_M$ , the coordinate of each data point

# Applications of PCA: Data compression

- An original example and its PCA reconstructions with different values of  $M$



An original example from the off-line digits data set together with its PCA reconstructions obtained by retaining  $M$  principal components for various values of  $M$ . As  $M$  increases the reconstruction becomes more accurate and would become perfect when  $M = D = 28 \times 28 = 784$ .

$$\tilde{\mathbf{x}}_n = \bar{\mathbf{x}} + \sum_{i=1}^M (\mathbf{x}_n^T \mathbf{u}_i - \bar{\mathbf{x}}^T \mathbf{u}_i) \mathbf{u}_i$$

# Applications of PCA: Data normalization

- PCA can make a more substantial normalization of data to give it zero mean and unit covariance
  - Different input variables become decorrelated
- This processing is called **whitening** or **sphering**
- 1. Solve the eigenvalues and eigenvectors of the covariance matrix  $S$

$$SU = UL$$

- where  $L$  is a  $D \times D$  diagonal matrix with elements  $\{\lambda_i\}$
- $U$  is a  $D \times D$  orthogonal matrix with columns given by  $\{\mathbf{u}_i\}$



# Applications of PCA: Data normalization


- 2. Transform each data point  $\mathbf{x}_n$  to

$$\mathbf{y}_n = \mathbf{L}^{-1/2} \mathbf{U}^T (\mathbf{x}_n - \bar{\mathbf{x}})$$

- The set  $\{\mathbf{y}_n\}$  has **zero mean**
- The set  $\{\mathbf{y}_n\}$  has an **identity covariance matrix**

$$\begin{aligned} \frac{1}{N} \sum_{n=1}^N \mathbf{y}_n \mathbf{y}_n^T &= \frac{1}{N} \sum_{n=1}^N \mathbf{L}^{-1/2} \mathbf{U}^T (\mathbf{x}_n - \bar{\mathbf{x}}) (\mathbf{x}_n - \bar{\mathbf{x}})^T \mathbf{U} \mathbf{L}^{-1/2} \\ &= \mathbf{L}^{-1/2} \mathbf{U}^T \mathbf{S} \mathbf{U} \mathbf{L}^{-1/2} = \mathbf{L}^{-1/2} \mathbf{L} \mathbf{L}^{-1/2} = \mathbf{I}. \end{aligned}$$


# Applications of PCA: Data normalization



- Left: Original data in a two-dimensional space
- Center: Two principle components and the corresponding eigenvalues
- Right: The whitening results
  - Zero mean and unit covariance

# PCA vs. FLD

- Comparisons between PCA and the Fisher linear discriminant
  - Both methods can be viewed as techniques for linear dimensionality reduction
  - PCA is unsupervised and depends only on the input data  $\{\mathbf{x}_n\}$
  - FLD depends on both data  $\{\mathbf{x}_n\}$  and class-label information  $\{t_n\}$
- Line: 1D subspace by PCA
  - Maximum variance
  - Minimum error
- Line: 1D subspace by FLD
  - Maximum data separation




# Multidimensional scaling (MDS)

- Given the **pairwise distances** between a set of  $N$  data, namely,  $\{\Delta_{ij}\}_{i,j=1}^N$ , we would like to seek  **$N$  vectors**  $\{\mathbf{y}_i\}_{i=1}^N$  so that  $\|\mathbf{y}_i - \mathbf{y}_j\| \approx \Delta_{ij}$

- Input

$$\begin{bmatrix} 0 & \Delta_{12} & \Delta_{13} & \Delta_{14} \\ \Delta_{12} & 0 & \Delta_{23} & \Delta_{24} \\ \Delta_{13} & \Delta_{23} & 0 & \Delta_{34} \\ \Delta_{14} & \Delta_{24} & \Delta_{34} & 0 \end{bmatrix}$$

- Output

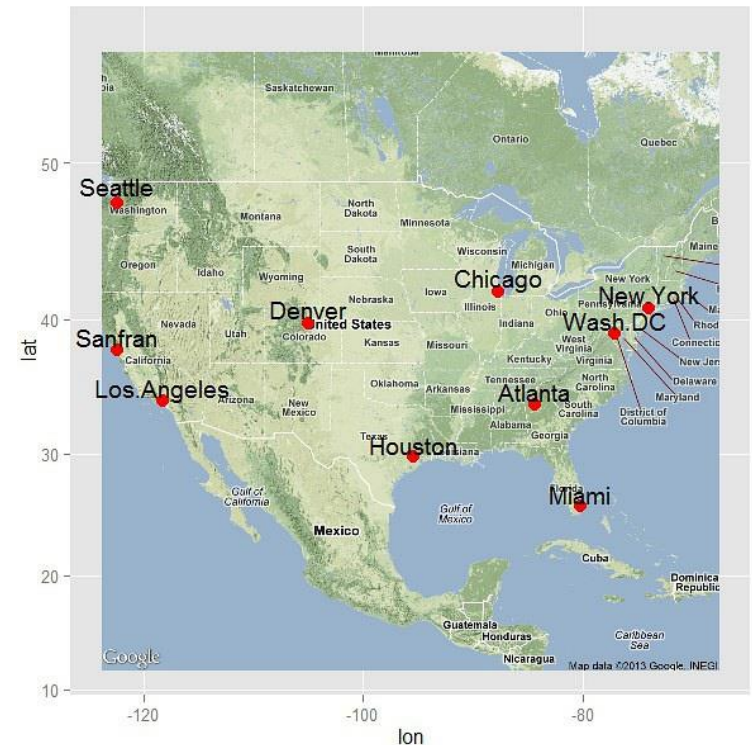


$$\|\mathbf{y}_i - \mathbf{y}_j\| \approx \Delta_{ij}$$

# An example of MDS



- The pairwise distances between 10 airports in US
- Airports' locations in the map

	A	B	C	D	E	F	G	H	I	J	K
1	ATL	ORD	DEN	HOU	LAX	MIA	JFK	SFO	SEA	IAD	
2	ATL	0	587	1212	701	1936	604	748	2139	2182	543
3	ORD	587	0	920	940	1745	1188	713	1858	1737	597
4	DEN	1212	920	0	879	831	1726	1631	949	1021	1494
5	HOU	701	940	879	0	1374	968	1420	1645	1891	1220
6	LAX	1936	1745	831	1374	0	2339	2451	347	959	2300
7	MIA	604	1188	1726	968	2339	0	1092	2594	2734	923
8	JFK	748	713	1631	1420	2451	1092	0	2571	2408	205
9	SFO	2139	1858	949	1645	347	2594	2571	0	678	2442
10	SEA	2182	1737	1021	1891	959	2734	2408	678	0	2329
11	IAD	543	597	1494	1220	2300	923	205	2442	2329	0




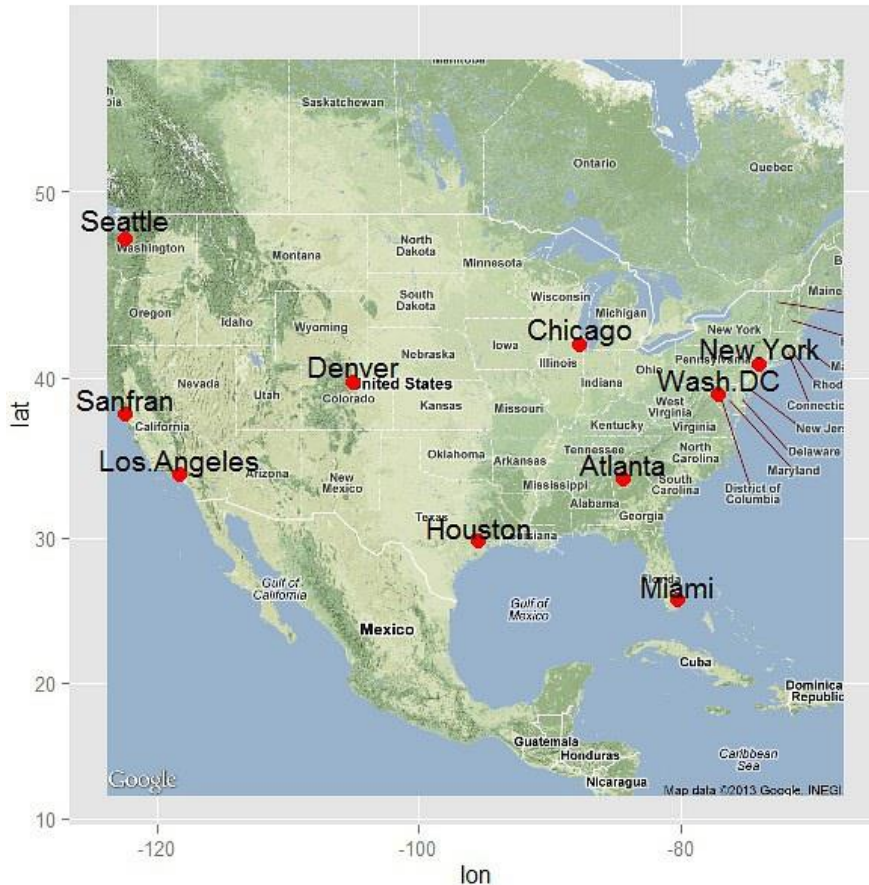
# An example of MDS

- The recovered 2D coordinates of those airports by MDS
- Horizontally mirror the map



# An example of MDS

- MDS can find the locations of these airports



# MDS: Two observations

- Observation 1: We can convert pairwise distances to pairwise inner products under some constraints
- Observation 2: The coordinate matrix can be derived from a Gram matrix,  $G_{ij} = \mathbf{x}_i \cdot \mathbf{x}_j$
- If  $\{\Delta_{ij}\}_{i,j=1}^N$  denote the pairwise Euclidean distances between zero mean vectors, then their inner products are

$$G_{ij} = \frac{1}{2} \left[ \frac{1}{N} \sum_{k=1}^N (\Delta_{ik}^2 + \Delta_{kj}^2) - \Delta_{ij}^2 - \frac{1}{N^2} \sum_{k,l=1}^N \Delta_{kl}^2 \right]$$






# Distance matrix to Gram matrix

$$G_{ij} = \frac{1}{2} \left[ \frac{1}{N} \sum_{k=1}^N (\Delta_{ik}^2 + \Delta_{kj}^2) - \Delta_{ij}^2 - \frac{1}{N^2} \sum_{k,l=1}^N \Delta_{kl}^2 \right]$$

- An example: three 1D points:



- Distance matrix:

$$[\Delta_{ij}] = \begin{bmatrix} 0 & 1 & 5 \\ 1 & 0 & 4 \\ 5 & 4 & 0 \end{bmatrix} \Rightarrow [\Delta_{ij}^2] = \begin{bmatrix} 0 & 1 & 25 \\ 1 & 0 & 16 \\ 25 & 16 & 0 \end{bmatrix}$$

- Inner product:

$$\begin{aligned} G_{12} &= \frac{1}{2} \left[ \frac{1}{3} \sum_k (\Delta_{1k}^2 + \Delta_{k2}^2) - \Delta_{12}^2 - \frac{1}{9} \sum_{k,l} \Delta_{kl}^2 \right] \\ &= \frac{1}{2} \left[ \frac{1}{3} (0 + 1 + 25 + 1 + 0 + 16) - 1 \right. \\ &\quad \left. - \frac{1}{9} (0 + 1 + 25 + 1 + 0 + 16 + 25 + 16 + 0) \right] \\ &= 2 \end{aligned}$$



# MDS: Objective and optimization

- Objective of MDS

$$\text{err}(\mathbf{y}) = \sum_{i,j} \left( G_{ij} - \mathbf{y}_i \cdot \mathbf{y}_j \right)^2$$

- Perform eigen-decomposition for Gram matrix  $G$

$$G = \sum_{\alpha=1}^N \lambda_{\alpha} \mathbf{v}_{\alpha} \mathbf{v}_{\alpha}^T \quad \text{where} \quad \lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_N \geq 0$$

- Optimal (low-rank) approximation

$$\mathbf{y}_{i\alpha} = \sqrt{\lambda_{\alpha}} v_{\alpha i} \quad \text{for} \quad \alpha = 1, 2, \dots, M$$


- Each data vector is constructed by **scaled, truncated eigenvectors**

# MDS: Embedding

- Eigen-decomposition

$$G = \sum_{\alpha=1}^N \lambda_{\alpha} \mathbf{v}_{\alpha} \mathbf{v}_{\alpha}^T \quad \text{where} \quad \lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_N \geq 0$$

- Embedding of the  $N$  data points



# MDS: Interpretation

- Embedding

$$\mathbf{y}_{i\alpha} = \sqrt{\lambda_\alpha} v_{\alpha i} \quad \text{for } \alpha = 1, 2, \dots, M$$

- Eigenvectors
  - Ordered, scaled, and truncated to yield low-dimensional embedding
- Eigenvalues
  - Measure how each dimension contributes to dot products
- Estimated dimensionality
  - Number of significant (nonnegative) eigenvalues

# PCA vs. MDS

- Perform eigen-decomposition
  - PCA: covariance matrix of size  $D \times D$
  - MDS: Gram matrix of size  $N \times N$
- Both PCA and MDS consider the  $M$  largest eigenvalues
- Eigenvectors in PCA serve as the projection matrix
- Scaled eigenvectors in MDS are used for embedding, the projected data vectors

# Summary of MDS

- Convert the distance matrix to Gram matrix

$$G_{ij} = \frac{1}{2} \left[ \frac{1}{N} \sum_{k=1}^N (\Delta_{ik}^2 + \Delta_{kj}^2) - \Delta_{ij}^2 - \frac{1}{N^2} \sum_{k,l=1}^N \Delta_{kl}^2 \right]$$

- Perform eigen-decomposition of the Gram matrix


$$G = \sum_{\alpha=1}^N \lambda_{\alpha} \mathbf{v}_{\alpha} \mathbf{v}_{\alpha}^T \quad \text{where} \quad \lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_N \geq 0$$

- Embedding via low rank approximation

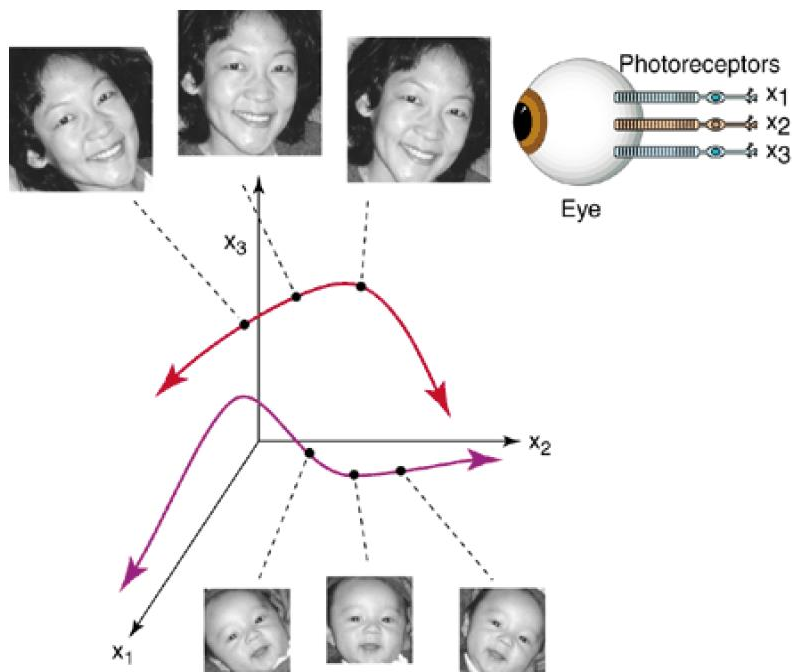
$$\mathbf{y}_{i\alpha} = \sqrt{\lambda_{\alpha}} v_{\alpha i} \quad \text{for} \quad \alpha = 1, 2, \dots, M$$

# Manifold (流形)

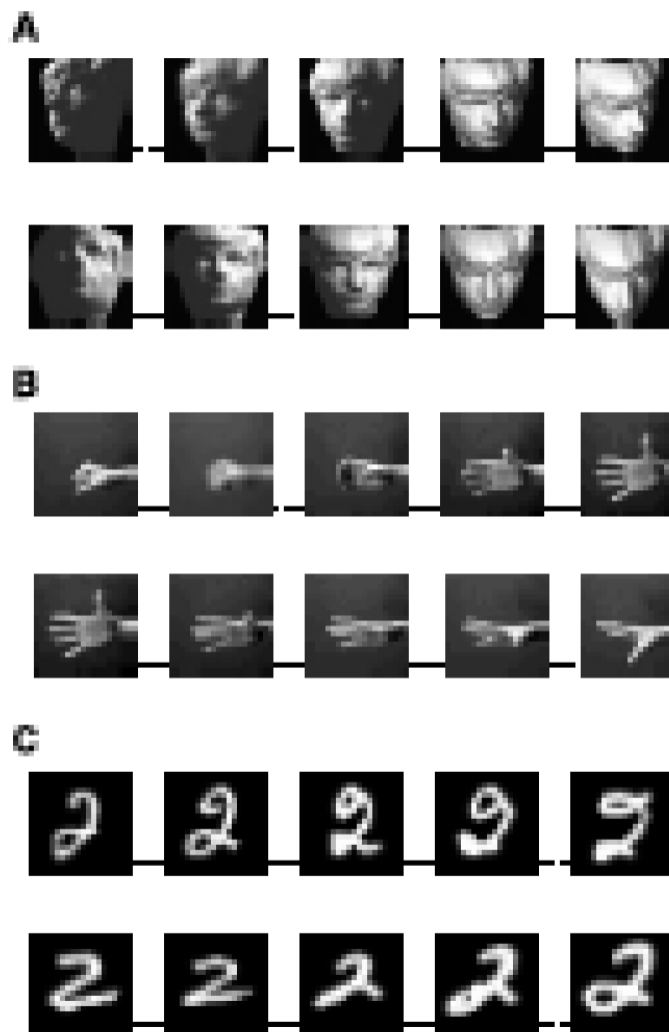
- In mathematics, a manifold is a topological space that locally resembles Euclidean space near each point
  - In the neighborhood of any point on a manifold, the space behaves just like it would in the neighborhood of any point in some  $n$ -dimensional Euclidean space
- A sphere in 3-dimensional Euclidean space
  - A very small patch from the sphere looks just like a piece of 2-dimensional Euclidean space
  - 2-manifold



# Image manifolds



Seung & Lee, 2000  
Tenebaum et al., 2000





# Manifold learning

- Given **high-dimensional data** sampled from a **low-dimensional manifold**, how to compute a faithful embedding?




# Isometric mapping (Isomap)

[Tenenbaum et al. Science 2000]

- Input: a data set of observations  $\{\mathbf{x}_n\}$  where  $n = 1, 2, \dots, N$  and  $\mathbf{x}_n$  in a  $D$ -dimensional space
- Output: a low-dimensional ( $M$ -dimensional) data representation  $\mathbf{y}_n$  for each input data point  $\mathbf{x}_n$ , where  $M \ll D$
- Goal: The **geodesic distance** in the original space can be preserved in the low-dimensional space
  - Nonlinear dimensionality reduction



# Geodesic distance

- The **geodesic distance** between two points on a manifold is the length of **the shortest connecting path along the manifold**
- Figure A: the geodesic distance between two points (circles)
- Figure B: An approximation of the geodesic distance based on sampled points
- Figure C: The outputs of Isomap, where geodesic distances can be approximately measured by Euclidean distances




# Step 1: Build an adjacency graph

- Adjacency graph
  - Vertices represent inputs
  - Undirected edges connect neighbors
- Neighborhood selection
  - Many options:
    - ◆  $k$  nearest neighbors
    - ◆  $\varepsilon$ -ball
    - ◆ Prior knowledge
  - Weight edges by local Euclidean distances



## Step 2: Estimate geodesics

- Dynamic programming
  - Weight edges by local Euclidean distances
  - Compute **shortest paths** through graph
- Dijkstra's algorithm for computing the shortest path between a pair of vertices
- Floyd's algorithm for computing the pair-wise shortest paths




## Step 3: Apply MDS

- Embedding
  - Given the pair-wise distances between data points, we can apply MDS to recover the data vectors of these points in a low-dimensional space where the distances are preserved

# Algorithm of Isomap

- Algorithm
  - (1)  $k$  nearest neighbors
  - (2) Shortest paths through graph (estimating geodesic distances)
  - (3) MDS on geodesic distances


$$\begin{bmatrix} 0 & \Delta_{12} & \Delta_{13} & \Delta_{14} \\ \Delta_{12} & 0 & \Delta_{23} & \Delta_{24} \\ \Delta_{13} & \Delta_{23} & 0 & \Delta_{34} \\ \Delta_{14} & \Delta_{24} & \Delta_{34} & 0 \end{bmatrix}$$



- Impact
  - Much simpler than earlier algorithms for manifold learning
  - Does it work?

# Visualization

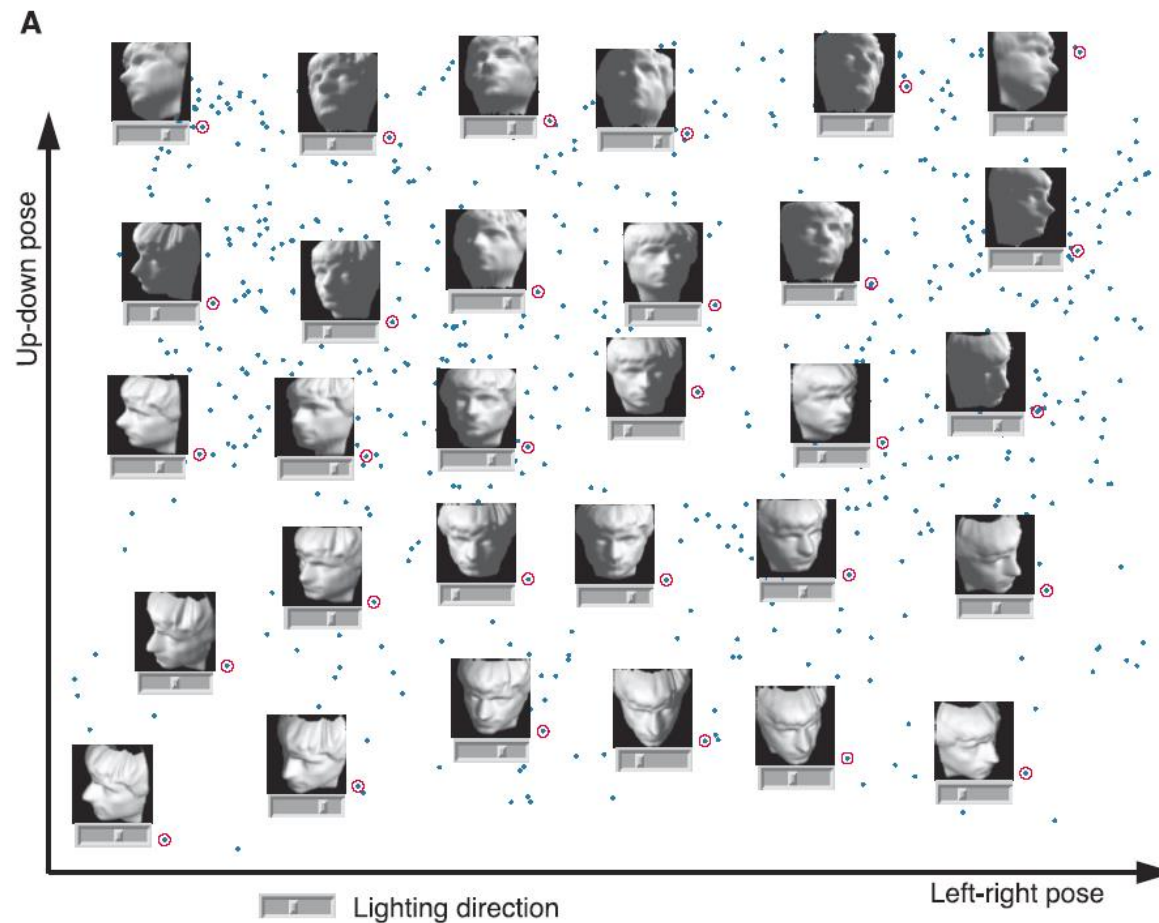
- Swiss roll dataset
  - 1024 data points in a 3D space
  - $k=12$  in the adjacency graph construction
  - 2-dimensional space where the Euclidean distance between a pair of data points approximates their geodesic distance on the manifold





# Visualization

- Face images
  - 64 x 64 (4096)-dim, 698 data points,  $k = 6$





# Locally linear embedding (LLE) [Roweis and Saul, Science 2000]

- MDS and Isomap
  - Preserve **global pairwise distances**
  - Construct large and dense matrices
  - Compute top eigenvectors
- Locally linear embedding (LLE)
  - Preserve **local geometric relationships**
  - Construct large and sparse matrices
  - Compute bottom eigenvectors

# How to exploit local linearity?

- Manifolds are globally nonlinear, but locally linear
- Map the inputs into a single continuous global coordinate system of lower dimensionality
  - Think globally, fit locally



# Locally linear embedding

- Input: a data set of observations  $\{\mathbf{x}_i\}$  where  $i = 1, 2, \dots, N$  and  $\mathbf{x}_i$  in a  $D$ -dimensional space
- Output: a low-dimensional ( $M$ -dimensional) data representation  $\mathbf{y}_i$  for each input data point  $\mathbf{x}_i$ , where  $M \ll D$
- Goal: The neighborhood of each point, i.e., local geometric structure, in the original space is preserved in the low-dimensional space
  - Nonlinear dimensionality reduction

# LLE algorithm


- Steps
  - (1) Nearest neighbor search for each input data point
  - (2) Least squares fitting for neighborhood representation
  - (3) Eigen-decomposition for embedding

# Step 1: Identify neighbors

- Identify the neighbors of each data point  $\mathbf{x}_i$ 
  - $k$  nearest neighbors
  - $\varepsilon$ -ball
  - Prior knowledge
- Assumptions
  - Data are densely sampled from a manifold
  - Sufficient neighboring data can be retrieved for each data point


# Neighborhood graph

- Assumption
  - Neighborhoods on the graph correspond to neighborhoods on the manifold



## Step 2: Compute reconstruction weights

- Characterize local geometry of each neighborhood by weights  $\{W_{ij}\}$  for each data  $\mathbf{x}_i$



- Compute the weights by **linearly reconstructing** each point  $\mathbf{x}_i$  from its neighbors



# Linear reconstructions


- Local linearity

- Neighbors lie on a **locally linear patch** of a manifold

- Reconstruction errors

- Least squared errors should be small

$$\Phi(W) = \sum_i \left| \mathbf{x}_i - \sum_j W_{ij} \mathbf{x}_j \right|^2$$



# Least squares fitting

- Objective

$$\Phi(W) = \sum_i \left| \mathbf{x}_i - \sum_j W_{ij} \mathbf{x}_j \right|^2$$

- Constraints

- Nonzero  $W_{ij}$  only if  $\mathbf{x}_j$  is a neighbor of  $\mathbf{x}_i$
- Weights must sum to one for each  $\mathbf{x}_i$ :  $\sum_j W_{ij} = 1$

- Solver: least squares fitting

- Local invariance:

- Optimal weights  $\{W_{ij}\}$  are invariant to rotation, translation, and scaling

## Step 3: Complete embedding

- Low dimensional representation
  - $\mathbf{x}_i$  in a  $D$ -dimensional space  $\rightarrow$   $\mathbf{y}_i$  in an  $M$ -dimensional space

- Objective: Minimize the reconstruction errors

$$\Psi(\mathbf{y}) = \sum_i \left| \mathbf{y}_i - \sum_j W_{ij} \mathbf{y}_j \right|^2$$

- Constraints

- The mean of the projected data is on the origin:  $\sum_i \mathbf{y}_i = \mathbf{0}$

- Impose unit covariance matrix:  $\frac{1}{N} \sum_{i=1}^N \mathbf{y}_i \mathbf{y}_i^T = \mathbf{I}$

# Solving an eigenvalue problem

- Quadratic form

$$\Psi(\mathbf{y}) = \sum_{i,j} (\mathbf{y}_i \cdot \mathbf{y}_j) \Psi_{ij} \quad \text{with } \Psi = (I - W)^T(I - W)$$

- The optimal embedding is given by the bottom  $M + 1$  eigenvectors
- Solution
  - Discard the bottom eigenvector  $[1 \ 1 \ \dots \ 1]^T$
  - The other eigenvectors are used to yield the low-dimensional data  $\{\mathbf{y}_i\}$



# Embedding

- After discarding the eigenvector  $[1 \ 1 \ \cdots \ 1]^T$ , the low-dimensional data  $\{\mathbf{y}_i\}$  are given below

	$\mathbf{v}_1$	$\mathbf{v}_2$		$\mathbf{v}_M$
$\mathbf{y}_1$	■	■	...	■
$\mathbf{y}_2$	■	■		■
	■	■		■
	■	■		■
	■	■		■
	■	■		■
	■	■		■
$\mathbf{y}_N$	■	■		■


# Summary of LLE

- Three steps
  - 1. Get  $k$ -nearest neighbors
  - 2. Compute weights
  - 3. Complete embedding

- Optimizations

$$\Phi(W) = \sum_i \left| \mathbf{x}_i - \sum_j W_{ij} \mathbf{x}_j \right|^2$$

$$\Psi(\mathbf{y}) = \sum_i \left| \mathbf{y}_i - \sum_j W_{ij} \mathbf{y}_j \right|^2$$

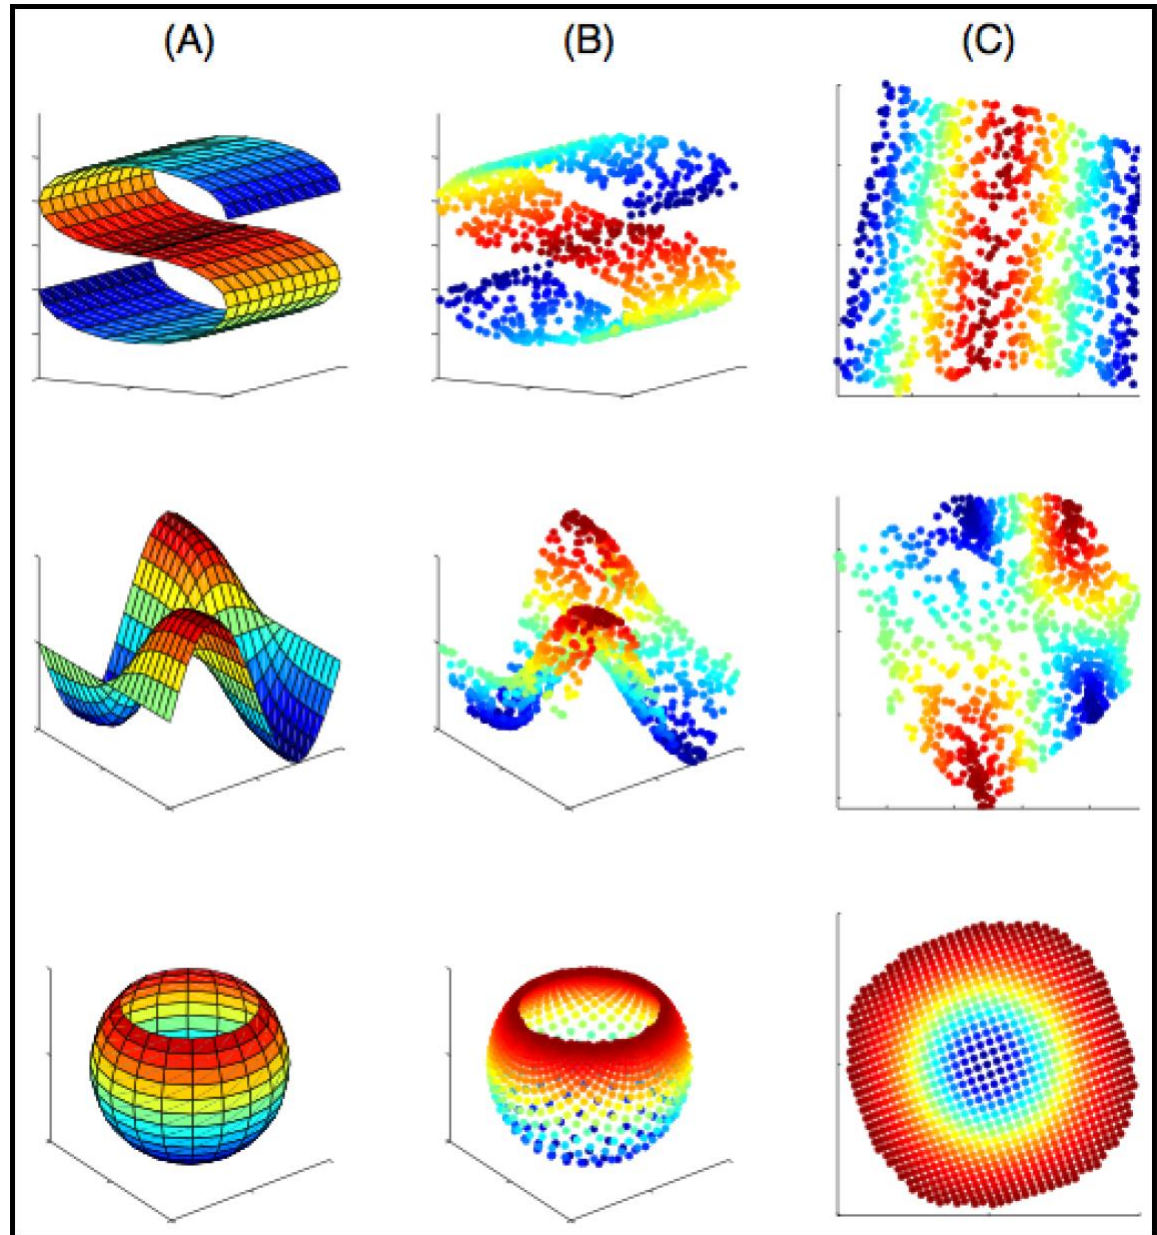


# Surfaces

$m = 1000$   
inputs

$k = 8$   
nearest  
neighbors

$D = 3$   
 $d = 2$   
dimensions




# Translated faces

$m = 961$   
images

$k = 4$   
nearest  
neighbors

$D = 3009$   
pixels

$d = 2$   
manifold






# Pose and expression

$m = 1965$   
images

$k = 12$   
nearest  
neighbors

$D = 560$   
pixels

$d = 2$   
(shown)



# Isomap vs. LLE

- Similarities
  - Nonlinear dimensionality reduction for manifold learning
  - Graph-based, spectral method
  - No local minima
  - Does not estimate dimensionality
- Differences
  - Constructs dense vs. sparse matrices
  - Preserves distances vs. local geometric structure

# References

- PCA
  - Chapter 12.1 in the PRML textbook
- MDS
  - [https://en.wikipedia.org/wiki/Multidimensional\\_scaling](https://en.wikipedia.org/wiki/Multidimensional_scaling)
- Isomap
  - J. B. Tenenbaum, V. de Silva and J. C. Langford. A global geometric framework for nonlinear dimensionality reduction. Science 290 (5500): 2319-2323, 2000.
- LLE
  - S. Roweis and L. Saul. Nonlinear dimensionality reduction by locally linear embedding. Science 290 (5500): 2323-2326 2000.

# Thank You for Your Attention!

THANK YOU FOR YOUR ATTENTION!

Yen-Yu Lin (林彥宇)

Email: [lin@cs.nctu.edu.tw](mailto:lin@cs.nctu.edu.tw)

URL: <https://www.cs.nctu.edu.tw/members/detail/lin>

